

Machine Learning

Ensembles: Bagging

Ensembles: Gradient Boosting

Ensembles: Ada Boost

Clustering

Ensemble methods

Why learn one classifier when you can learn many?

Ensemble: combine many predictors

- (Weighted) combinations of predictors
- May be same type of learner or different



Who wants to be a millionaire?

Various options for getting help:



Simple ensembles

“Committees”

- Unweighted average / majority vote:
- Take several trained models, and report their average prediction


Weighted averages

- Up-weight “better” predictors
- Ex: Classes: +1 , -1 , weights alpha:

$$\hat{y}_1 = f_1(x_1, x_2, \dots)$$

$$\hat{y}_2 = f_2(x_1, x_2, \dots) \quad \Rightarrow \quad \hat{y}_e = \text{sign}(\sum \alpha_i \hat{y}_i)$$

...



“Stacked” ensembles

Train a “predictor of predictors”

- Treat individual predictors as features

$$\hat{y}_1 = f_1(x_1, x_2, \dots)$$

$$\hat{y}_2 = f_2(x_1, x_2, \dots) \quad \Rightarrow \quad \hat{y}_e = f_e(\hat{y}_1, \hat{y}_2, \hat{y}_3, \hat{y}_4, \dots)$$


...

- Similar to multi-layer perceptron idea
- For example, if f_e is a linear classifier:
 - Weighted vote: $\hat{y}_e = \text{sign}(\sum_i \alpha_i \hat{y}_i)$, but with learned weights
- Can train stacked learner f_e on validation data
 - Avoids giving high weight to overfit models

Mixtures of experts

Can make weights depend on x

- Weight $\alpha_z(x)$ indicates “expertise”
- Combine using weighted average (or even just pick largest)



Mixture of three linear predictor experts

Weighted average:

$$f(x; \omega, \theta) = \sum_z \alpha_z(x; \omega) f_z(x; \theta_z)$$

Weights: (multi) logistic regression

$$\alpha_z(x; \omega) = \frac{\exp(x \cdot \omega^z)}{\sum_c \exp(x \cdot \omega^c)}$$

If loss, learners, weights are all differentiable, can train jointly...

Machine Learning

Ensembles: Bagging

Ensembles: Gradient Boosting

Ensembles: Ada Boost

Clustering

Ensemble methods

Why learn one classifier when you can learn many?

- “Committee”: learn K classifiers, average their predictions

Bagging = bootstrap aggregation

- Learn many classifiers, each with only part of the data
- Combine through model averaging

Remember overfitting: “memorize” the data

- Used test data to see if we had gone too far
- Cross-validation
 - Make many splits of the data for train & test
 - Each of these defines a classifier
 - Typically, we use these to check for overfitting
 - Instead of checking if we overfit, we combine these classifiers to produce a better classifier



Bagging

Bootstrap

- Create a random subset of data by sampling
- Draw m' of the m samples, with replacement (some variants w/o)
 - Some data left out; some data repeated several times

Bagging

- Repeat K times
 - Create a training set of $m' < m$ examples
 - Train a classifier on the random training set
- To test, run each trained classifier
 - Each classifier votes on the output, take majority
 - For regression: each regressor predicts, take average


Some complexity control: harder for each to memorize data

- Each learner has data set where data points are missing, thus memorization is suppressed
- Doesn't work for linear models (average of linear functions is linear function...)
- Perceptrons OK (linear + threshold = nonlinear)

Bias / variance

“The world”

Data we observe




Test
Error

We only see a little bit of data

Can decompose error into two parts

- Bias – error due to model choice
 - Inability of model class to represent the best function
 - Gets better with more model complexity
- Variance – randomness due to data size
 - Better w/ more data, worse w/ complexity
- *Bagging reduces variance and increases bias*









Bagged decision trees

Randomly resample data

Learn a decision tree for each

- No max depth = very flexible class of functions
- Learner is low bias, but high variance
- Sampling:
 - simulates “equally likely” data sets we **could have** observed
 - train decision tree on every sampled data set




Bagged decision trees

Average over collection

- Classification: majority vote

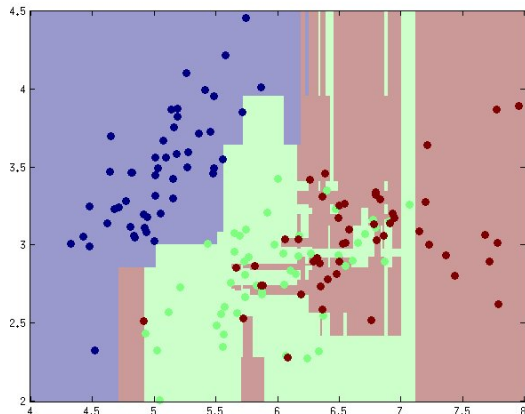
Reduces memorization effect

- Not every predictor sees each data point
- Lowers effective “complexity” of the overall average
- Usually, better generalization performance
- Intuition: reduces variance while keeping bias low




Full data set


Avg of 5 trees



Avg of 25 trees



Avg of 100 trees



Bagging in Python

```
# Load data set X, Y for training the ensemble...
m,n = X.shape
classifiers = [ None ] * nBag           # Allocate space for learners
for i in range(nBag):
    ind = np.floor( m * np.random.rand(nUse)) # Bootstrap sample a data set:
    Xi, Yi = X[ind,:], Y[ind]                # select the data at those indices
    classifiers[i] = ml.MyClassifier(Xi, Yi)   # Train a model on data Xi, Yi
```

```
# test on data Xtest
mTest = Xtest.shape[0]
predict = np.zeros( (mTest, nBag) )      # Allocate space for predictions from each model
for i in range(nBag):
    predict[:,i] = classifiers[i].predict(Xtest) # Apply each classifier

# Make overall prediction by majority vote
predict = np.mean(predict, axis=1) > 0 # if +1 vs -1
```

Random forests

Bagging applied to decision trees

Problem

- With lots of data, we usually learn the same classifier
- Averaging over these doesn't help!

Introduce extra variation in learner

- At each step of training, only allow a subset of features
- Enforces diversity (“best” feature not available)
- Keeps bias low (every feature available eventually)
- Average over these learners (majority vote)

```
# in FindBestSplit(X,Y):  
for each of a subset of features  
    for each possible split  
        Score the split (e.g. information gain)  
Pick the feature & split with the best score  
Recurse on left & right splits
```

Summary

Ensembles: collections of predictors

- Combine predictions to improve performance

Bagging

- “Bootstrap aggregation”
- Reduces complexity of a model class prone to overfit
- In practice: Resample the data many times
 - For each, generate a predictor on that resampling
- Plays on bias / variance trade off
- **Price:** more computation per prediction

Machine Learning

Ensembles: Bagging

Ensembles: Gradient Boosting

Ensembles: Ada Boost

Clustering

Ensembles


Weighted combinations of predictors

“Committee” decisions

- Trivial example
- Equal weights (majority vote / unweighted average)
- Might want to weight unevenly – up-weight better predictors

Bagging

- Bootstrapping: subsampling with replacement
- Train classifiers on bootstraps, average the results
- Reduces complexity of a model class prone to overfit
 - Plays on bias / variance trade off



Ensembles


Weighted combinations of predictors

“Committee” decisions

- Trivial example
- Equal weights (majority vote / unweighted average)
- Might want to weight unevenly – up-weight better predictors

Boosting

- Focus new learners on examples that others get wrong
- Train learners sequentially
- Errors of early predictions indicate the “hard” examples
- Focus later predictions on getting these examples right
- Combine the whole set in the end
- Convert many “weak” learners into a complex predictor



Gradient boosting

Learn a regression predictor

$$f_1(x^{(i)}) \approx y^{(i)}$$

Compute the error residual


$$\epsilon^{(i)} = y^{(i)} - f_1(x^{(i)})$$

Learn to predict the residual

$$f_2(x^{(i)}) \approx \epsilon^{(i)}$$


Learn a simple predictor...

$$f_1(x^{(i)}) \approx y^{(i)}$$



Then try to correct its errors

$$\epsilon^{(i)} = y^{(i)} - f_1(x^{(i)})$$



Gradient boosting

Learn a regression predictor $f_1(x^{(i)}) \approx y^{(i)}$

Compute the error residual $\epsilon^{(i)} = y^{(i)} - f_1(x^{(i)})$



Learn to predict the residual $f_2(x^{(i)}) \approx \epsilon^{(i)}$

Combining gives a better predictor...

$$\Rightarrow f_1(x^{(i)}) + f_2(x^{(i)}) \approx y^{(i)}$$

Can try to correct its errors also, & repeat

$$\epsilon_2^{(i)} = y^{(i)} - f_1(x^{(i)}) - f_2(x^{(i)}) \quad \dots$$



Gradient boosting


Learn sequence of predictors

Sum of predictions is increasingly accurate


Predictive function is increasingly complex

$$y^{(i)} \approx \sum_z f_z(x^{(i)})$$

Data & prediction function



Error residual



...

Gradient boosting

Make a set of predictions $\hat{y}[i]$

The “error” in our predictions is $J(y, \hat{y})$

- For MSE: $J(.) = \sum (y[i] - \hat{y}[i])^2$

We can “adjust” \hat{y} to try to reduce the error

- $\hat{y}[i] = \hat{y}[i] + \alpha f[i]$
- $f[i] \approx \nabla J(y, \hat{y}) = (y[i] - \hat{y}[i])$ for MSE

Each learner is estimating the gradient of the loss function

Gradient descent: take sequence of steps to reduce J

- Sum of predictors, weighted by step size α

Gradient boosting in Python

```
# Load data set X, Y ...
learner = [None] * nBoost # storage for ensemble of models
alpha = [1.0] * nBoost # and weights of each learner
mu = Y.mean()           # often start with constant "mean" predictor
dY = Y - mu             # subtract this prediction away
for k in range( nBoost ):
    learner[k] = ml.MyRegressor( X, dY )    # regress to predict residual dY using X
    alpha[k] = 1.0      # alpha: "learning rate" or "step size"
    # smaller alphas need more classifiers, but may predict better given enough of them

    # compute the residual given our new prediction:
    dY = dY - alpha[k] * learner[k].predict(X)
```

```
# test on data Xtest
mTest = Xtest.shape[0]
predict = np.zeros( (mTest,) ) + mu # Allocate space for predictions & add 1st (mean)
for k in range(nBoost):
    predict += alpha[k] * learner[k].predict(Xtest) # Apply next predictor & accum
```

Summary

Ensemble methods

- Combine multiple classifiers to make “better” one
- Committees, average predictions
- Can use weighted combinations
- Can use same or different classifiers

Gradient Boosting

- Use a simple regression model to start
- Subsequent models predict the error residual of the previous predictions
- Overall prediction given by a weighted sum of the collection

Demo Time

http://arogozhnikov.github.io/2016/06/24/gradient_boosting_explained.html

Machine Learning

Ensembles: Bagging

Ensembles: Gradient Boosting

Ensembles: Ada Boost

Clustering

Ensembles

Weighted combinations of classifiers

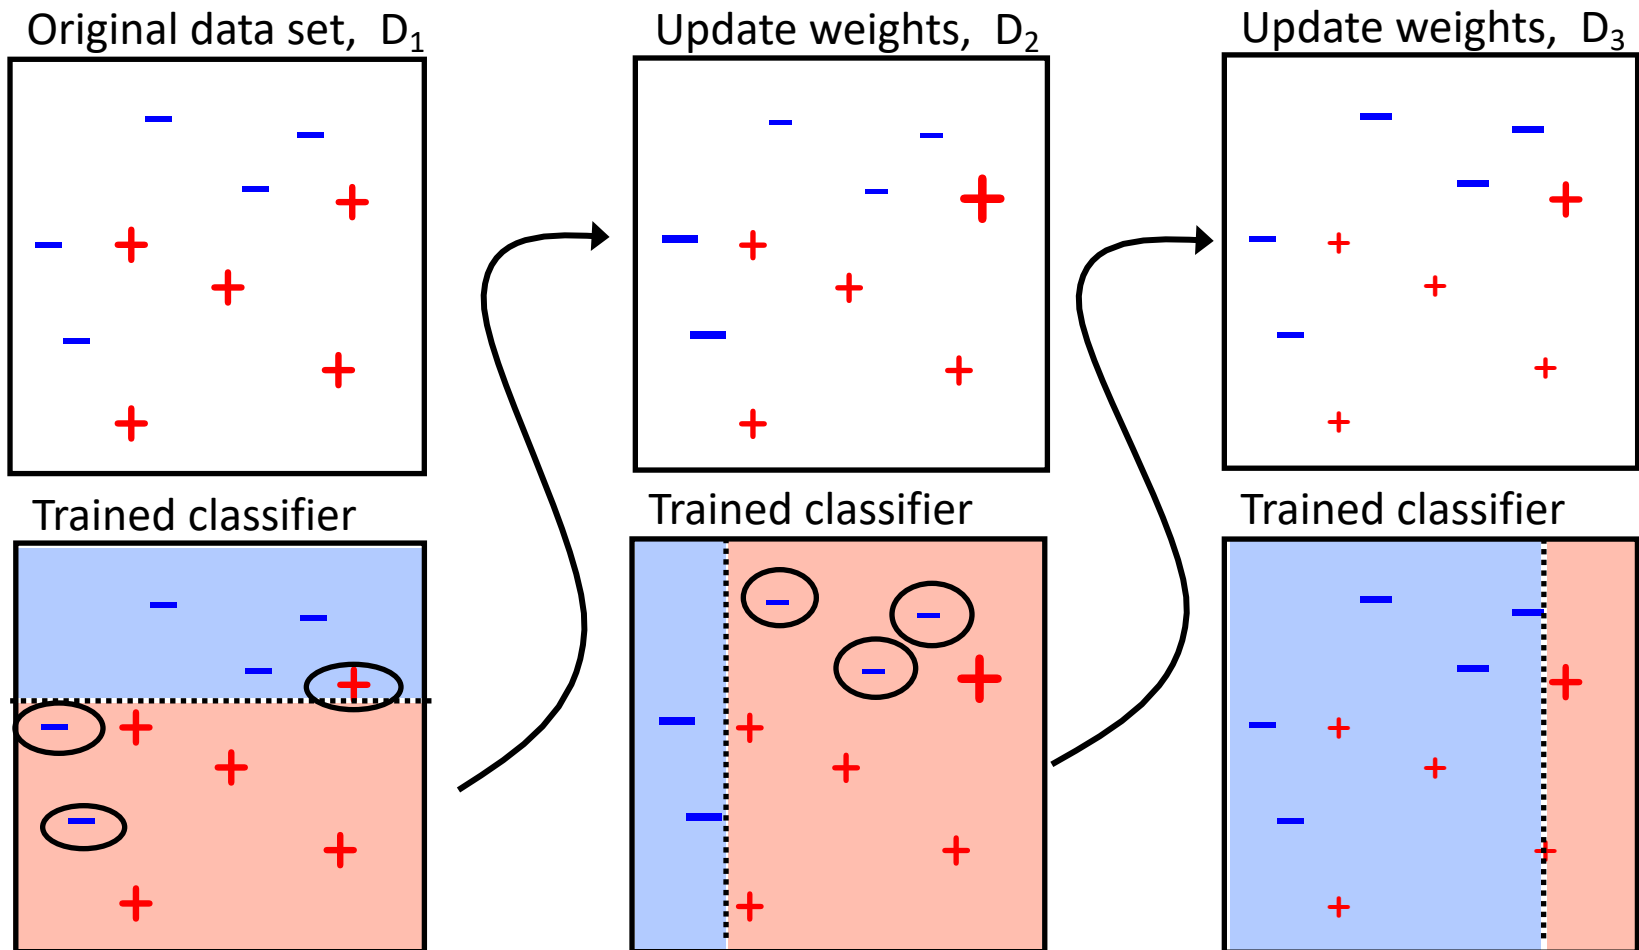
“Committee” decisions

- Trivial example
- Equal weights (majority vote)
- Might want to weight unevenly – up-weight good experts

Boosting

- Focus new experts on examples that others get wrong
- Train experts sequentially
- Errors of early experts indicate the “hard” examples
- Focus later classifiers on getting these examples right
- Combine the whole set in the end
- Convert many “weak” learners into a complex classifier

Boosting example



Minimizing Weighted Error

So far we've mostly minimized unweighted error

Minimizing weighted error is no harder!

Unweighted average loss:

$$J(\theta) = \frac{1}{m} \sum_i J_i(\theta, x^{(i)})$$

Weighted average loss:

$$J(\theta) = \sum_i w_i J_i(\theta, x^{(i)})$$

For any loss (logistic MSE, hinge, ...)

$$J(\theta, x^{(i)}) = (\sigma(\theta x^{(i)}) - y^{(i)})^2$$

$$J(\theta, x^{(i)}) = \max[0, 1 - y^{(i)} \theta x^{(i)}]$$

For e.g. decision trees, compute **weighted** information gain:

$p(+1) \propto$ total weight of data with class +1


$p(-1) \propto$ total weight of data with class -1 $\Rightarrow H(p) = \text{entropy}$

Boosting example

Weight each classifier and combine them:

$$.33 * \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \\ \hline \end{array} + .57 * \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \\ \hline \end{array} + .42 * \begin{array}{|c|} \hline \text{blue} \\ \hline \text{red} \\ \hline \end{array} \geq 0$$

Combined classifier



1-node decision trees
“decision stumps”
very simple classifiers

AdaBoost: “Adaptive boosting”

```
# Load data set X, Y ... ; Y assumed +1 / -1
for i in range(nBoost):
    learner[i] = ml.MyClassifier( X,Y, weights=wt ) # train a weighted classifier
    Yhat = learner[i].predict(X)
    e = wt.dot( Y != Yhat ) # compute weighted error rate
    alpha[i] = 0.5 * np.log( (1-e)/e )
    wt *= np.exp( -alpha[i] * Y * Yhat ) # update weights
    wt /= wt.sum() # and normalize them
```

Notes

- $e > .5$ means classifier is not better than random guessing
- if $Y == Yhat$, $Y * Yhat > 0$ and weights decrease
- Otherwise, they increase

```
# Final classifier:
predict = np.zeros( (mTest,) )
for i in range(nBoost):
    predict += alpha[i] * learner[i].predict(Xtest) # compute contribution of each
predict = np.sign(predict) # and convert to +1 / -1 decision
```

Summary

Ensemble methods

- Combine multiple classifiers to make “better” one
- Committees, majority vote
- Weighted combinations
- Can use same or different classifiers

Boosting

- Train sequentially; later predictors focus on mistakes by earlier

Boosting for classification (e.g., [AdaBoost](#))

- Use results of earlier classifiers to know what to work on
- Weight “hard” examples so we focus on them more
- Example: Viola-Jones for face detection