

```

# to prevent unnecessary warnings
import warnings
warnings.simplefilter(action='ignore', category=FutureWarning)

# TensorFlow and tf.keras
import tensorflow as tf

# Helper libraries
import numpy as np
import matplotlib.pyplot as plt
import os
import subprocess
import cv2
import json
import requests
from tqdm import tqdm

print(tf.__version__)

2.8.2

fashion_mnist = tf.keras.datasets.fashion_mnist
(train_images, train_labels), (test_images, test_labels) =
fashion_mnist.load_data()

class_names = ['T-shirt/top', 'Trouser', 'Pullover', 'Dress', 'Coat',
'Sandal', 'Shirt', 'Sneaker', 'Bag', 'Ankle boot']

print('\nTrain_images.shape: {}, of {}'.format(train_images.shape,
train_images.dtype))
print('Test_images.shape: {}, of {}'.format(test_images.shape,
test_images.dtype))

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/train-labels-idx1-ubyte.gz
32768/29515 [=====] - 0s 0us/step
40960/29515 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/train-images-idx3-ubyte.gz
26427392/26421880 [=====] - 0s 0us/step
26435584/26421880 [=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/t10k-labels-idx1-ubyte.gz
16384/5148
[=====]
=====] - 0s 0us/step
Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/t10k-images-idx3-ubyte.gz
4423680/4422102 [=====] - 0s 0us/step
4431872/4422102 [=====] - 0s 0us/step

```

```
Train_images.shape: (60000, 28, 28), of uint8
Test_images.shape: (10000, 28, 28), of uint8
```

#Reshape data

```
train_images_3ch = np.stack([train_images]*3, axis=-1)
test_images_3ch = np.stack([test_images]*3, axis=-1)

print('\nTrain_images.shape: {}, of {}'.format(train_images_3ch.shape,
train_images_3ch.dtype))
print('Test_images.shape: {}, of {}'.format(test_images_3ch.shape,
test_images_3ch.dtype))
```

#imagenet competition was for colored images hence 3 channels

```
Train_images.shape: (60000, 28, 28, 3), of uint8
Test_images.shape: (10000, 28, 28, 3), of uint8
```

#Resizing image

###The minimum image size expected by the VGG model is 32x32 so we need to resize our images

```
import cv2

def resize_image_array(img, img_size_dims):
    img = cv2.resize(img, dsize=img_size_dims,
        interpolation=cv2.INTER_CUBIC)
    img = np.array(img, dtype=np.float32)
    return img

%%time

IMG_DIMS = (32, 32)

train_images_3ch = np.array([resize_image_array(img,
img_size_dims=IMG_DIMS) for img in train_images_3ch])
test_images_3ch = np.array([resize_image_array(img,
img_size_dims=IMG_DIMS) for img in test_images_3ch])

print('\nTrain_images.shape: {}, of {}'.format(train_images_3ch.shape,
train_images_3ch.dtype))
print('Test_images.shape: {}, of {}'.format(test_images_3ch.shape,
test_images_3ch.dtype))

Train_images.shape: (60000, 32, 32, 3), of float32
Test_images.shape: (10000, 32, 32, 3), of float32
```

CPU times: user 1.5 s, sys: 663 ms, total: 2.16 s
Wall time: 2.21 s

#Build CNN Model

```
# define input shape
INPUT_SHAPE = (32, 32, 3)

# get the VGG19 model
vgg_layers = tf.keras.applications.vgg19.VGG19(weights='imagenet',
include_top=False,
input_shape=INPUT_SHAPE)
#we are adding our own layer
vgg_layers.summary()
#include_top=False = we will add our own last layer
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg19/vgg19_weights_tf_dim_ordering_tf_kernels_notop.h5
80142336/80134624 [=====] - 0s 0us/step
80150528/80134624 [=====] - 0s 0us/step
Model: "vgg19"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 32, 32, 3)]	0
block1_conv1 (Conv2D)	(None, 32, 32, 64)	1792
block1_conv2 (Conv2D)	(None, 32, 32, 64)	36928
block1_pool (MaxPooling2D)	(None, 16, 16, 64)	0
block2_conv1 (Conv2D)	(None, 16, 16, 128)	73856
block2_conv2 (Conv2D)	(None, 16, 16, 128)	147584
block2_pool (MaxPooling2D)	(None, 8, 8, 128)	0
block3_conv1 (Conv2D)	(None, 8, 8, 256)	295168
block3_conv2 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv3 (Conv2D)	(None, 8, 8, 256)	590080
block3_conv4 (Conv2D)	(None, 8, 8, 256)	590080
block3_pool (MaxPooling2D)	(None, 4, 4, 256)	0
block4_conv1 (Conv2D)	(None, 4, 4, 512)	1180160

block4_conv2 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv3 (Conv2D)	(None, 4, 4, 512)	2359808
block4_conv4 (Conv2D)	(None, 4, 4, 512)	2359808
block4_pool (MaxPooling2D)	(None, 2, 2, 512)	0
block5_conv1 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv2 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv3 (Conv2D)	(None, 2, 2, 512)	2359808
block5_conv4 (Conv2D)	(None, 2, 2, 512)	2359808
block5_pool (MaxPooling2D)	(None, 1, 1, 512)	0

```

=====
Total params: 20,024,384
Trainable params: 20,024,384
Non-trainable params: 0
=====

```

#Fine tuning

```

# Fine-tune all the layers
for layer in vgg_layers.layers:
    layer.trainable = True #we are training layers

# Check the trainable status of the individual layers
for layer in vgg_layers.layers:
    print(layer, layer.trainable)

<keras.engine.input_layer.InputLayer object at 0x7f0a33341450> True
<keras.layers.convolutional.Conv2D object at 0x7f0a25136d50> True
<keras.layers.convolutional.Conv2D object at 0x7f0a24e65510> True
<keras.layers.pooling.MaxPooling2D object at 0x7f0a25126d50> True
<keras.layers.convolutional.Conv2D object at 0x7f0a250e2250> True
<keras.layers.convolutional.Conv2D object at 0x7f0a250e5b50> True
<keras.layers.pooling.MaxPooling2D object at 0x7f0a25070350> True
<keras.layers.convolutional.Conv2D object at 0x7f0a2506cd10> True
<keras.layers.convolutional.Conv2D object at 0x7f0a32cf1850> True
<keras.layers.convolutional.Conv2D object at 0x7f0a2507e810> True
<keras.layers.convolutional.Conv2D object at 0x7f0a25082bd0> True
<keras.layers.pooling.MaxPooling2D object at 0x7f0a2506c9d0> True
<keras.layers.convolutional.Conv2D object at 0x7f0a25089710> True
<keras.layers.convolutional.Conv2D object at 0x7f0a25092290> True

```

```

<keras.layers.convolutional.Conv2D object at 0x7f0a250890d0> True
<keras.layers.convolutional.Conv2D object at 0x7f0a2509b110> True
<keras.layers.pooling.MaxPooling2D object at 0x7f0a250a3390> True
<keras.layers.convolutional.Conv2D object at 0x7f0a250a2510> True
<keras.layers.convolutional.Conv2D object at 0x7f0a25084110> True
<keras.layers.convolutional.Conv2D object at 0x7f0a25032fd0> True
<keras.layers.convolutional.Conv2D object at 0x7f0a2503af10> True
<keras.layers.pooling.MaxPooling2D object at 0x7f0a2502cc90> True

```

#Build CNN on the top of VGG19

```

# define sequential model
model = tf.keras.models.Sequential()

# Add the vgg convolutional base model
model.add(vgg_layers)

# add flatten layer
model.add(tf.keras.layers.Flatten())

# add dense layers with some dropout
model.add(tf.keras.layers.Dense(256, activation='relu'))
model.add(tf.keras.layers.Dropout(rate=0.3))
model.add(tf.keras.layers.Dense(256, activation='relu'))
model.add(tf.keras.layers.Dropout(rate=0.3))

# add output layer
model.add(tf.keras.layers.Dense(10, activation='softmax'))

# compile model
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=2e-5),
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

# view model layers
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384
flatten (Flatten)	(None, 512)	0
dense (Dense)	(None, 256)	131328
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 256)	65792

dropout_1 (Dropout)	(None, 256)	0
dense_2 (Dense)	(None, 10)	2570

```
=====
Total params: 20,224,074
Trainable params: 20,224,074
Non-trainable params: 0
=====
```

```
EPOCHS = 100
train_images_3ch_scaled = train_images_3ch / 255.
es_callback = tf.keras.callbacks.EarlyStopping(monitor='val_loss',
patience=2,
restore_best_weights=True,
verbose=1)
```


```
history = model.fit(train_images_3ch_scaled, train_labels,
batch_size=32,
callbacks=[es_callback],
validation_split=0.1, epochs=EPOCHS,
verbose=1)
```

```
Epoch 1/100
869/1688 [=====>.....] - ETA: 56:29 - loss: 0.6488
- accuracy: 0.7827
```

#Accuracy graphs

```
import pandas as pd
fig, ax = plt.subplots(1, 2, figsize=(10, 4))

history_df = pd.DataFrame(history.history)
history_df[['loss', 'val_loss']].plot(kind='line', ax=ax[0])
history_df[['accuracy', 'val_accuracy']].plot(kind='line', ax=ax[1]);
```



#Test data accuracy

```
test_images_3ch_scaled = test_images_3ch / 255.
predictions = model.predict(test_images_3ch_scaled)
predictions[:5]

array([[8.2340867e-10, 2.8423059e-09, 2.0353642e-10, 5.0417581e-09,
        4.4183140e-10, 2.5759527e-07, 9.0209801e-10, 1.1363073e-05,
        1.3139400e-09, 9.9998832e-01],
       [2.4492384e-04, 5.2904752e-06, 9.9845862e-01, 2.3523331e-05,
        5.6046224e-04, 2.1584758e-05, 6.6686596e-04, 9.2018054e-06,
        6.5140521e-06, 2.9754356e-06],
       [1.3240979e-11, 1.0000000e+00, 2.7550812e-11, 1.3596733e-09,
        1.4097771e-11, 3.4155168e-11, 5.0390820e-11, 1.3885454e-11,
        3.7343105e-13, 1.0608974e-10],
       [9.7635571e-09, 9.9999964e-01, 1.6221936e-08, 3.1688836e-07,
        7.5511704e-09, 1.3278602e-08, 2.6905701e-08, 8.8272527e-09,
        5.8504795e-10, 3.3598564e-08],
       [4.9692589e-01, 7.8504533e-04, 1.2435036e-02, 5.7563242e-03,
        3.6659278e-03, 2.6308002e-03, 4.7291115e-01, 1.2741705e-03,
        2.1356090e-03, 1.4800023e-03]], dtype=float32)

prediction_labels = np.argmax(predictions, axis=1)
prediction_labels[:5]

array([9, 2, 1, 1, 0])

from sklearn.metrics import confusion_matrix, classification_report
import pandas as pd

print(classification_report(test_labels, prediction_labels,
target_names=class_names))
pd.DataFrame(confusion_matrix(test_labels, prediction_labels),
index=class_names, columns=class_names)
```

	precision	recall	f1-score	support
T-shirt/top	0.87	0.89	0.88	1000
Trouser	0.99	0.99	0.99	1000
Pullover	0.88	0.92	0.90	1000
Dress	0.92	0.93	0.93	1000
Coat	0.87	0.91	0.89	1000
Sandal	0.99	0.98	0.99	1000
Shirt	0.84	0.74	0.79	1000
Sneaker	0.98	0.97	0.97	1000
Bag	0.99	0.99	0.99	1000
Ankle boot	0.96	0.98	0.97	1000
accuracy			0.93	10000
macro avg	0.93	0.93	0.93	10000
weighted avg	0.93	0.93	0.93	10000

	T-shirt/top	Trouser	Pullover	Dress	Coat	Sandal
Shirt \						
T-shirt/top	888	3	14	25	1	0
Trouser	2	990	0	6	0	0
Pullover	14	1	915	7	33	0
Dress	11	4	11	929	24	0
Coat	3	1	42	14	909	0
Sandal	0	0	0	0	0	983
Shirt	100	2	54	26	72	0
Sneaker	0	0	0	0	0	4
Bag	1	0	1	1	1	1
Ankle boot	0	0	0	0	0	2

	Sneaker	Bag	Ankle boot
T-shirt/top	0	4	0
Trouser	0	1	0
Pullover	0	0	0
Dress	0	1	3
Coat	0	1	0
Sandal	11	0	6
Shirt	0	4	0
Sneaker	966	0	30

Bag	0	994	1
Ankle boot	13	0	984

```
print(test_labels[:100])
```

```
[9 2 1 1 6 1 4 6 5 7 4 5 7 3 4 1 2 4 8 0 2 5 7 9 1 4 6 0 9 3 8 8 3 3 8
0 7
 5 7 9 6 1 3 7 6 7 2 1 2 2 4 4 5 8 2 2 8 4 8 0 7 7 8 5 1 1 2 3 9 8 7 0
2 6
 2 3 1 2 8 4 1 8 5 9 5 0 3 2 0 6 5 3 6 7 1 8 0 1 4 2]
```

```
##F1 score is 0.93
```