

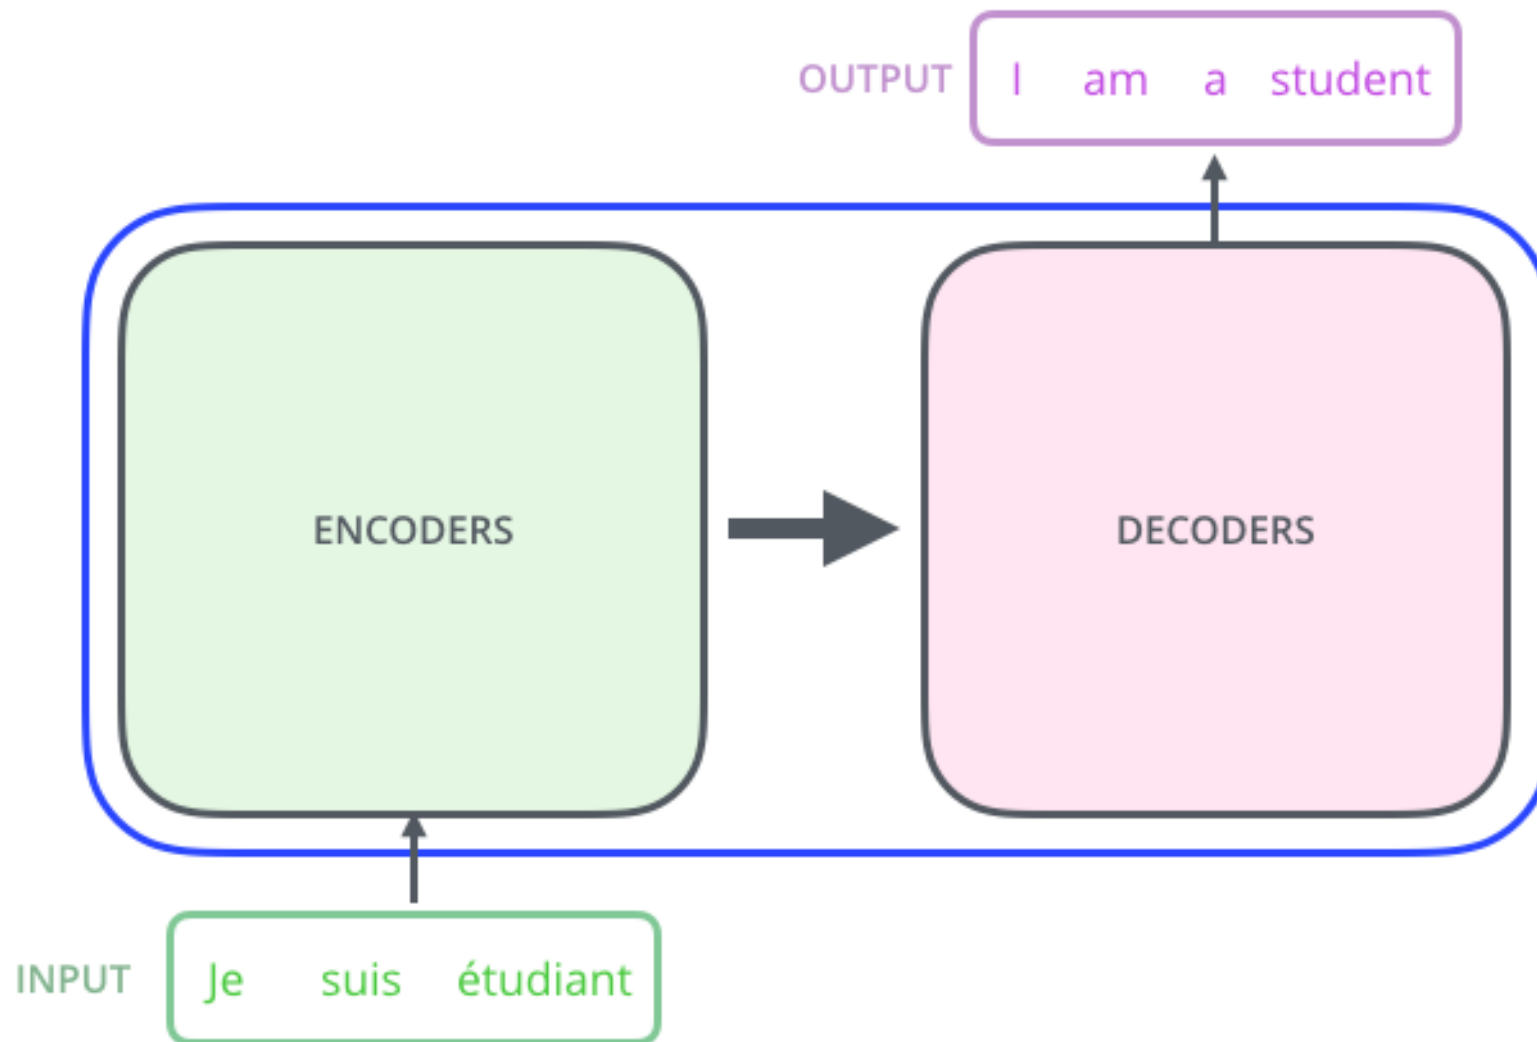
Transformer

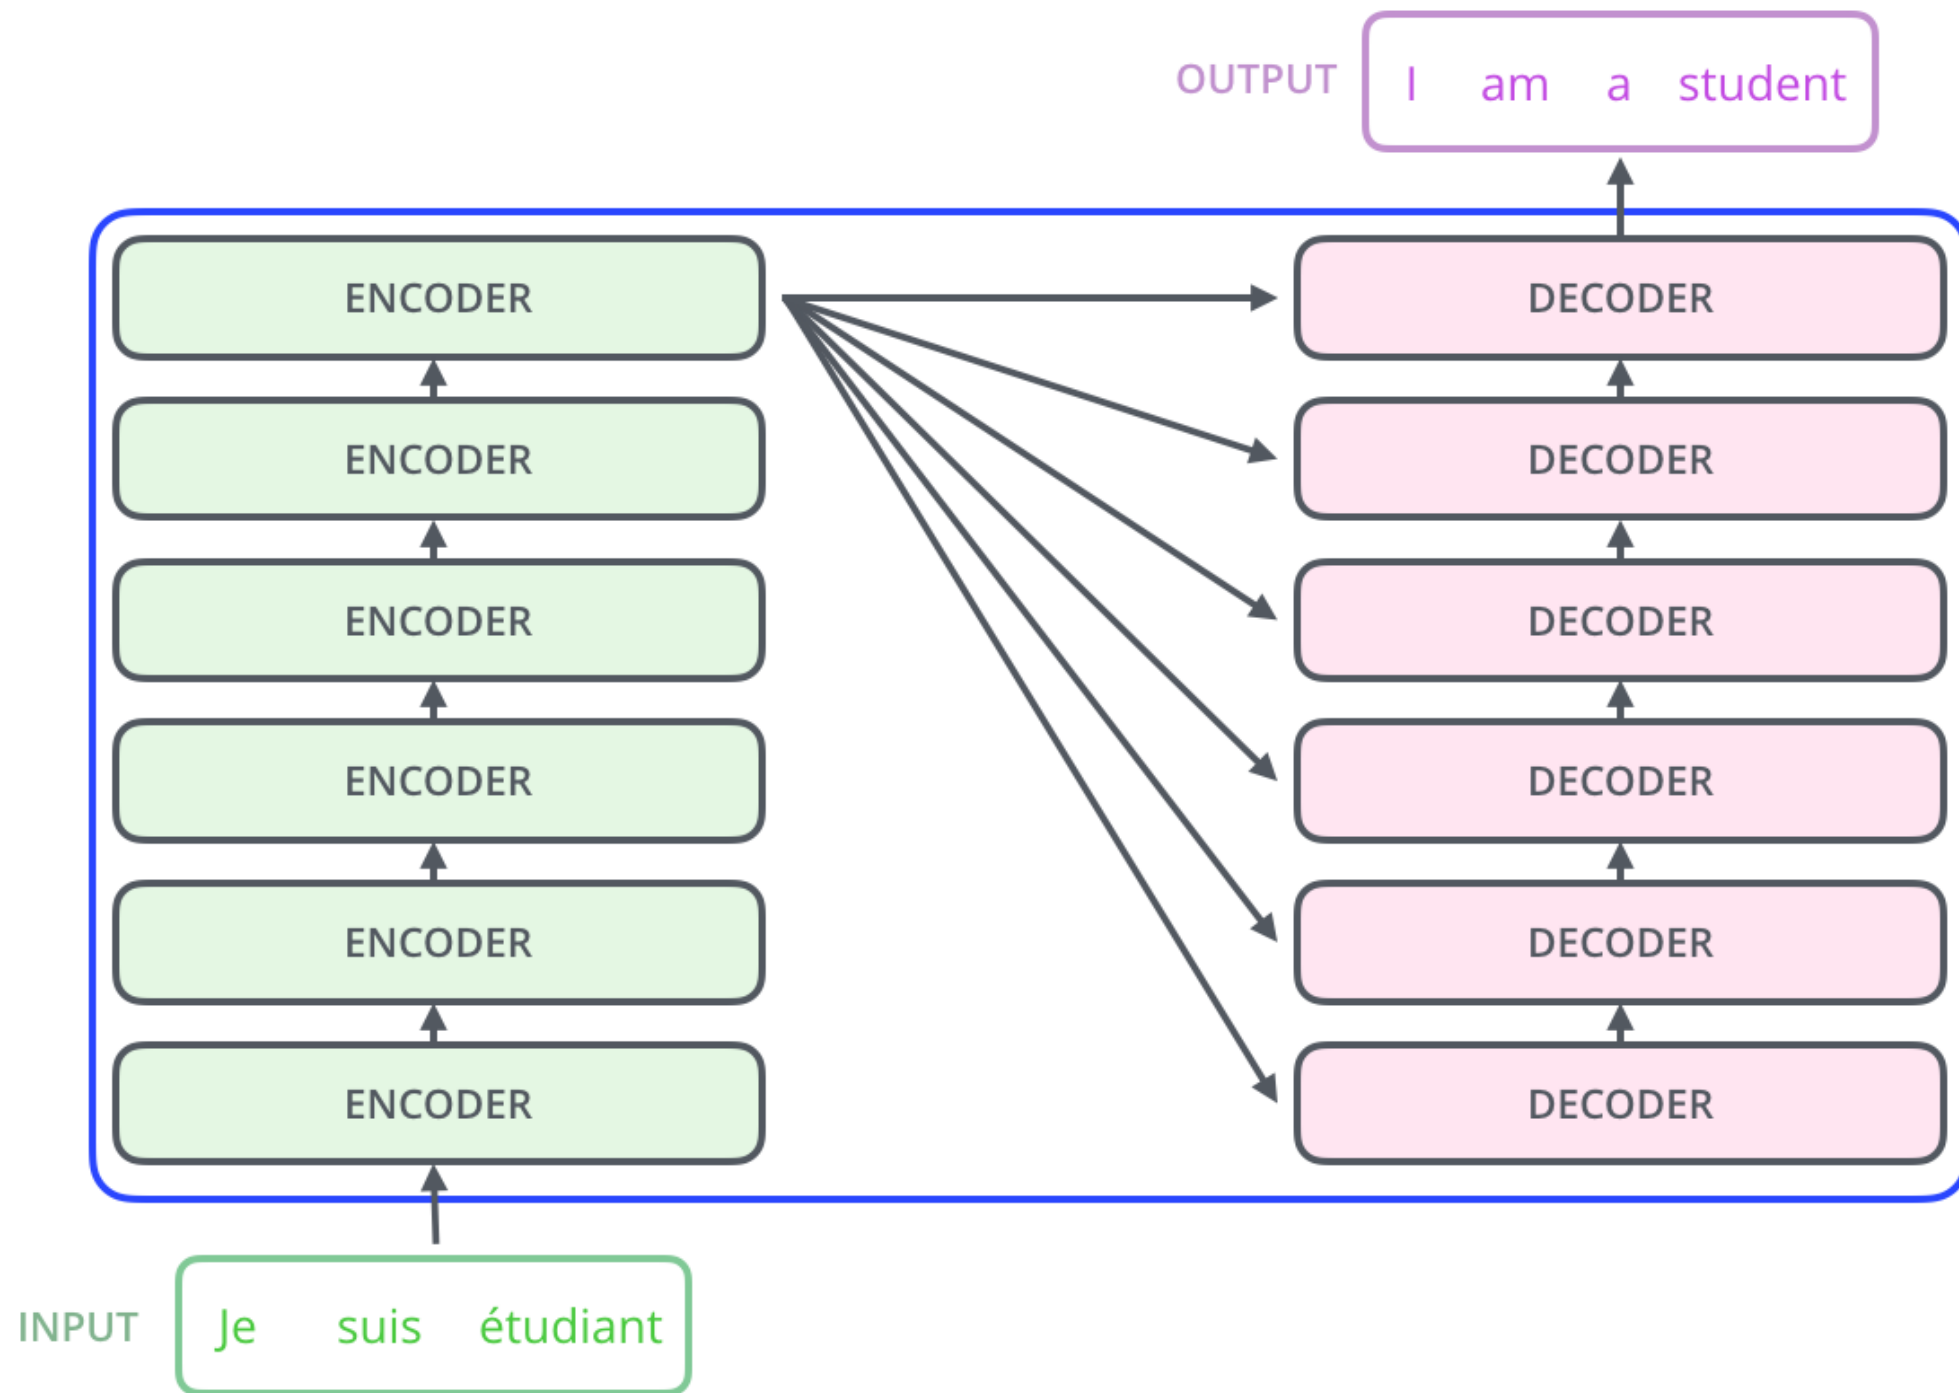
(Attention is All You Need)

Transformer (A High-Level Look)

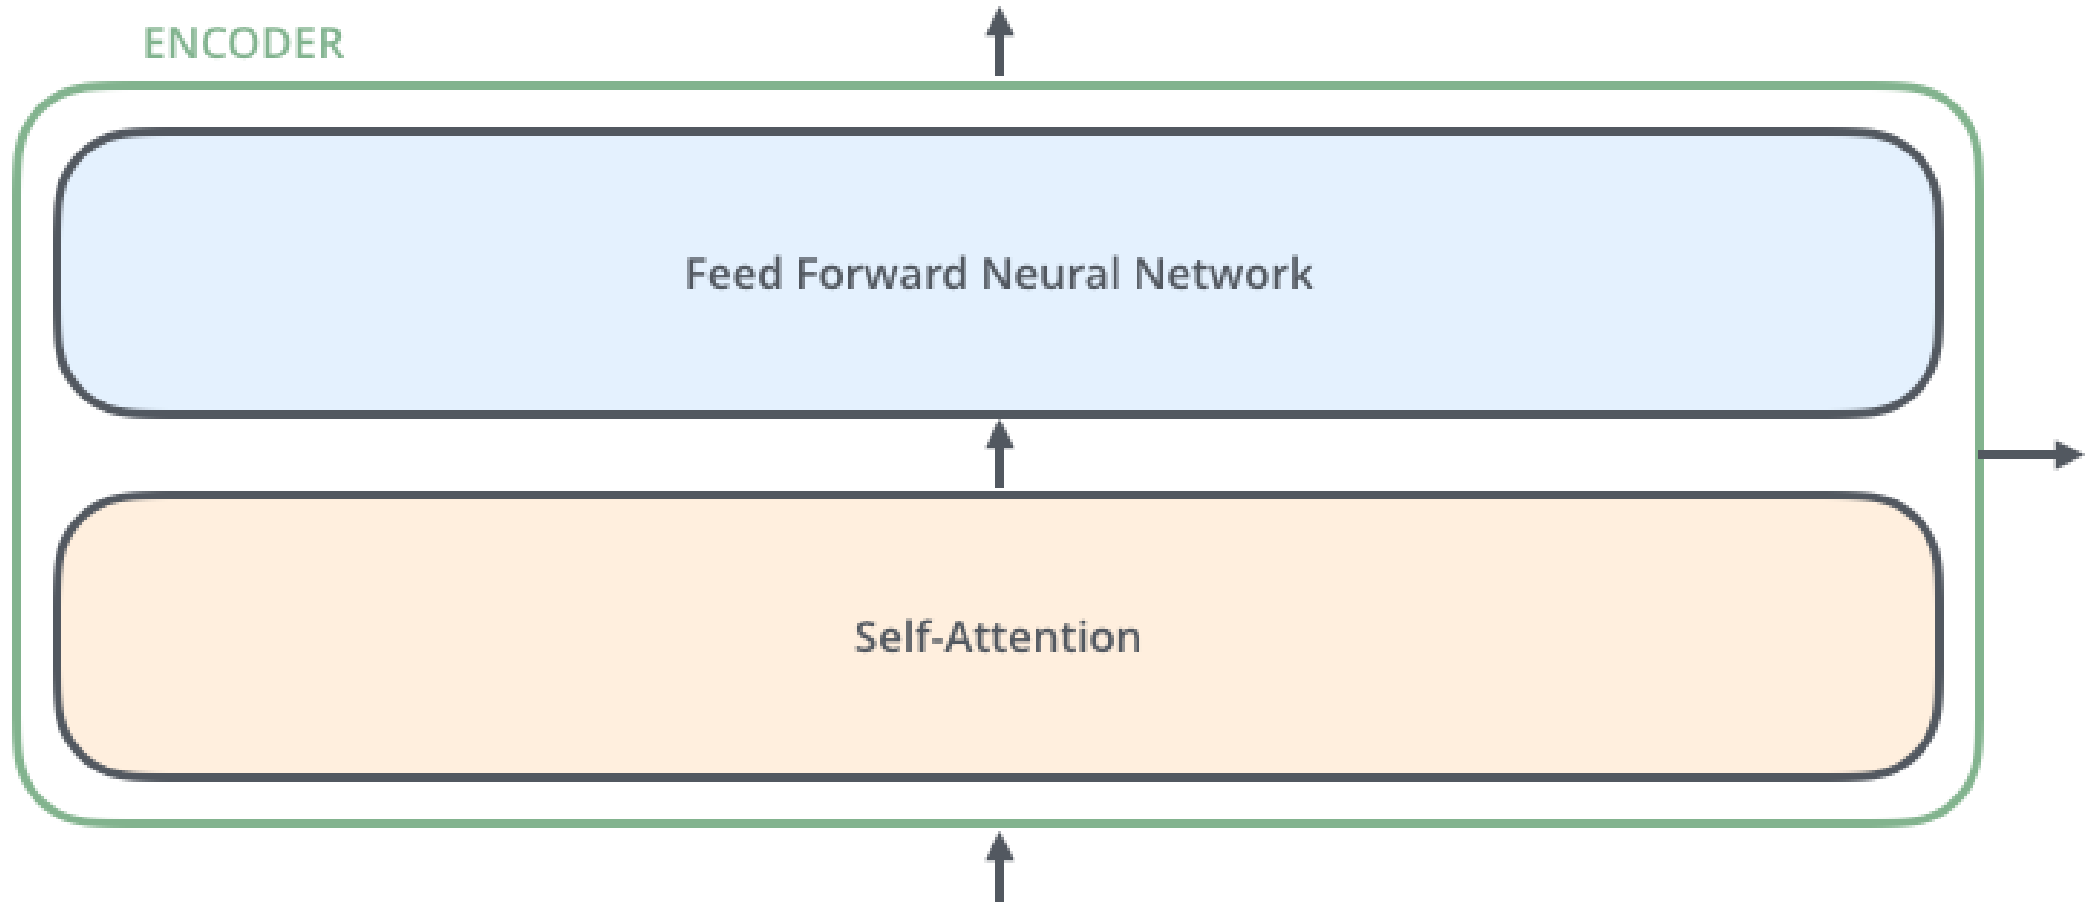


Transformer (A High-Level Look)

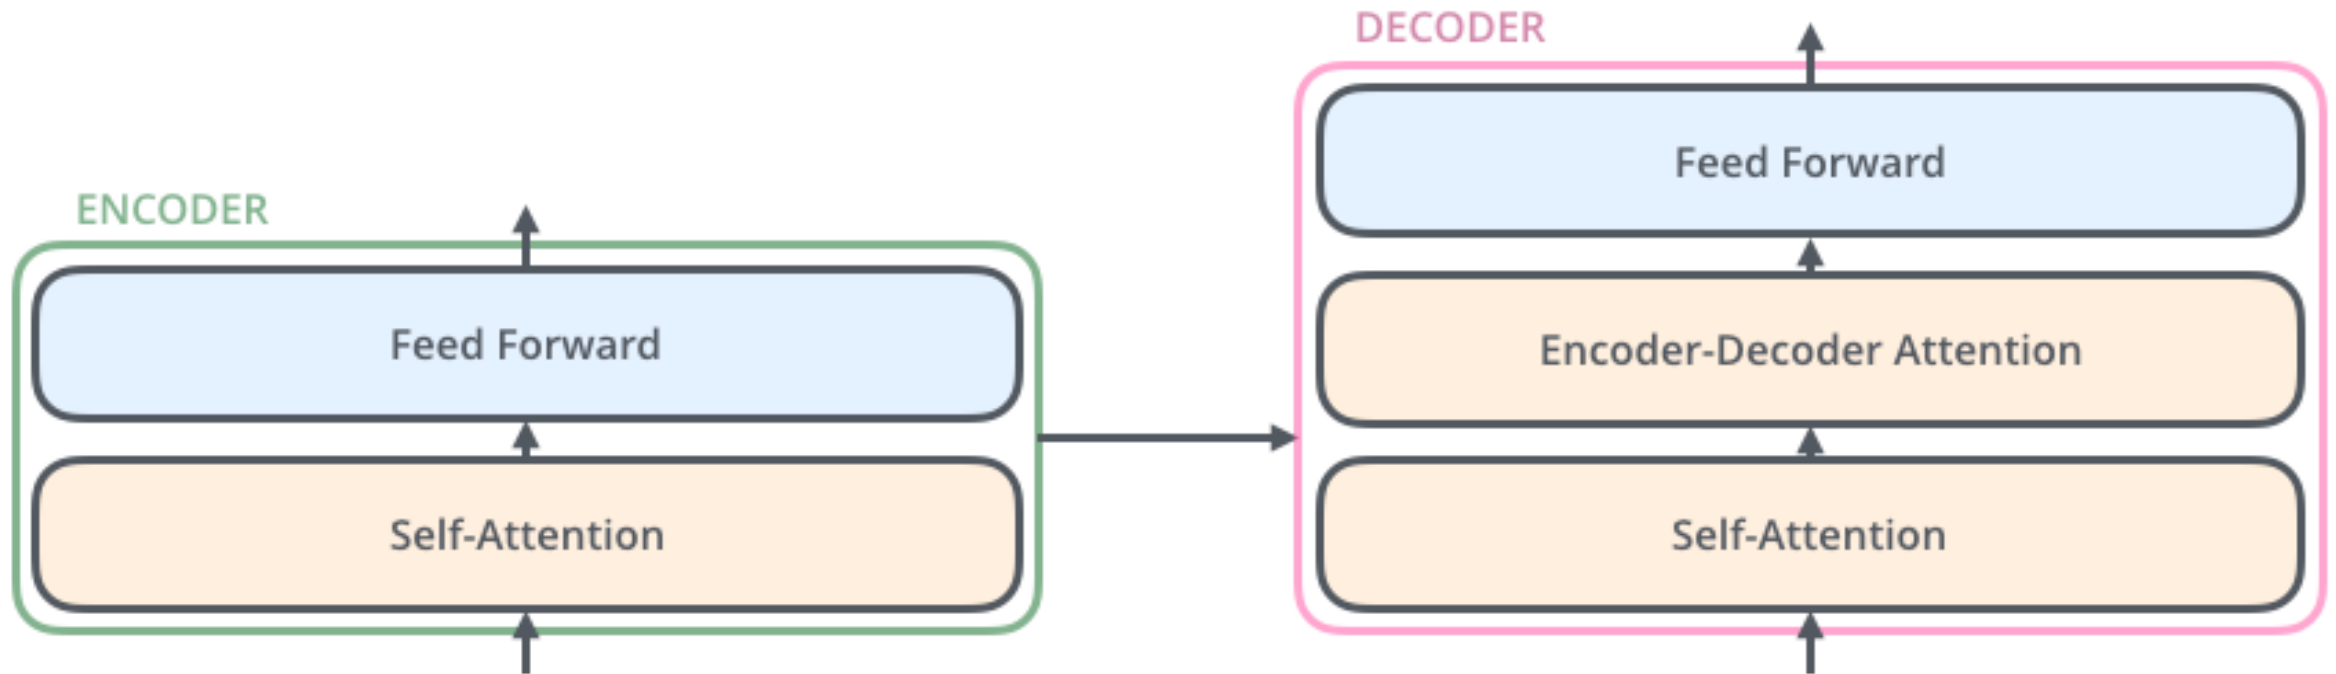




The encoders are all identical in structure. Each one is broken down into two sub-layers:




The encoders are all identical in structure. Each one is broken down into two sub-layers:




Each word is embedded into a vector of size 512. We'll represent those vectors with these simple boxes.



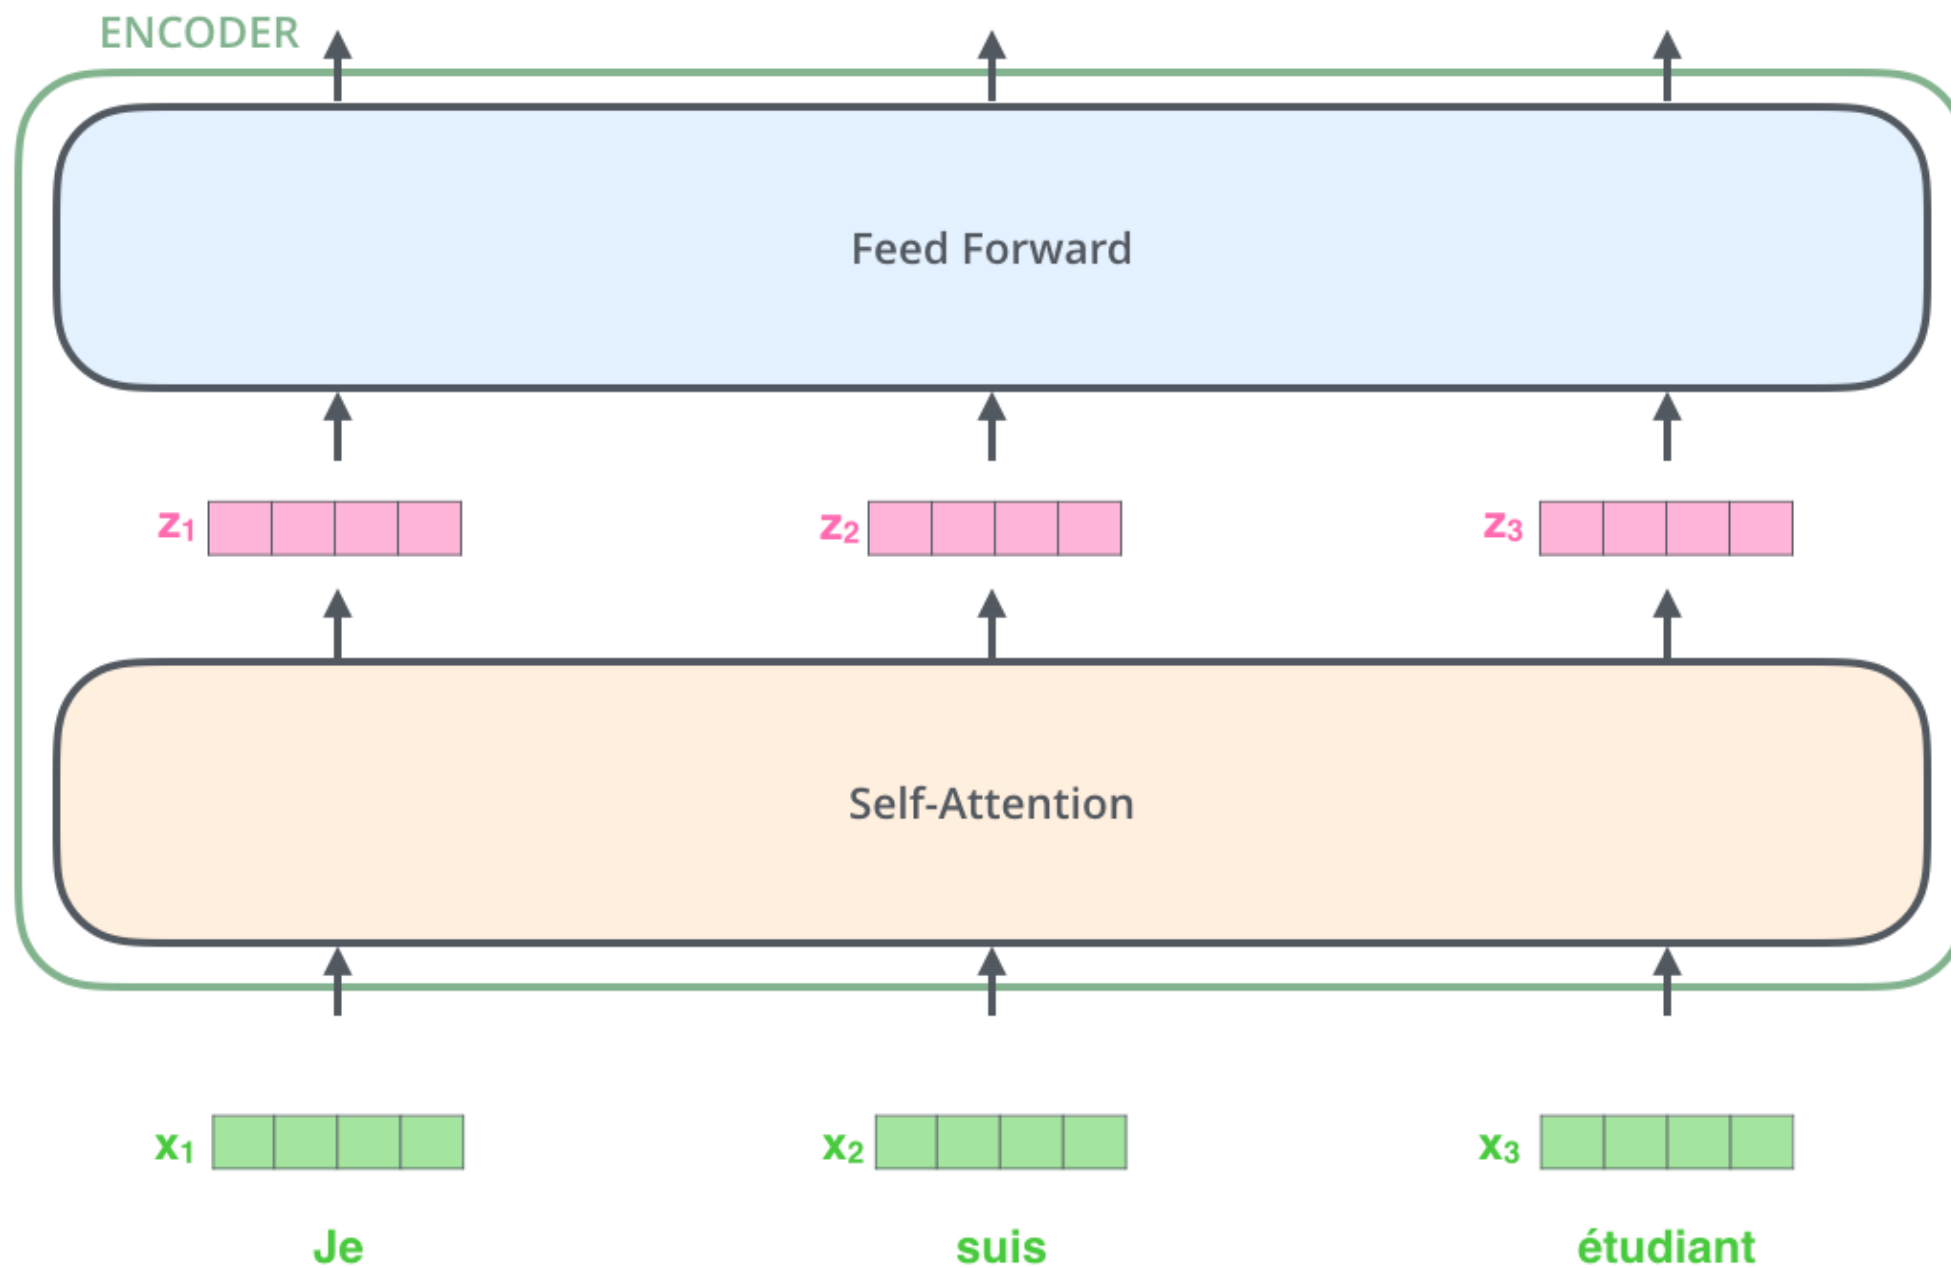
Je



suis

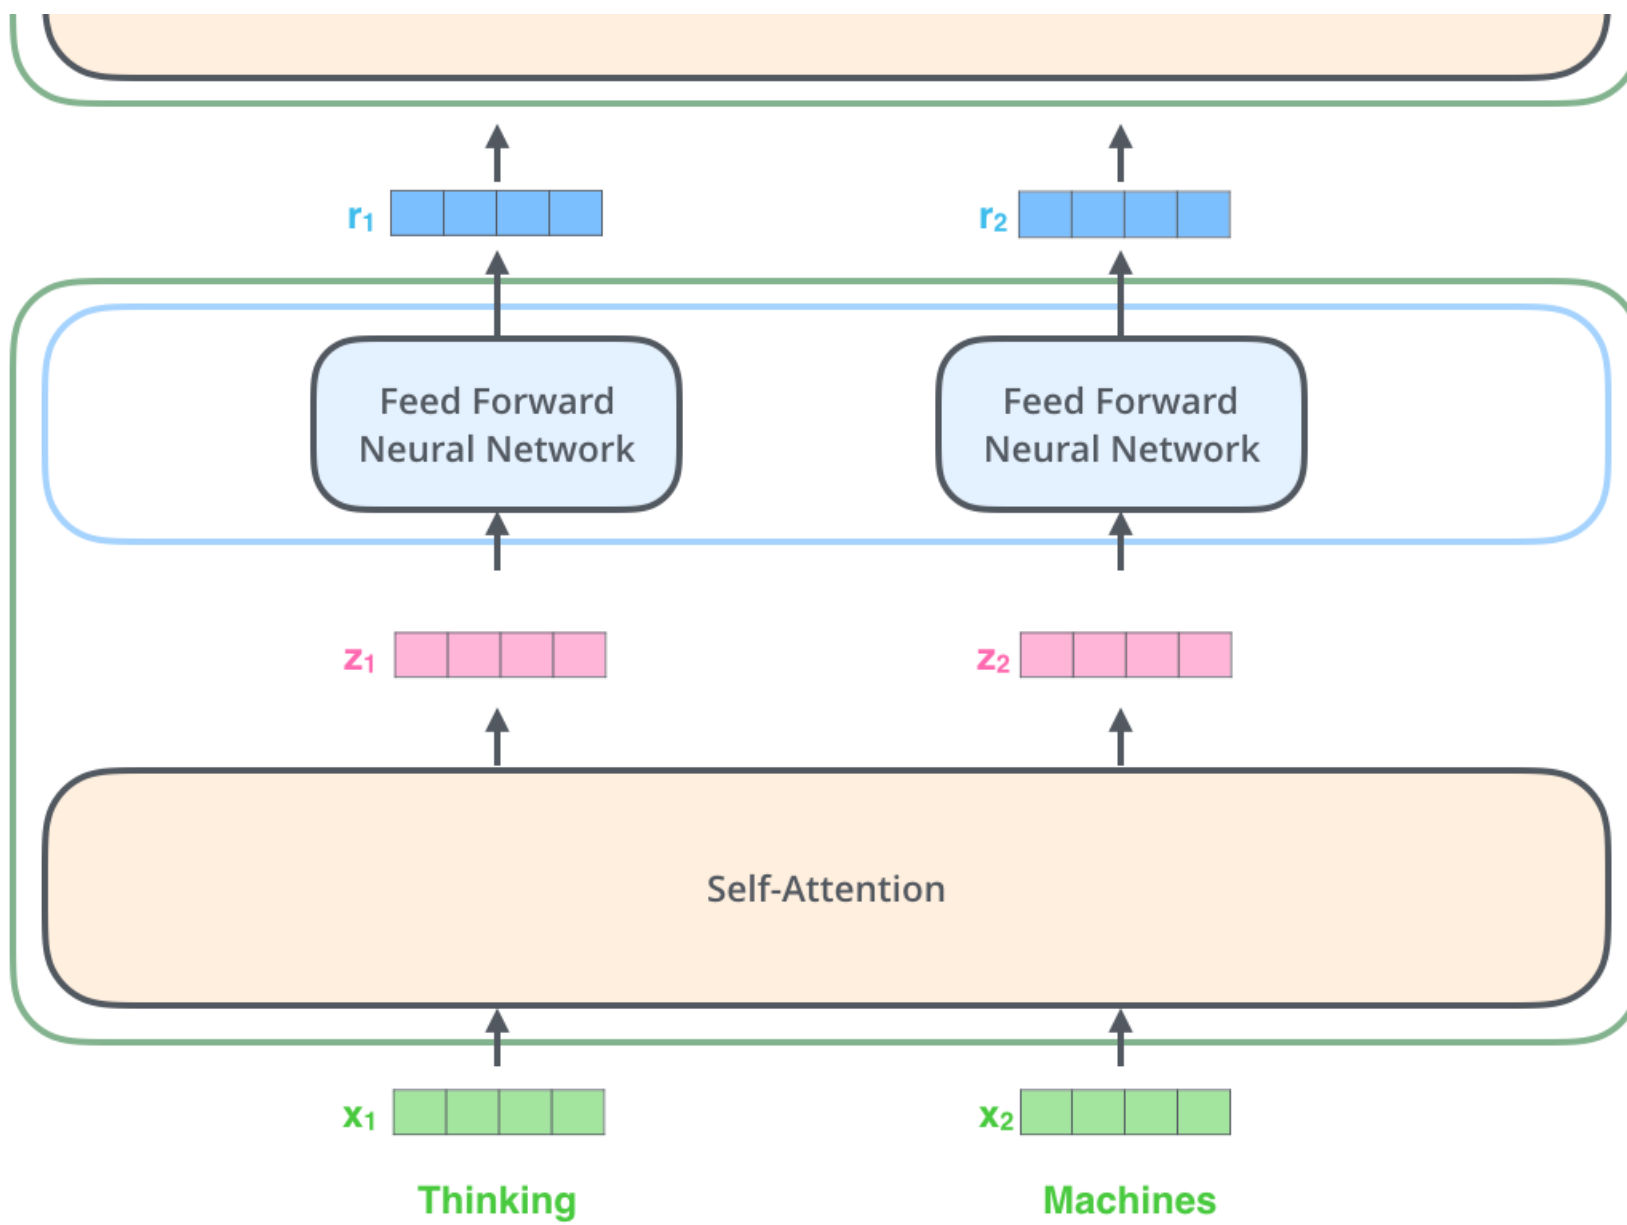


étudiant

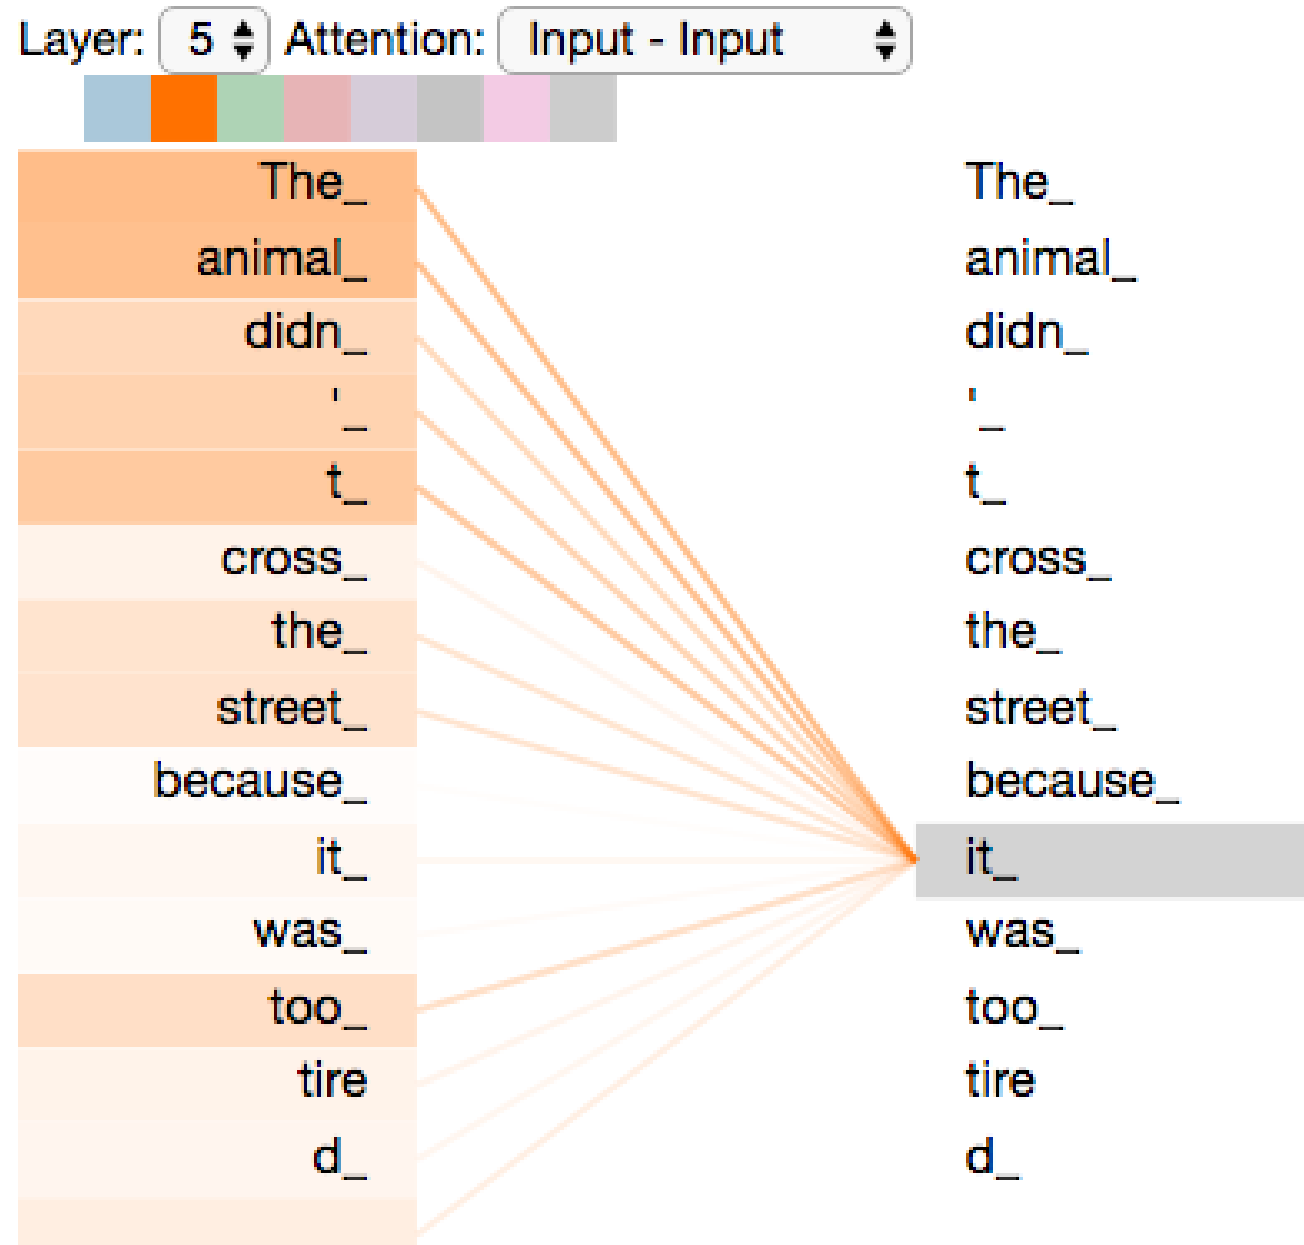


ENCODER #2

ENCODER #1



The animal didn't cross the street because it was too tired”





Input




Thinking

Machines

Embedding





Queries



W^Q

Keys



W^K

Values



W^V


Input

Thinking


Machines

Embedding

x_1




x_2




Queries

q_1




q_2




Keys

k_1



k_2



Values

v_1



v_2



Score

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$

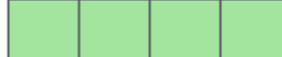
Input

Thinking

Machines

Embedding

x_1 

x_2 

Queries

q_1 

q_2 

Keys

k_1 

k_2 

Values

v_1 

v_2 

Score

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$

Divide by 8 ($\sqrt{d_k}$)

14

12

Softmax

0.88

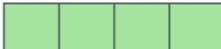
0.12

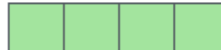
Input

Thinking

Machines

Embedding

x_1 

x_2 

Queries

q_1 

q_2 

Keys

k_1 

k_2 

Values

v_1 

v_2 

Score

$$q_1 \cdot k_1 = 112$$

$$q_1 \cdot k_2 = 96$$

Divide by 8 ($\sqrt{d_k}$)

14

12

Softmax

0.88

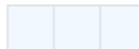
0.12

Softmax

X

Value

v_1 


v_2 

Sum

z_1 


z_2 

X




\times

W^Q




$=$

Q




X




\times

W^K




$=$

K




X




\times

W^V



$=$

V



The self-attention calculation in matrix form

$$\text{softmax} \left(\frac{\begin{matrix} \text{Q} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix} \times \begin{matrix} \text{K}^T \\ \begin{array}{|c|c|} \hline \square & \square \\ \hline \square & \square \\ \hline \square & \square \\ \hline \end{array} \end{matrix}}{\sqrt{d_k}} \right) \begin{matrix} \text{V} \\ \begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array} \end{matrix}$$

=

Z

$\begin{array}{|c|c|c|} \hline \square & \square & \square \\ \hline \square & \square & \square \\ \hline \end{array}$


The Beast With Many Heads

Multi Head Attention

- It expands the model's ability to focus on different positions
- It gives the attention layer multiple “representation subspaces”



X

Thinking
Machines





ATTENTION HEAD #0

Q_0





W_0^Q

K_0



W_0^K


V_0



W_0^V


ATTENTION HEAD #1

Q_1





W_1^Q

K_1



W_1^K


V_1



W_1^V

X

Thinking
Machines




Calculating attention separately in
eight different attention heads




ATTENTION
HEAD #0

Z_0



ATTENTION
HEAD #1


Z_1




...

ATTENTION
HEAD #7


Z_7



1) Concatenate all the attention heads




3) The result would be the Z matrix that captures information from all the attention heads. We can send this forward to the FFNN



2) Multiply with a weight matrix W^O that was trained jointly with the model

X



1) This is our input sentence*

2) We embed each word*


3) Split into 8 heads. We multiply X or R with weight matrices

4) Calculate attention using the resulting $Q/K/V$ matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

Thinking
Machines

X




W_0^Q




Q_0

K_0

V_0



Z_0




W_1^Q




Q_1

K_1

V_1



Z_1



...

...

...


W_7^Q




Q_7

K_7

V_7




Z_7



W^O




Z

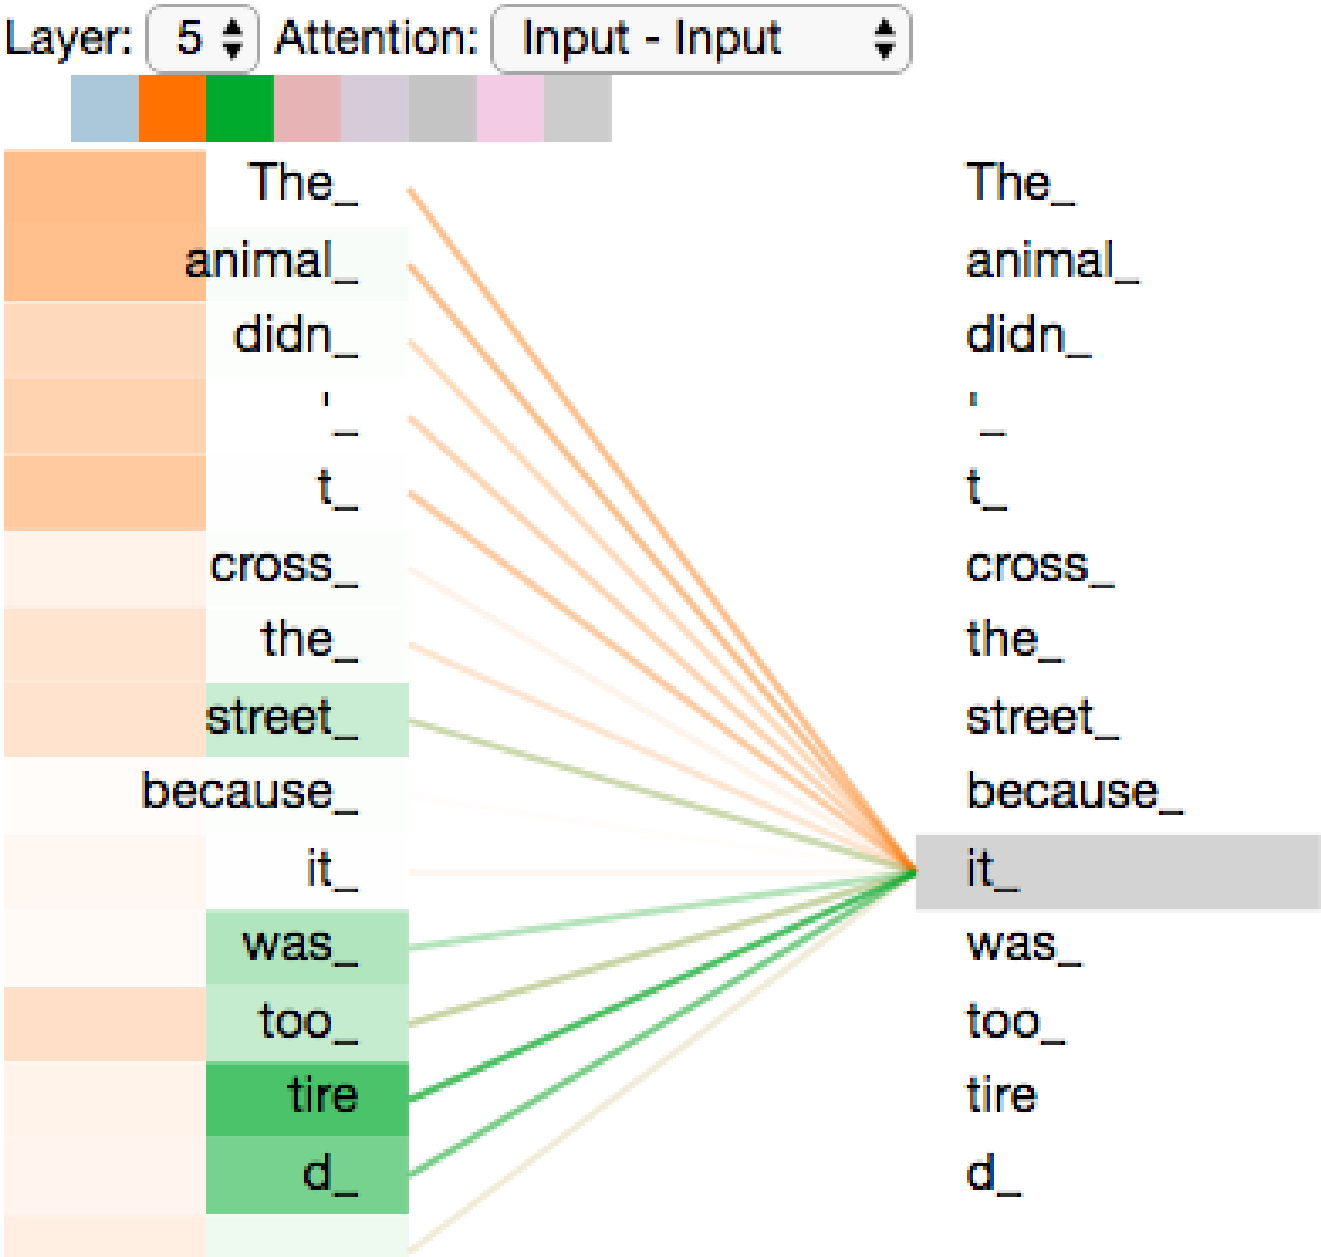


* In all encoders other than #0, we don't need embedding. We start directly with the output of the encoder right below this one

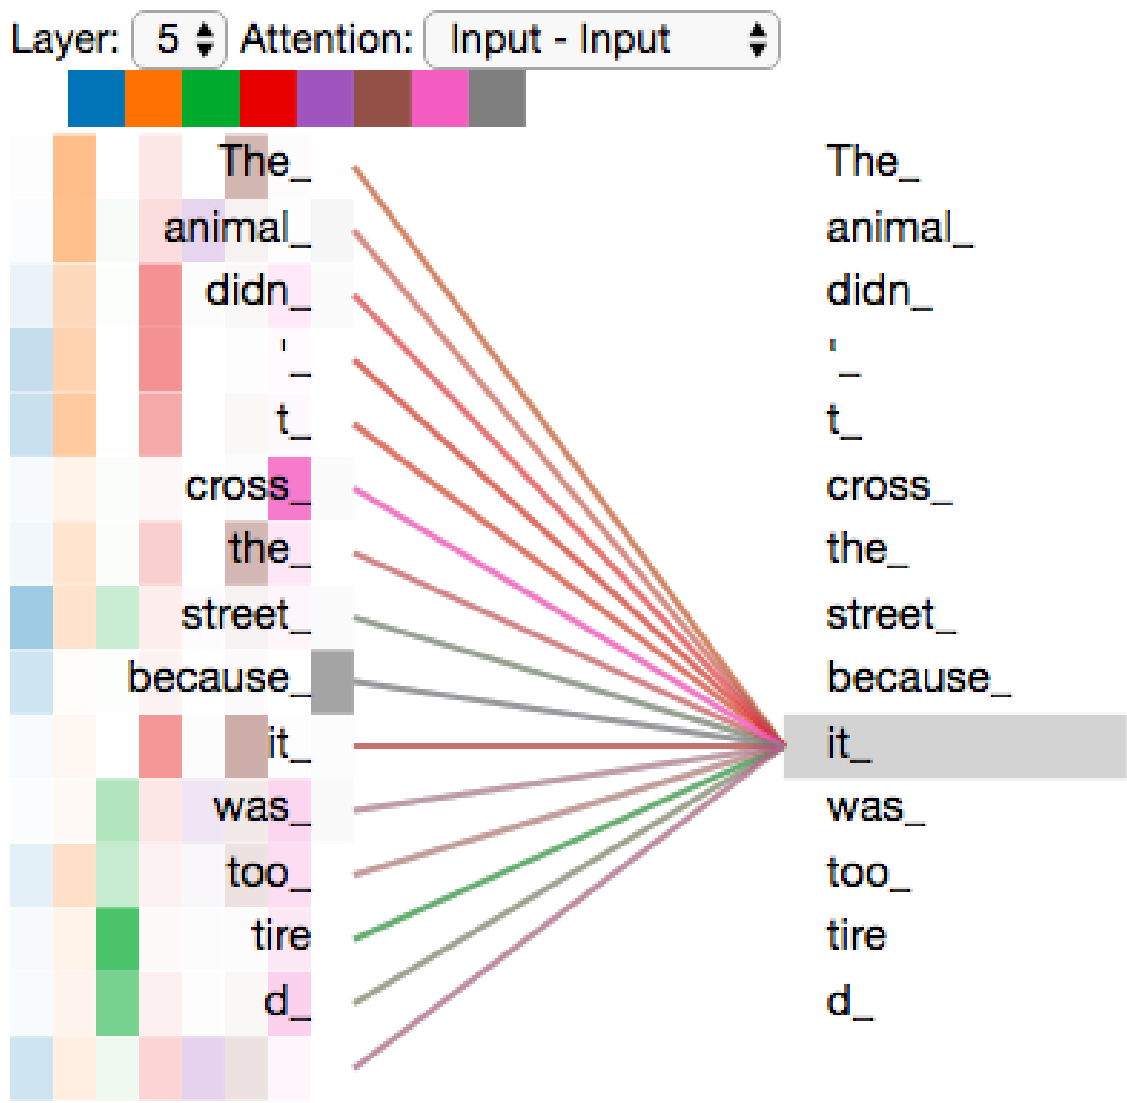
R



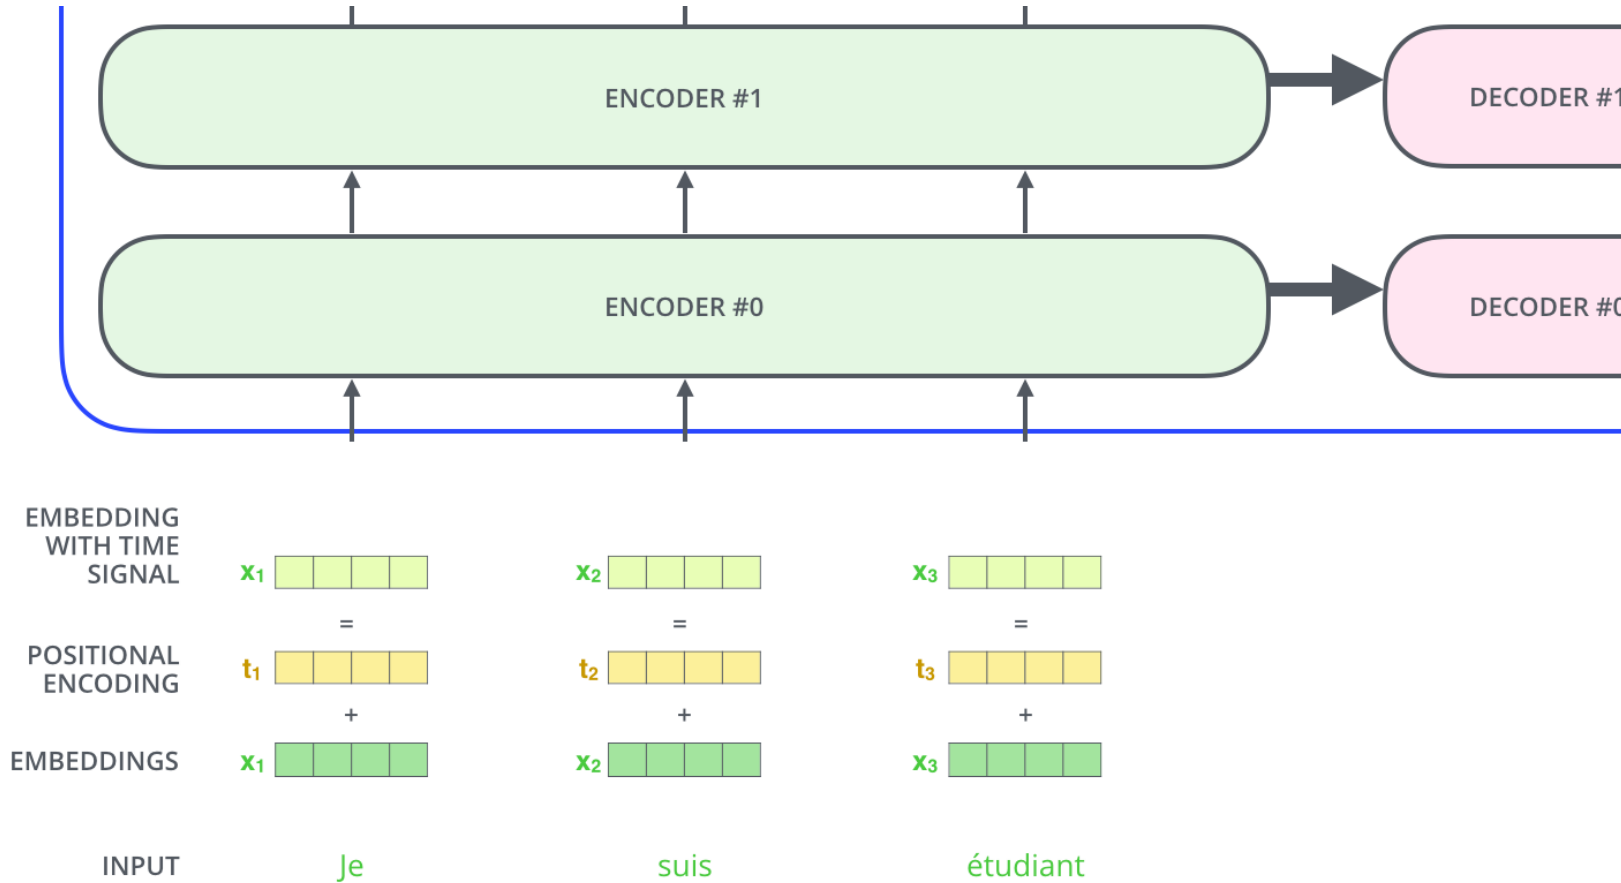
Multi Head Attention



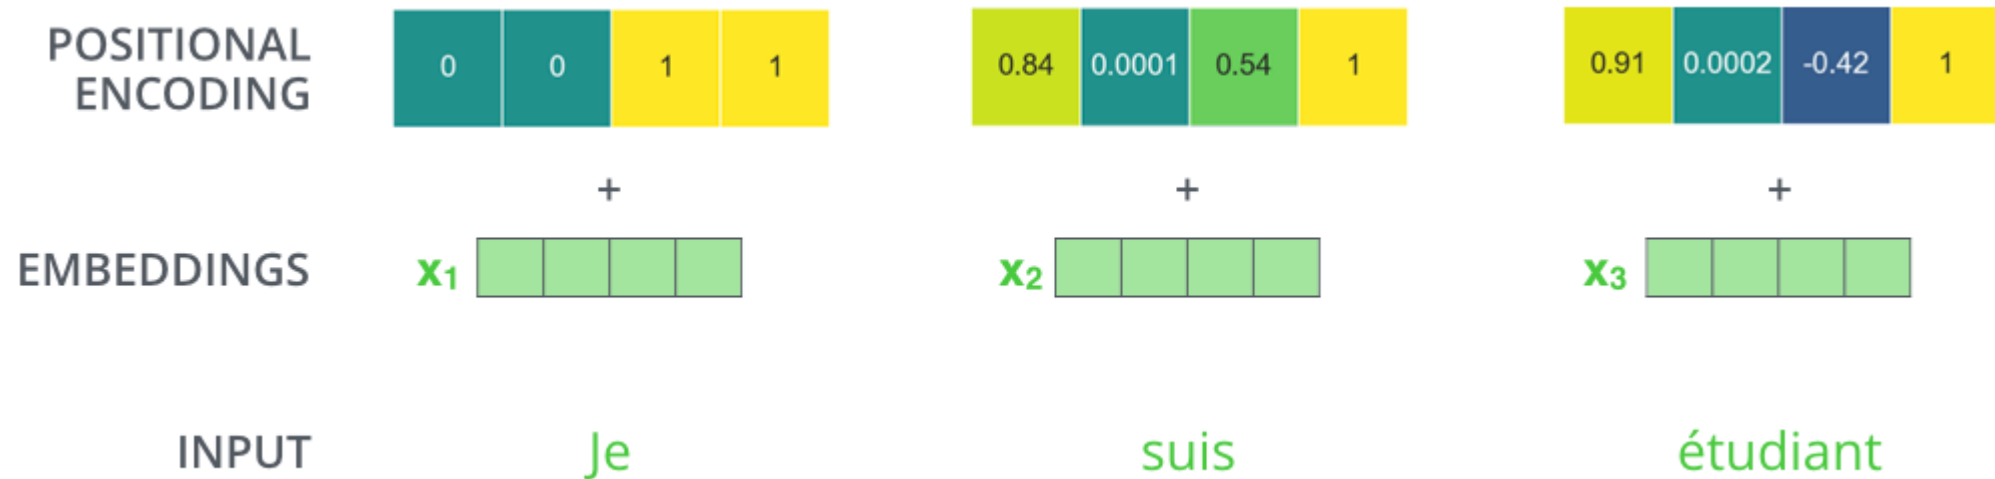
Multi Head Attention



Representing The Order of The Sequence Using Positional Encoding



Representing The Order of The Sequence Using Positional Encoding



Positional Embedding


$$\begin{aligned} e'_w &= e_w + \left[\sin \left(\frac{pos}{10000^0} \right), \cos \left(\frac{pos}{10000^0} \right), \sin \left(\frac{pos}{10000^{2/4}} \right), \cos \left(\frac{pos}{10000^{2/4}} \right) \right] \\ &= e_w + \left[\sin(pos), \cos(pos), \sin \left(\frac{pos}{100} \right), \cos \left(\frac{pos}{100} \right) \right] \end{aligned}$$

$$PE(pos, 2i) = \sin \left(\frac{pos}{10000^{2i/d_{model}}} \right),$$


$$PE(pos, 2i + 1) = \cos \left(\frac{pos}{10000^{2i/d_{model}}} \right).$$

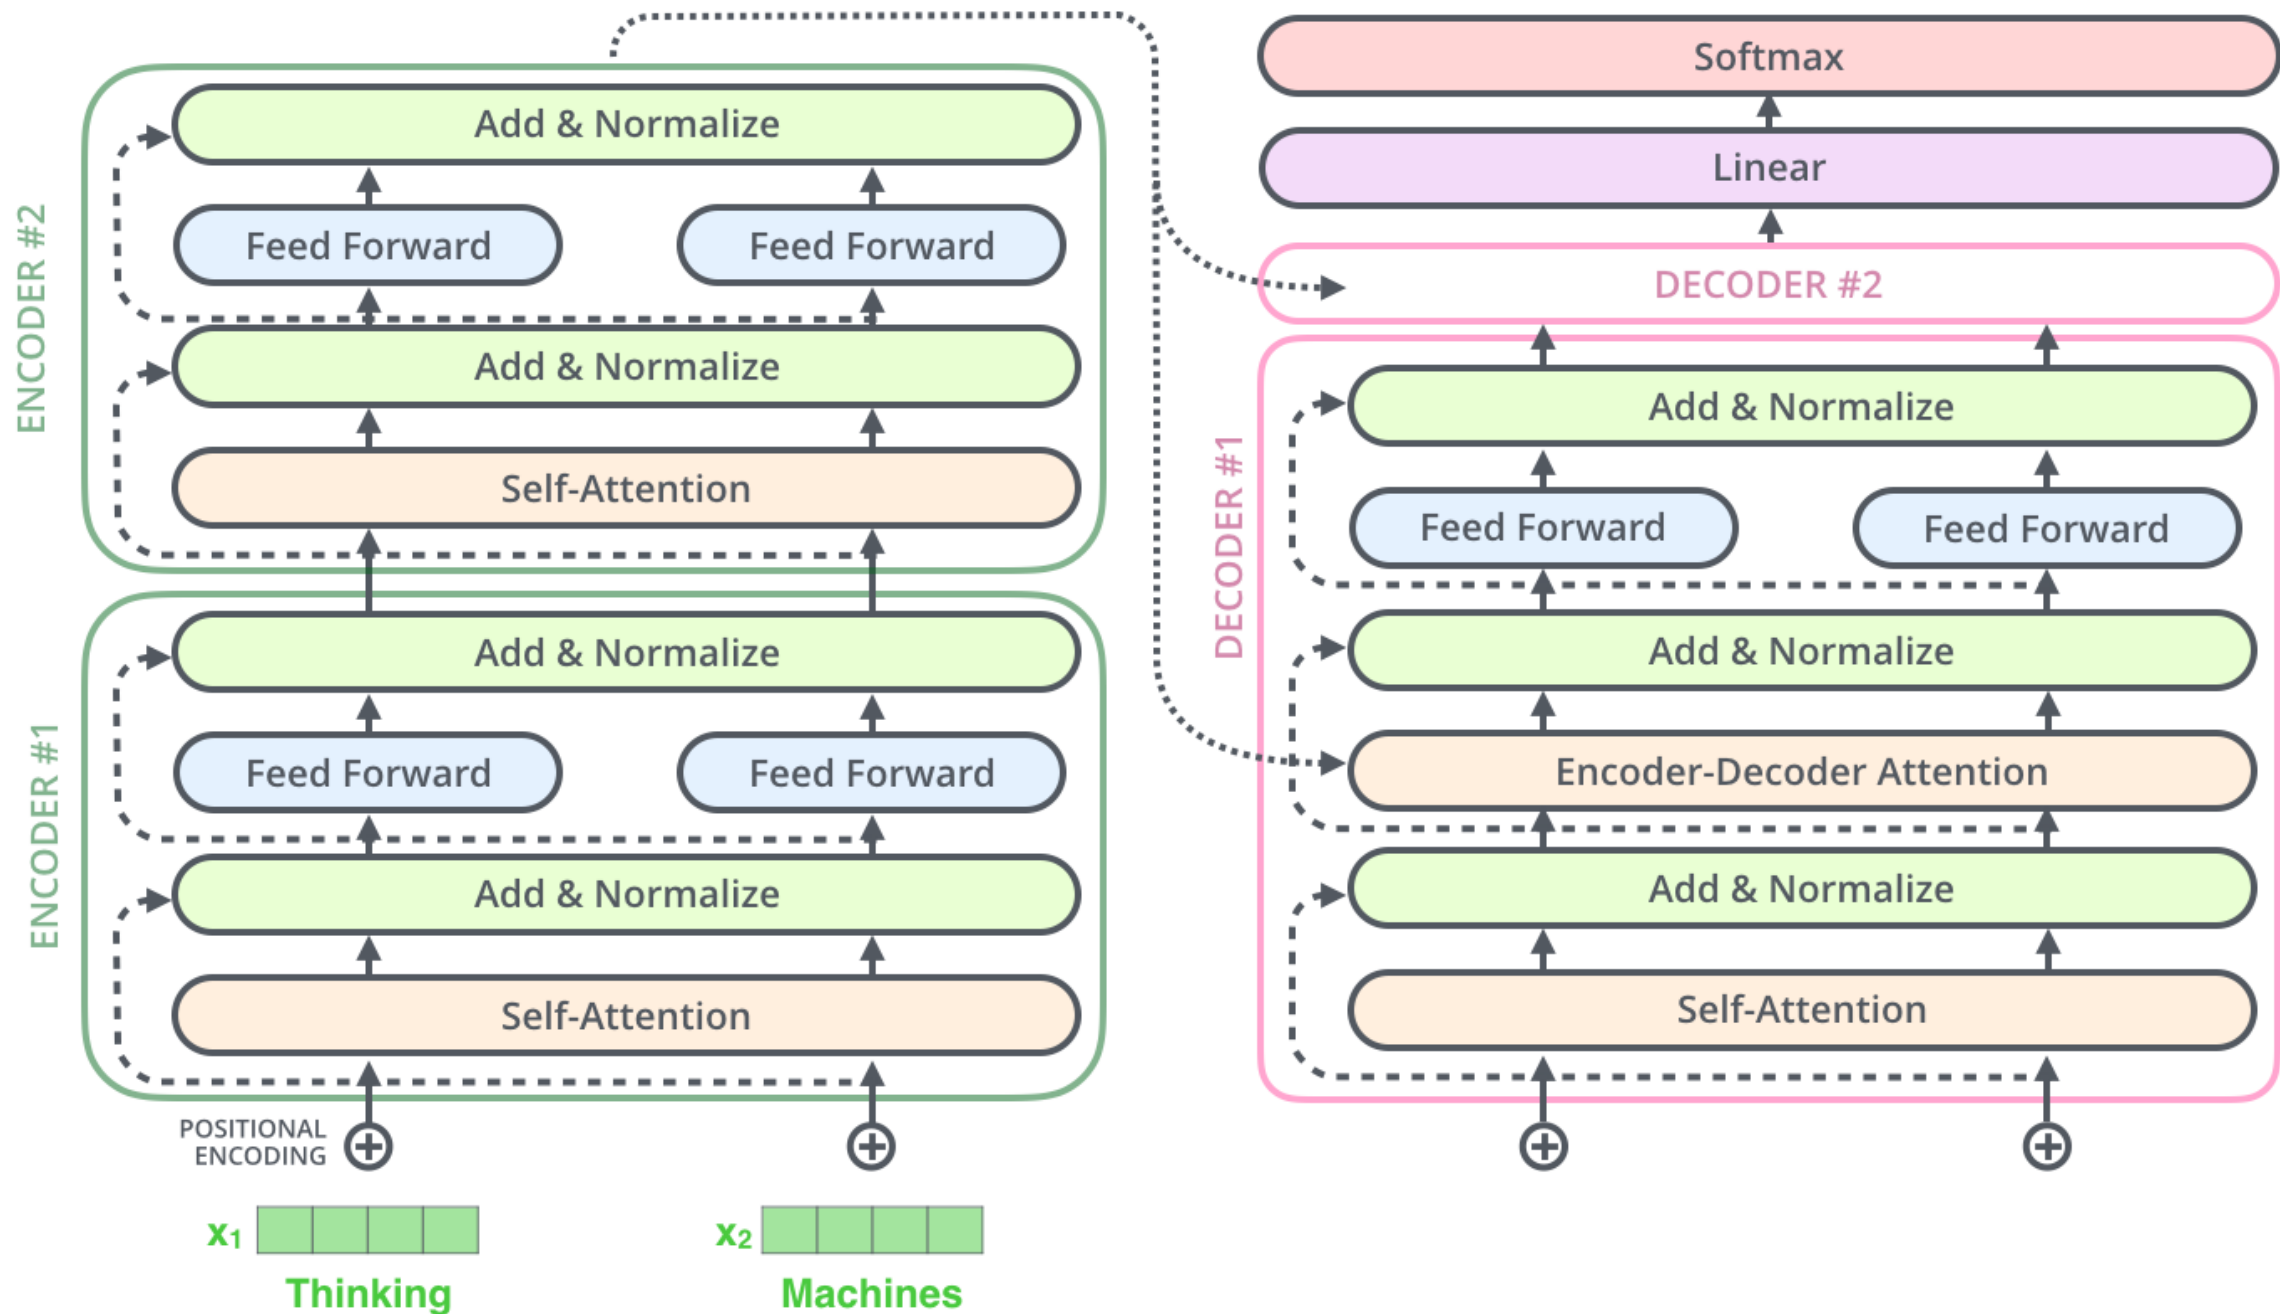
pos = position of a word in window e.g 1,2,3...
i = 512 vector size
d = dimension..idk

Positional Embedding



Positional Embedding

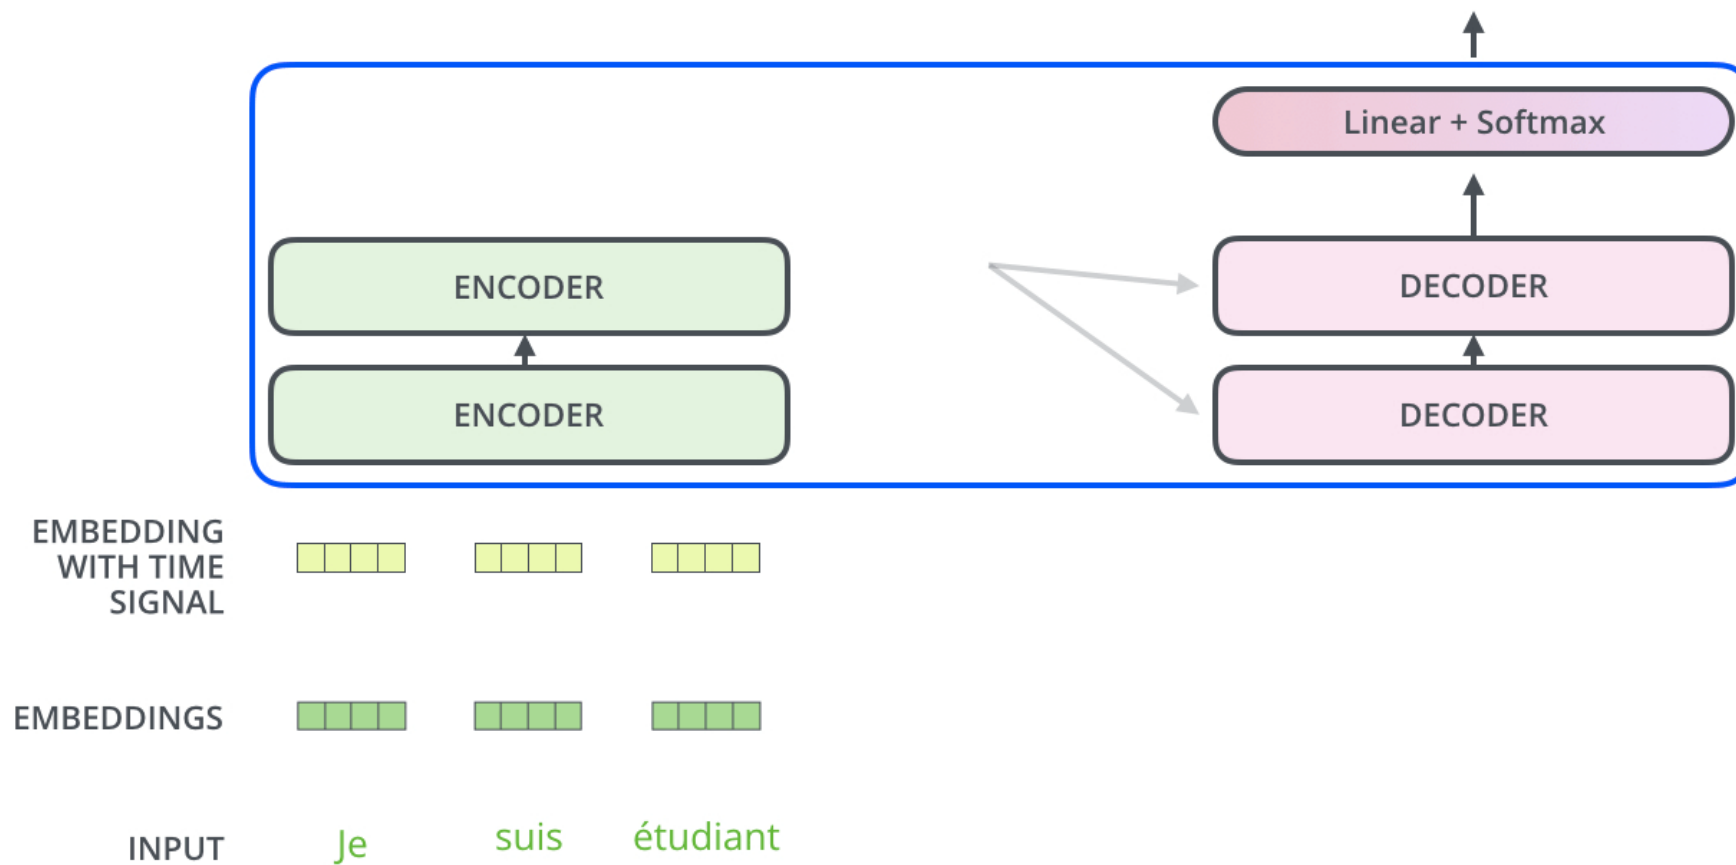




The Decoder Side

Decoding time step: 1 2 3 4 5 6

OUTPUT



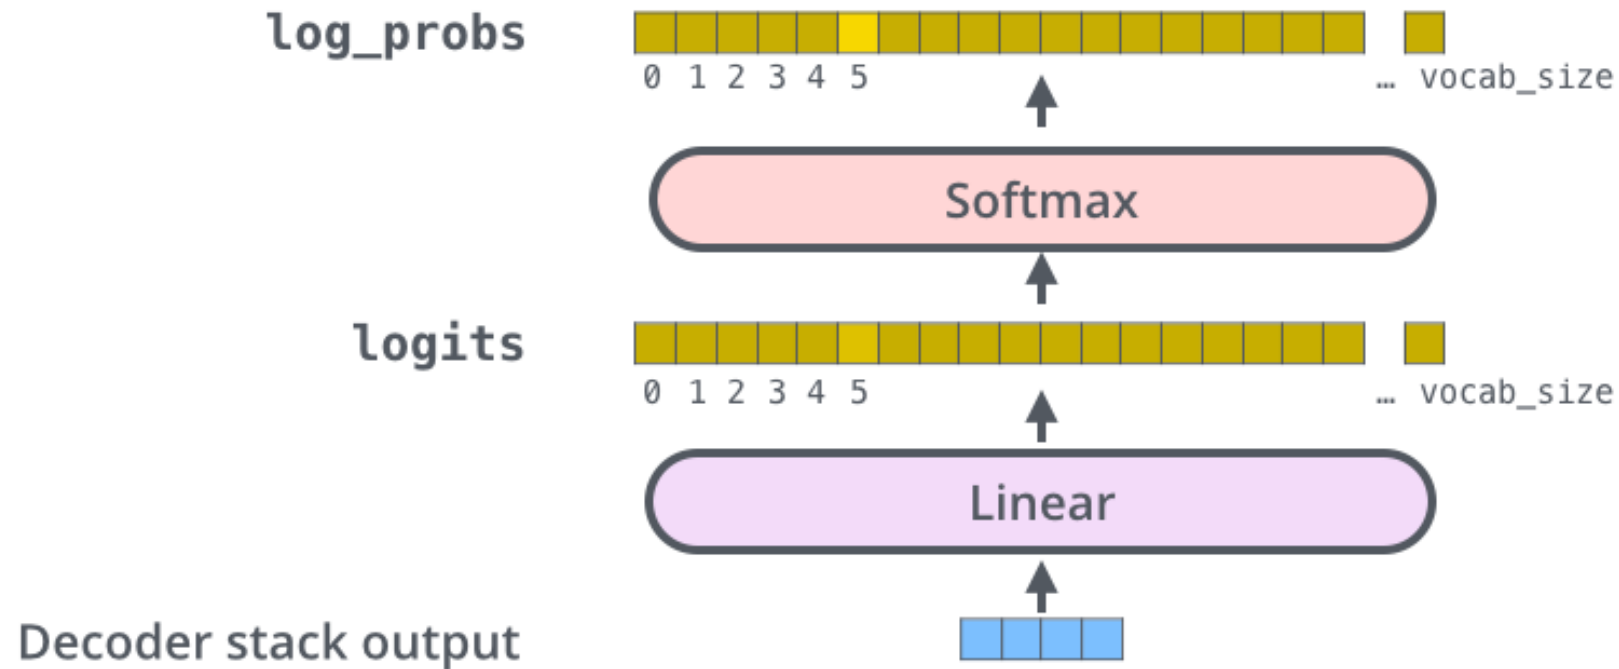
Attention in Decoder

- The self attention layers in the decoder operate in a slightly different way than the one in the encoder
- In the decoder, the self-attention layer is only allowed to attend to earlier positions in the output sequence. This is done by masking future positions (setting them to **-inf**) before the softmax step in the self-attention calculation
- The “Encoder-Decoder Attention” layer works just like multiheaded self-attention, except it creates its Queries matrix from the layer below it, and takes the Keys and Values matrix from the output of the encoder stack

The Final Linear and Softmax Layer

Which word in our vocabulary
is associated with this index?

Get the index of the cell
with the highest value
(**argmax**)



Recap Of Training

Output Vocabulary

WORD	a	am	I	thanks	student	<eos>
INDEX	0	1	2	3	4	5

One-hot encoding of the word "am"

0.0	1.0	0.0	0.0	0.0	0.0
-----	-----	-----	-----	-----	-----

Untrained Model Output



Correct and desired output



a

am

I

thanks

student

<eos>




The Loss Function

- Each probability distribution is represented by a vector of width `vocab_size` (6 in our toy example, but more realistically a number like 3,000 or 10,000)
- The first probability distribution has the highest probability at the cell associated with the word “i”
- The second probability distribution has the highest probability at the cell associated with the word “am”
- And so on, until the fifth output distribution indicates


Target Model Outputs

Output Vocabulary: a am I thanks student <eos>



Trained Model Outputs

Output Vocabulary: a am I thanks student <eos>



Beam Search



References

- [Jay Alammar](https://jalammar.github.io/illustrated-transformer/) <https://jalammar.github.io/illustrated-transformer/>
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. **Attention is all you need**. In *Proceedings of the 31st International Conference on Neural Information Processing Systems (NIPS'17)*,