

Parallel and Distributed Computing

CS3006

Lecture 18

Hybrid Programming (MPI + OpenMP)


15th May 2024

Dr. Rana Asif Rehman

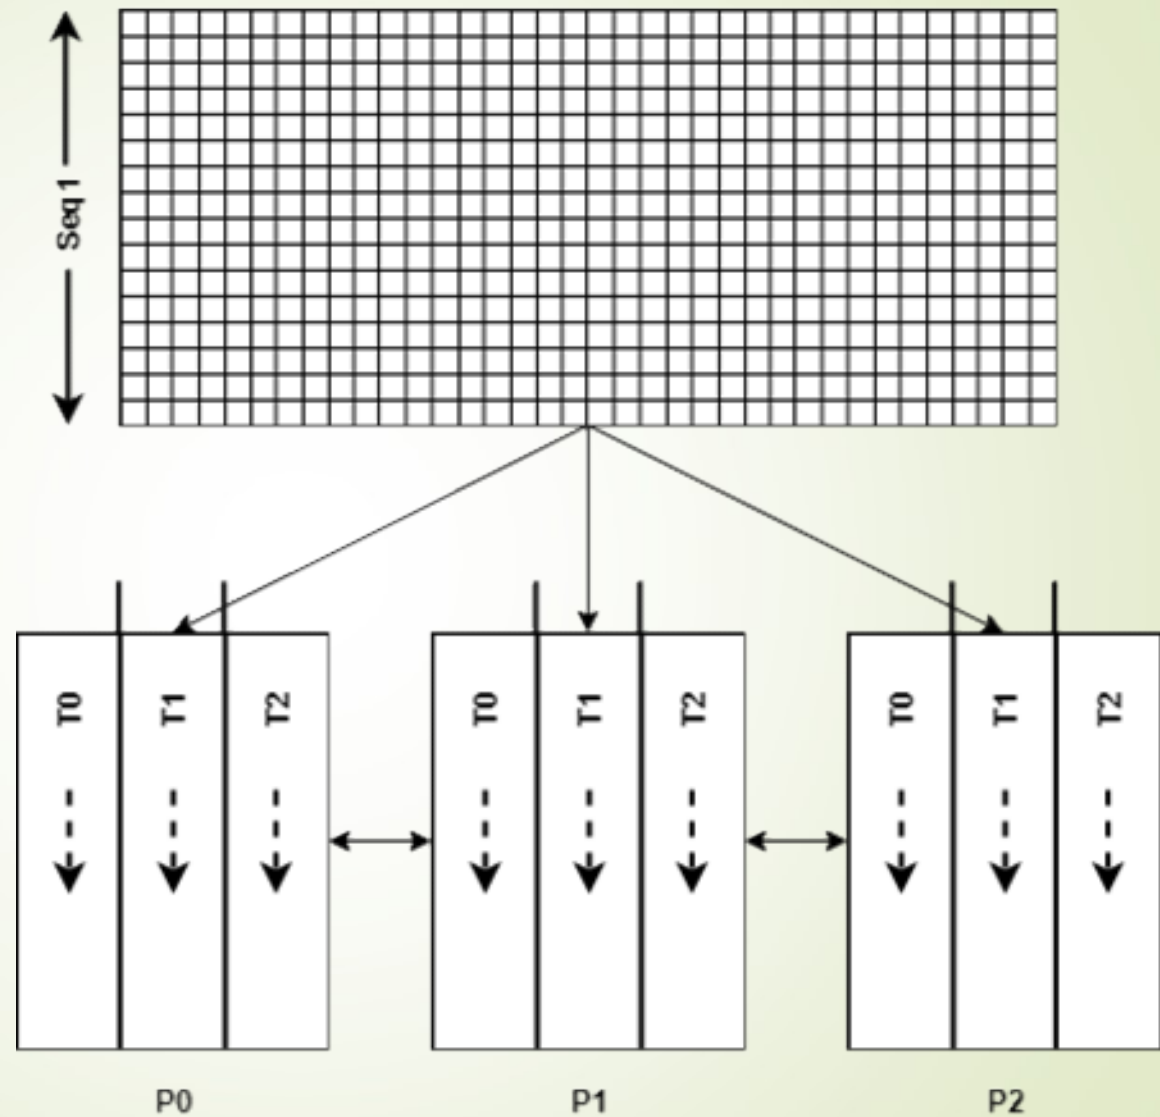
Hybrid Programming with MPI and OpenMP

- Consider your cluster of workstations consists of different multicore processors.
- Launching a single MPI process for each of the workstations, may not optimally use the available resources.
- We can combine **MPI and OpenMP** to increase utilization of the resources by:
 - Using single MPI process per workstation
 - And generating different number of threads in each workstation.
 - Each workstation can have different number of threads, according to their infrastructures.

Hybrid Programming with MPI and OpenMP



**Single MPI
process per
workstation and
multiple threads**



Hybrid Programming with MPI and OpenMP [Use cases]

5

- The simplest and safe way to combine MPI with OpenMP is to **never use the MPI calls inside the OpenMP parallel regions.**
- MPI level communication must only be done by the master thread.

```
main(int argc, char **argv) {  
    ...  
    MPI_Init(&argc, &argv);  
    ... // master thread only --> MPI calls here  
    #pragma omp parallel  
    {  
        .. // team of threads --> no MPI calls here  
    }  
    ... // master thread only --> MPI calls here  
    MPI_Finalize();  
}
```

Hybrid Programming with MPI and OpenMP [Helloworld.c program]

```
MPI_Init(&argc, &argv);
    int P, name_len;
    MPI_Comm_size(MPI_COMM_WORLD, &P);
    int process_id;
    MPI_Comm_rank(MPI_COMM_WORLD, &process_id);
    char processor_name[MPI_MAX_PROCESSOR_NAME];
    MPI_Get_processor_name(processor_name, &name_len);
omp_set_dynamic(0);
    omp_set_num_threads(omp_get_num_procs());
    int thread_id;
#pragma omp parallel private(thread_id)
{
    thread_id = omp_get_thread_num();
    printf("From Process%d thread_id %d out of %d mapped on\n", process_id, thread_id,
        omp_get_num_threads(), processor_name);
}
MPI_Finalize();
```

Hybrid Programming with MPI and OpenMP [Helloworld.c program]

Important: At header files include both `mpi.h` and `omp.h`

Compiling the program:

```
mpicc -fopenmp ./helloworld.c -o ./a.out
```

Executing the program:

```
mpiexec -np 2 ./a.out
```

```
From Process#1 thread_id 2 out of 4 mapped on processor:Haier-PC Hello Hybrid
husnain8721@Haier-PC:/mnt/c/hybrid$ mpiexec -np 2 ./a.out
From Process#0 thread_id 0 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#0 thread_id 1 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#1 thread_id 3 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#1 thread_id 1 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#1 thread_id 2 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#0 thread_id 3 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#1 thread_id 0 out of 4 mapped on processor:Haier-PC Hello Hybrid
From Process#0 thread_id 2 out of 4 mapped on processor:Haier-PC Hello Hybrid
```

Hybrid Programming with MPI and OpenMP

8

- Till now we have only seen solutions that don't require inter-workstation (OR MPI level) communication in the parallel region
- What if we try to use MPI calls within openmp parallel regions?
 - Default calls for MPI-1 primitive routines are not thread safe
 - May result in strange outputs
- So, solution?
 - Use `MPI_Init_thread()` with **`MPI_THREAD_MULTIPLE`** as support level instead of simple `MPI_init()` call for starting the program

Hybrid Programming with MPI and OpenMP

```
int MPI_Init_thread(int *argc, char ***argv, int required, int *provided)
```

- MPI_Init_thread() initializes the MPI execution environment (similar to MPI_Init())
- defines the support level for multithreading:
 - **required** is the aimed support level
 - **provided** is the support level provided by the MPI implementation

Available Support levels:

- **MPI_THREAD_SINGLE** – only one thread will execute (the same as initializing the environment with MPI_Init())
- **MPI_THREAD_FUNNELED** – only the master thread can make MPI calls
- **MPI_THREAD_SERIALIZED** – all threads can make MPI calls, but only one thread at a time can be in such state
- **MPI_THREAD_MULTIPLE** – all threads can make simultaneous MPI calls without any constraints

Hybrid Programming with MPI and OpenMP

MPI_THREAD_FUNNELED

- With support level MPI_THREAD_FUNNELED only the master thread can make MPI calls.
- One way to ensure this is to protect the MPI calls with the **omp master** directive.

```
#pragma omp parallel
{
    ...
    #pragma omp barrier // explicit barrier at entrance
    #pragma omp master // only the master thread makes the MPI call
    mpi_call();
    #pragma omp barrier // explicit barrier at exit
    ...
}
```

Hybrid Programming with MPI and OpenMP

MPI_THREAD_SERIALIZED

- Only one thread at a time will make calls to the MPI library, but all threads are eligible to make such calls as long as they do not do so at the same time

```
#pragma omp parallel private(tid) num_threads(2)
{
    tid = omp_get_thread_num();
    if(tid== 0){ mpi_call(); }
    #pragma omp barrier // barrier so that other threads don't enter next mpi
    call
    if(tid== 1){ mpi_call(); }
    #pragma omp barrier // explicit barrier

    ...
}
```

Hybrid Programming with MPI and OpenMP

MPI_THREAD_MULTIPLE

- All the threads can make simultaneous MPI calls without any constraints
- No need of any barrier synchronizations among the threads
- MPI library is then responsible for thread safety within that library, and for any libraries that it in turn uses

```
#pragma omp parallel private(tid) num_threads(2)
{
    ....
    tid = omp_get_thread_num();
    mpi_call(); //all the threads will call this mpi routine parallely
    ...
}
```

Hybrid Programming with MPI and OpenMP

MPI_THREAD_MULTIPLE

- How to identify the thread number that is involved in the communication?
- You can pass **threadID** as message **tag**. The receiving process then can identify each message from the source uniquely.

```
#pragma omp parallel private(tid) num_threads(2)
{
...
    tid= omp_get_thread_num();
    if (my_rank == 0) // MPI process 0 sends 2 messages
        MPI_Send(a, 1, MPI_INT, 1, tid, MPI_COMM_WORLD);
    else if (my_rank == 1) // MPI process 1 receives 2 messages
        MPI_Recv(b, 1, MPI_INT, 0, tid, MPI_COMM_WORLD, &status);
...
}
```

Hybrid Programming with MPI and OpenMP

`hybrid_multiple_support_sendRecv.c`

- Program uses `MPI_THREAD_MULTIPLE` as level of support
- The program basically implements a naïve hybrid one-to-all broadcast of a buffer containing 4 integers.
- Source process uses 4 threads, where each thread is responsible of sending one integer to all the other processes
- While receiving processes also use 4 threads, each which receives a single integer from the source process and copies it into correct position in the receive buffer.

Hybrid Programming with MPI and OpenMP

hybrid_multiple_support_sendRecv.c OUTPUT

- For 2 MPI processes and sendbuff=[2,4,6,8]
process_id:0-->thread#0 is sending 2 to process 1
process_id:0-->thread#1 is sending 4 to process 1
process_id:0-->thread#2 is sending 6 to process 1
process_id:0-->thread#3 is sending 8 to process 1
process-id:1 Received 2 from thread#0 of source
process-id:1 Received 4 from thread#1 of source
process-id:1 Received 6 from thread#2 of source
process-id:1 Received 8 from thread#3 of source
recvbuff=2 4 6 8

Hybrid Programming with MPI and OpenMP

hybrid_multiple_support_sendRecv.c OUTPUT

➤ For 4 MPI processes [**Sorce Process output**]

process_id:0-->thread#0 is sending 2 to process 1
process_id:0-->thread#0 is sending 2 to process 2
process_id:0-->thread#0 is sending 2 to process 3

process_id:0-->thread#1 is sending 4 to process 1
process_id:0-->thread#1 is sending 4 to process 2
process_id:0-->thread#1 is sending 4 to process 3

process_id:0-->thread#2 is sending 6 to process 1
process_id:0-->thread#2 is sending 6 to process 2
process_id:0-->thread#2 is sending 6 to process 3

process_id:0-->thread#3 is sending 8 to process 1
process_id:0-->thread#3 is sending 8 to process 2
process_id:0-->thread#3 is sending 8 to process 3

Hybrid Programming with MPI and OpenMP

hybrid_multiple_support_sendRecv.c OUTPUT

- For 4 MPI processes **[receiving processes output]**

process-id:1 Received 2 from thread#0 of source
 process-id:1 Received 4 from thread#1 of source
 process-id:1 Received 6 from thread#2 of source
 process-id:1 Received 8 from thread#3 of source

process-id:2 Received 2 from thread#0 of source
 process-id:2 Received 4 from thread#1 of source
 process-id:2 Received 6 from thread#2 of source
 process-id:2 Received 8 from thread#3 of source

process-id:3 Received 2 from thread#0 of source
 process-id:3 Received 6 from thread#2 of source
 process-id:3 Received 4 from thread#1 of source
 process-id:3 Received 8 from thread#3 of source

Hybrid Programming with MPI and OpenMP

hybrid_multiple_support _sendRecv.c OUTPUT

➤ Actual output snippet

```
usnain8721@Haier-PC:/mnt/c/hybrid$ mpiexec -np 4 ./a
rocess_id:0-->thread#3 is sending 8 to process 1
rocess_id:0-->thread#3 is sending 8 to process 2
rocess_id:0-->thread#3 is sending 8 to process 3
rocess_id:0-->thread#0 is sending 2 to process 1
rocess_id:0-->thread#0 is sending 2 to process 2
rocess_id:0-->thread#0 is sending 2 to process 3
rocess_id:0-->thread#1 is sending 4 to process 1
rocess_id:0-->thread#1 is sending 4 to process 2
rocess_id:0-->thread#1 is sending 4 to process 3
rocess-id:1 Received 8 from thread#3 of source
rocess-id:2 Received 8 from thread#3 of source
rocess-id:2 Received 2 from thread#0 of source
rocess-id:2 Received 4 from thread#1 of source
rocess-id:3 Received 4 from thread#1 of source
rocess_id:0-->thread#2 is sending 6 to process 1
rocess_id:0-->thread#2 is sending 6 to process 2
rocess_id:0-->thread#2 is sending 6 to process 3
rocess-id:1 Received 2 from thread#0 of source
rocess-id:1 Received 4 from thread#1 of source
rocess-id:1 Received 6 from thread#2 of source
rocess-id:2 Received 6 from thread#2 of source
rocess-id:3 Received 2 from thread#0 of source
rocess-id:3 Received 8 from thread#3 of source
rocess-id:3 Received 6 from thread#2 of source
ecvbuff=2      4      6      8
```

19

Questions



References

20

1. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (2017). *Introduction to parallel computing*. Redwood City, CA: Benjamin/Cummings.