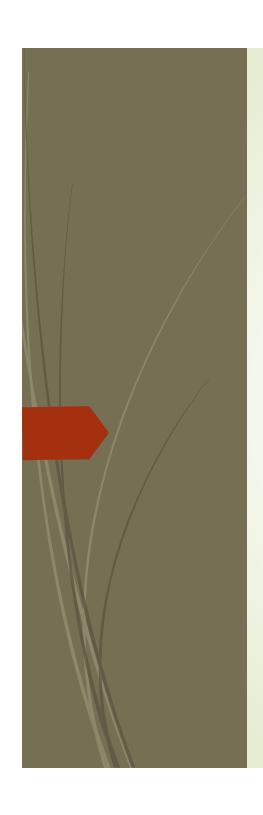
# Parallel and Distributed Computing CS3006

Lecture 11 **Basic Communication Operations**18th March 2024

Dr. Rana Asif Rehman

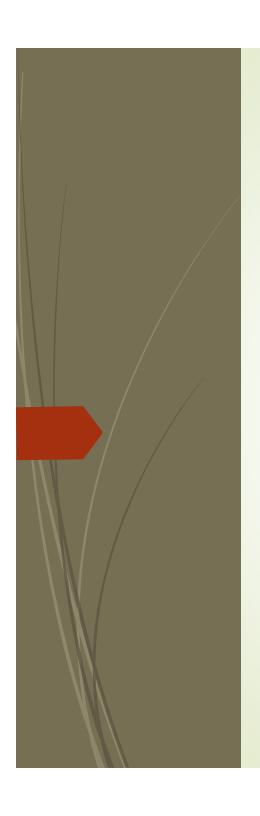


#### **Preliminaries**

- Exchanging the data is fundamental requirement for most of the parallel algorithms
- $-t_s + mt_w$  the simplified communication cost model :-
  - Over distributed memory infrastructure
  - Assuming the cut-through routing
- The chapter is about commonly used basic communication patterns over the different interconnections
  - We shall drive communication costs of these operations on different interconnections.

#### **Assumptions for the Operations**

- Interconnections support cut-through routing
- Communication time between any pair of nodes in the network is same (regardless of the number of intermediate nodes)
- Links are bi-directional
  - The directly connected nodes can simultaneously send messages of m words without any congestion
- Single-port communication model
  - A node can send on only one of its links at a time
  - A node can receive on only one of its links at a time
- However, a node can receive a message while sending another message at the same time on the same or a different link.
- Consider p=2<sup>i</sup> nodes



# One-to-All Broadcast and All-to-One Reduction

(One-to-All Broadcast and All-to-One Reduction)

#### One-to-All Broadcast

- A single process sends identical data to all other processes.
  - Initially one process has data of m size.
  - After broadcast operation, each of the processes have own copy of the m size.

#### **All-to-One Reduction**

- Dual of one-to-all broadcast
- The m-sized data from all processes are combined through an associative operator
- accumulated at a single destination process into one buffer of size m

(One-to-All Broadcast and All-to-One Reduction)



**Figure 4.1** One-to-all broadcast and all-to-one reduction.

(One-to-All Broadcast and All-to-One Reduction)

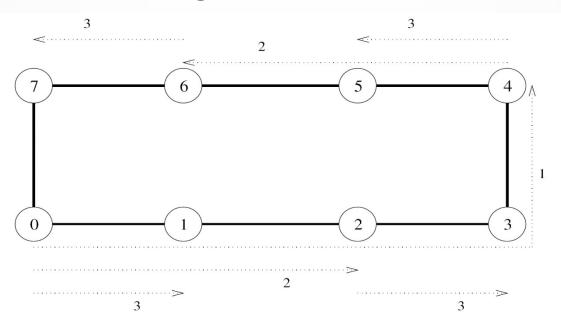
#### **Linear Array or Ring**

- Naïve solution
  - sequentially send p 1 messages from the source to the other p 1 processes
    - Bottle necks, and underutilization of communication network
  - Solution?
- Recursive doubling
  - Source process sends the massage to another process
  - In next communication phase both the processes can simultaneously propagate the message

(One-to-All Broadcast and All-to-One Reduction)

#### **Linear Array or Ring**

Recursive Doubling Broadcast

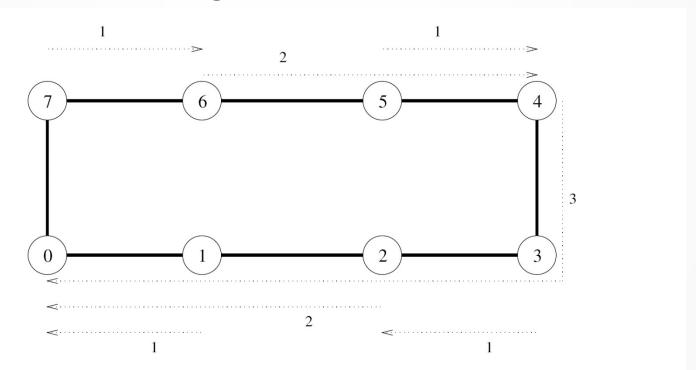


**Figure 4.2** One-to-all broadcast on an eight-node ring. Node 0 is the source of the broadcast. Each message transfer step is shown by a numbered, dotted arrow from the source of the message to its destination. The number on an arrow indicates the time step during which the message is transferred.

(One-to-All Broadcast and All-to-One Reduction)

#### **Linear Array or Ring**

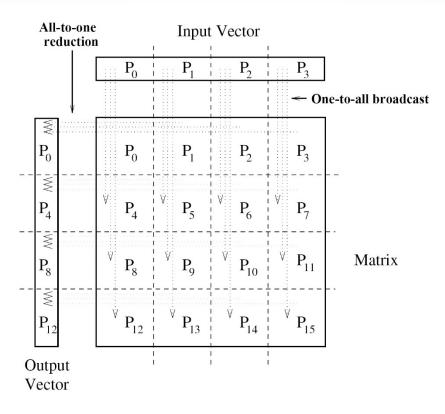
Recursive Doubling Reduction



**Figure 4.3** Reduction on an eight-node ring with node 0 as the destination of the reduction.

(One-to-All Broadcast and All-to-One Reduction)

Matrix-Vector Multiplication (An Application)



**Figure 4.4** One-to-all broadcast and all-to-one reduction in the multiplication of a  $4 \times 4$  matrix with a  $4 \times 1$  vector.

(One-to-All Broadcast and All-to-One Reduction)

#### Mesh

- We can regard each row and column of a square mesh of p nodes as a linear array of nodes
- Communication algorithms on the mesh are simple extensions of their linear array counterparts

#### Broadcast and Reduction

- Two step breakdown:
  - The operation is performed along one by treating the row as linear array
  - II. Then the all the columns are treated similarly

(One-to-All Broadcast and All-to-One Reduction)

Mesh (Broadcast and Reduction)

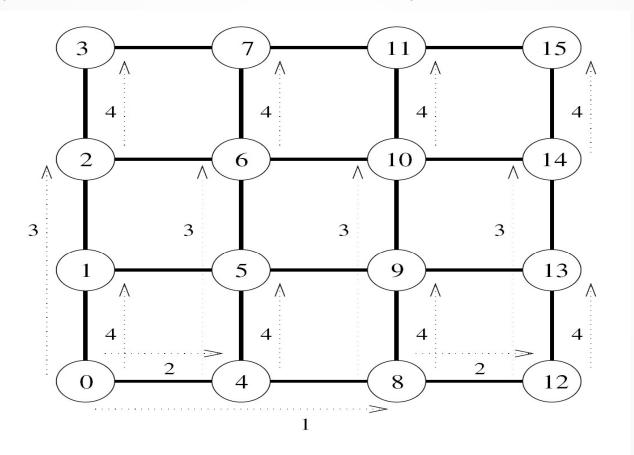


Figure 4.5 One-to-all broadcast on a 16-node mesh.

(One-to-All Broadcast and All-to-One Reduction)

#### **Balanced Binary Tree**

Broadcast

**Figure 4.7** One-to-all broadcast on an eight-node tree.

(One-to-All Broadcast and All-to-One Reduction)

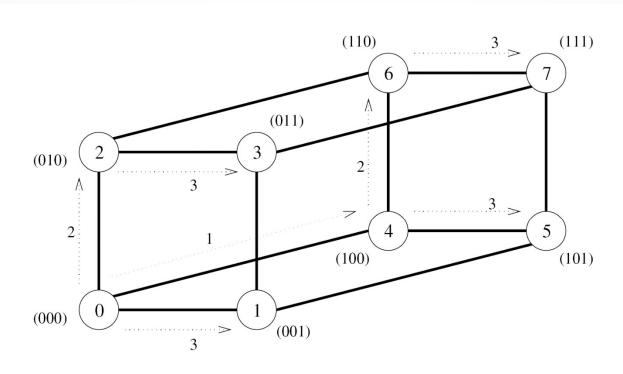
#### Hypercube

- Broadcast
  - Source node first send data to one node in the highest dimension
  - The communication successively proceeds along lower dimensions in the subsequent steps
  - The algorithm is same as used for linear array
    - But, here changing order of dimension does not congest the network

(One-to-All Broadcast and All-to-One Reduction)

#### Hypercube

Broadcast



**Figure 4.6** One-to-all broadcast on a three-dimensional hypercube. The binary representations of node labels are shown in parentheses.

## (One-to-All Broadcast and All-to-One Reduction)

```
1.
     procedure ONE_TO_ALL_BC(d, my_id, X)
     begin
        mask := 2^d - 1;
                                  /* Set all d bits of mask to 1 */
         for i := d - 1 downto 0 do /* Outer loop */
            mask := mask \text{ XOR } 2^i; /* Set bit i of mask to 0 */
5.
6.
            if (my\_id \text{ AND } mask) = 0 then /* If lower i bits of my\_id are 0 */
               if (my\_id \text{ AND } 2^i) = 0 then
7.
8.
                   msg\_destination := my\_id XOR 2^{i};
9.
                   send X to msg_destination;
10.
               else
                   msg\_source := my\_id \text{ XOR } 2^i;
11.
12.
                   receive X from msg_source;
13.
               endelse:
14.
            endif:
15.
         endfor;
     end ONE_TO_ALL_BC
16.
```

**Algorithm 4.1** One-to-all broadcast of a message X from node 0 of a d-dimensional p-node hypercube ( $d = \log p$ ). AND and XOR are bitwise logical-and and exclusive-or operations, respectively.

### (One-to-All Broadcast and All-to-One Reduction)

```
1.
      procedure GENERAL_ONE_TO_ALL_BC(d, my_id, source, X)
2.
      begin
         my\_virtual\_id := my\_id \text{ XOR } source;
         mask := 2^d - 1:
4.
5.
         for i := d - 1 downto 0 do /* Outer loop */
             mask := mask \text{ XOR } 2^i; /* Set bit i of mask to 0 */
6.
             if (my\_virtual\_id \text{ AND } mask) = 0 then
7.
8.
                if (my\_virtual\_id \text{ AND } 2^i) = 0 then
9.
                    virtual\_dest := my\_virtual\_id XOR 2^i;
                    send X to (virtual_dest XOR source);
10.
         /* Convert virtual_dest to the label of the physical destination */
11.
                else
12.
                    virtual\_source := mv\_virtual\_id XOR 2^i;
                    receive X from (virtual_source XOR source);
13.
         /* Convert virtual_source to the label of the physical source */
14.
                endelse;
15.
         endfor:
     end GENERAL_ONE_TO_ALL_BC
16.
```

**Algorithm 4.2** One-to-all broadcast of a message *X* initiated by *source* on a *d*-dimensional hypothetical hypercube. The AND and XOR operations are bitwise logical operations.

### (One-to-All Broadcast and All-to-One Reduction)

```
procedure ALL_TO_ONE_REDUCE(d, my\_id, m, X, sum)
1.
2.
         begin
3.
              for j := 0 to m - 1 do sum[j] := X[j];
4.
              mask := 0:
5.
              for i := 0 to d - 1 do
                   /* Select nodes whose lower i bits are 0 */
                  if (my\_id \text{ AND } mask) = 0 then
6.
7.
                       if (my\_id \text{ AND } 2^i) \neq 0 then
8.
                            msg\_destination := my\_id XOR 2^i;
9.
                            send sum to msq\_destination;
10.
                       else
11.
                            msq\_source := my\_id XOR 2^i;
12.
                            receive X from msq\_source;
13.
                            for j := 0 to m-1 do
14.
                                 sum[j] := sum[j] + X[j];
15.
                       endelse:
                   mask := mask XOR 2^i; /* Set bit i of mask to 1 */
16.
17.
              endfor:
         end ALL_TO_ONE_REDUCE
18.
```

(One-to-All Broadcast and All-to-One Reduction)

#### **Cost Estimation**

- Broadcast needs log(p) point-to-point simple message transfer steps.
- Message size of each transfer is m
- lacktriangle Time for each of the transfers is:  $t_s + mt_w$

Hence cost for log(p)transfers= $T = (t_s + mt_w) \log p$ 

## Questions



Parallel and Distributed Computing (CS3006) - Spring 2024

## References

- 1. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (1994). *Introduction to parallel computing* (Vol. 110). Redwood City, CA: Benjamin/Cummings.
- 2. Quinn, M. J. Parallel Programming in C with MPI and OpenMP,(2003).