



Parallel and Distributed Computing

CS3006

Lecture 9

OpenMP-II

6th March 2024

Dr. Rana Asif Rehman

Review of OpenMP Library Functions

➤ Controlling Number of Threads and Processors

- `void omp_set_num_threads (int num_threads);`
- `int omp_get_num_threads ();`
- `int omp_get_max_threads ();`
- `int omp_get_thread_num ();`
- `int omp_get_num_procs ();`
- `int omp_in_parallel ();`

➤ Controlling and Monitoring Thread Creation

- `void omp_set_dynamic (int dynamic_threads);`
- `int omp_get_dynamic ();`
- `void omp_set_nested (int nested);`
- `int omp_get_nested ();`



OpenMP

#pragma omp directive [clause list]

OpenMP Directives

- `#pragma omp parallel`
- `#pragma omp for`

One more thing to note

Difference between *omp for* and *omp parallel*

```
1 #pragma omp parallel
2 {
3   #pragma omp for
4   for (i = 0 < i < n; i++) {
        //omp for schedules/distributes iterations between the threads
5       /* body of parallel for loop */
6   }
7 }
```

is same as

```
1 #pragma omp parallel for
3   for (i = 0 < i < n; i++) {
4       /* body of parallel for loop */
6   }
```

Some Useful Clauses in OpenMP


- A clause is an optional, additional component to a pragma
- **Private:** The private clause directs the compiler to make one or more variables private

```
int k=3;
#pragma omp parallel for default(shared) private(j) shared(k)
for (i = 0; i < n; i++)
    for (j = 0; j < n; j++)
        a[i][j] = MIN(a[i][j], a[i][k]+tmp) ;
```

Comments:

- Here the private variable j is undefined -
 - when this parallel construct is entered
 - when this parallel construct is exited

Some Useful Clauses in OpenMP

- 
- ➔ **firstprivate:** It directs the compiler to create private variables having initial values identical to the value of the variable controlled by the master thread as the loop is entered.

```
s = complex_function();  
#pragma omp parallel for firstprivate(s) num_threads(2)  
for (i = 0 ; i < n ; i++) (  
    s = s*omp_get_thread_num();  
    printf("S is %d at thread#%d\n", s,omp_get_thread_num());  
}
```

Some useful clauses




➡ lastprivate:

consider the following code

```
s = complex_function();  
#pragma omp parallel for private(j) firstprivate(s)  
for (i = 0 ; i < n ; i ++ ) {  
    s += 1  
}  
printf("s after join:%d\n",s); //undefined value
```

Some Useful Clauses in OpenMP

- 
- ➔ **lastprivate:** used to copy back to the master thread's copy of the variable, the private copy of the variable from the thread that executed the last iteration.

```
s = complex_function();  
#pragma omp parallel for private(j) firstprivate(s) lastprivate(s)  
for (i = 0 ; i < n ; i++) (  
    s += 1;  
}  
printf("s after join:%d\n",s); //value of s as it was for last iteration of the loop
```


Reduction clause

- Reductions are so common that OpenMP provides support for them
- May add reduction clause to **parallel for** pragma
- Specify reduction operation and reduction variable
- OpenMP takes care of storing partial results in private variables and combining partial results after the loop
- The reduction clause has this syntax:
reduction (<op> :<variable>)

- **Operators**

- + Sum
- * Product
- & Bitwise and
- | Bitwise or
- ^ Bitwise exclusive or
- && Logical and
- || Logical or

Reduction clause



```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x) reduction(+:area)
for (i = 0; i < n; i++) {
    x = (i+0.5)/n;
    area += 4.0/(1.0 + x*x);
}
pi = area / n;
```

Conditional Parallelism Clause

- **if Clause:** The if clause gives us the ability to direct the compiler to insert code that determines at run-time whether the loop should be executed in parallel or not.
- **The clause has this syntax:** `if (<scalar expression>)`

```
double area, pi, x;
int i, n;
...
area = 0.0;
#pragma omp parallel for private(x) reduction(+:area) if (n>5000)
for (i = 0; i < n; i++) {
    x = (i+0.5)/n;
    area += 4.0/(1.0 + x*x);
}
pi = area / n;
```

Scheduling Loops (a clause)

- **Scheduling** the loops means dividing number of iterations between the processes.

- Syntax of schedule clause

`schedule (<type>[, <chunk>])`

- Schedule type is required but, chunk size optional
- A chunk is a contiguous range of iterations
 - Increasing chunk size reduces scheduling overhead and may increase cache hit rate [due to operations on contiguous memory locations]
 - Decreasing chunk size allows finer balancing of workloads

Scheduling Loops

1. **Static:** `schedule(static[, chunk-size])`

- Splits the iteration space into equal chunks of size `chunk-size` and assigns them to threads in a round-robin fashion.
- When no `chunk-size` is specified, the iteration space is split into as many chunks as there are threads (i.e., size of each is $n/\text{tot.threads}$) and one chunk is assigned to each thread.
- Decision about work division is done before actually executing the code.
- Results in lower scheduling overhead. But, can cause load-imbalance if all processors are not of same compute-capability.

Scheduling Loops

1. Static: `schedule(static[, chunk-size])`

Example when reducing chunk size improves load-balancing

```
#pragma omp parallel for private (j) schedule(static, 1)
for (i = 0; i < n; i++)
    for (j = i; j < n; j++)
        a[i][j] = complex_func(i, j);
```

Scheduling Loops

2. **Dynamic:** `schedule(dynamic[, chunk-size])`

- The iteration space is partitioned into chunks given by chunk-size
- Initially every thread is assigned single chunk. The decision for remaining iteration chunks is done on run-time
- This means chunk is assigned to threads as they become idle.
- This takes care of the temporal imbalances resulting from static scheduling.
- If no chunk-size is specified, it defaults to a single iteration per chunk

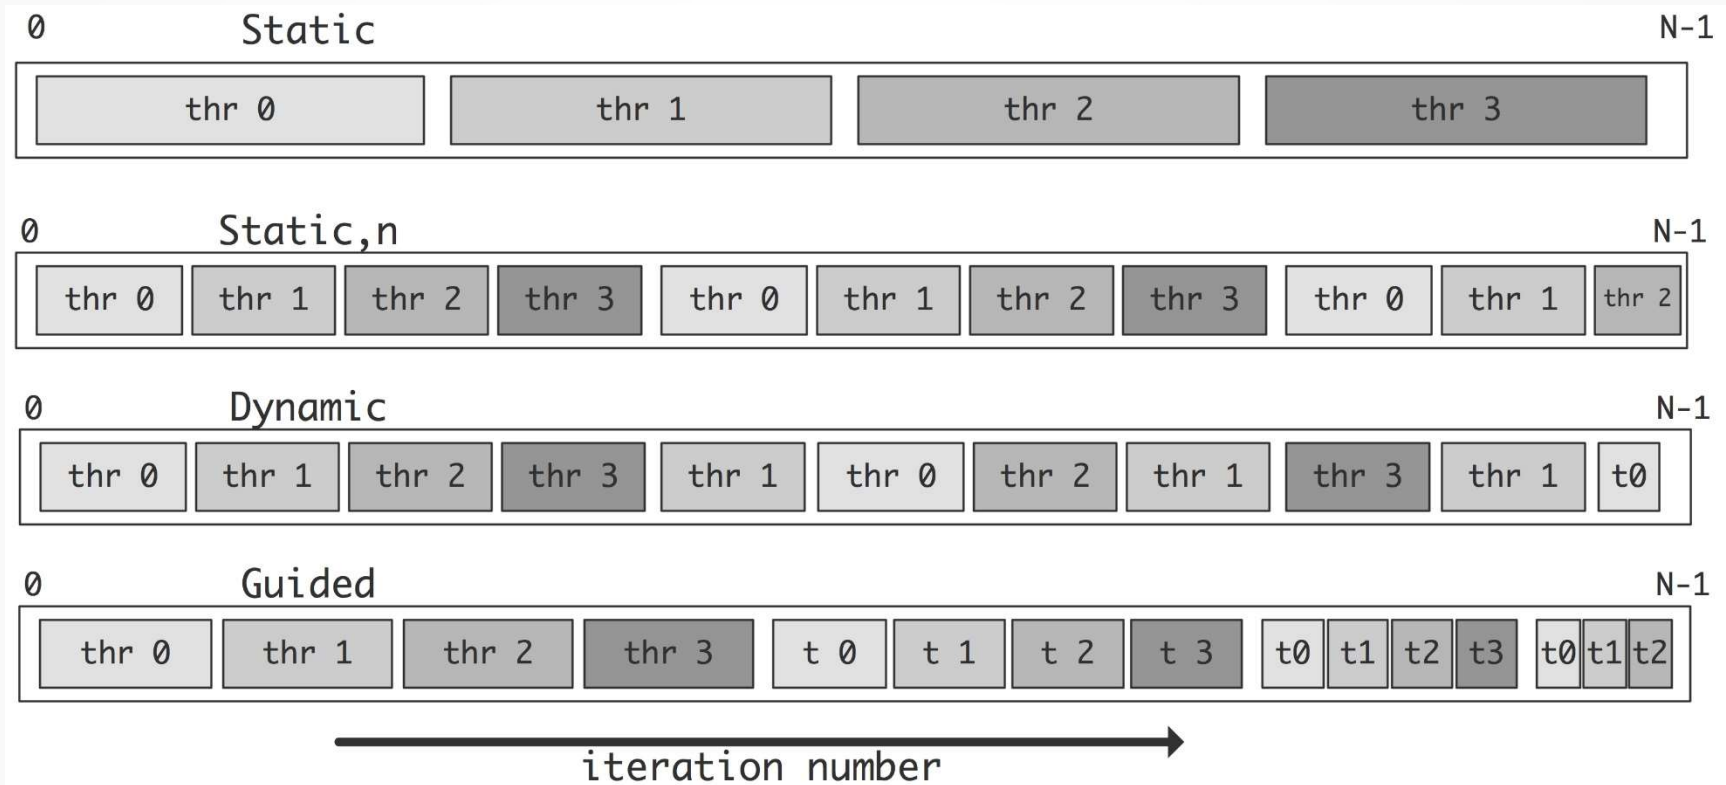
Scheduling Loops

3. Guided:

- **schedule(guided, C)**: dynamic allocation of chunks to tasks using guided self-scheduling heuristic. Initial chunks are bigger, later chunks are smaller, **minimum** chunk size is C.
- **schedule(guided)**: guided self-scheduling with minimum chunk size 1

4. **schedule(runtime)**: schedule chosen at run-time based on value of OMP_SCHEDULE env variable.

Scheduling Loops(Summary)





No Wait Clause

- In order to avoid implicit barrier
- A thread can easily move to next after completed its assign task/iterations

#pragma omp parallel for nowait

Functional / Task Parallelism in OpenMP





Functional/Task Parallelism

#pragma omp sections [clause list]


Functional/Task Parallelism

- If your code is based on different segments or sections that can be executed in parallel.
- Also known to as task parallelism


```
v = alpha();  
w = beta();  
x = gamma(v,w);  
y = delta();
```

```
printf("%6.2f\n",epsilon(x,y));
```

- Can execute alpha, beta, delta parallelly
- Remaining ones are executed sequentially according to the dependency.



parallel sections, section pragmas



```
#pragma omp parallel sections
{
    #pragma omp section //[optional for 1st block]
    v = alpha();
    #pragma omp section
    W = beta();
    #pragma omp section
    y = delta();
}
x = gamma(v,w);
printf("%6.2f\n", epsilon(x,y));
```

- #pragma omp parallel sections creates a team of threads which executes the sections in the region parallelly
 - Sections that can be executed parallel are preceded by 'omp section' pragma.


Functional Parallelism

Another approach


```
v = alpha();  
w = beta();  
x = gamma(v,w);  
y = delta();
```

```
printf("%6.2f\n",epsilon(x,y));
```


- Execute alpha and beta in parallel.
- Execute gamma and delta in parallel



omp sections **pragma**

- 
- Appears inside a parallel block of code
 - This pragma distributes enclosed sections among the threads in the team
 - The difference between *omp parallel sections* and *omp sections* is that,
 - **Omp parallel sections** generate its own team of threads
 - While simple **omp sections** pragma uses existing team of threads and distributes section among the threads
 - If multiple sections pragmas are inside one parallel block, may reduce fork/join costs

sections pragma



```
#pragma omp parallel num_threads(2)
{
    #pragma omp sections
    {
        #pragma omp section //optional
        v = alpha();
        #pragma omp section
        w = beta();
    } // here an implicit barrier exists

    #pragma omp sections
    {
        x = gamma(v, w);
        #pragma omp section
        y = delta();
    }
}

printf ("%6.2f\n", epsilon(x,y));
```

Questions



Parallel and Distributed Computing
(CS3006) - Spring 2024



References

1. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (2017). *Introduction to parallel computing*. Redwood City, CA: Benjamin/Cummings.