

# Parallel and Distributed Computing

## CS3006

Lecture 6

**Decomposition Techniques**

21st February 2024

Dr. Rana Asif Rehman

# Agenda

- A Quick Review
- **Decomposition Techniques**
  - Recursive
  - Data-decomposition
  - Exploratory
  - Speculative
  - Hybrid

# Quick Review to the Previous Lecture

3

- Parallel Algorithm Design Life Cycle
- Tasks, Decomposition, and Task-dependency graphs
- Granularity
  - Fine-grained
  - Coarse-grained
- Concurrency
  - Max-degree of concurrency
  - Critical path length
  - Average-degree of concurrency

# Decomposition Techniques

4

- The process of decomposing larger problems into smaller tasks for concurrent executions, is known to as decomposition.
- The techniques that facilitate this decomposition are known to as decomposition techniques.
- **Common techniques:**
  - Recursive
  - Data-decomposition
  - Exploratory decomposition
  - Speculative decomposition
  - Hybrid
- Recursive and data decompositions are relatively general purpose
- Exploratory and speculative are special purpose in nature

# Decomposition Techniques

5

## 1. Recursive Task Decomposition

- Recursive decomposition is a method for inducing concurrency in the problems that can be solved using divide and conquer strategy
- Divides each problem into a set of independent sub-problems
- Each one of these subproblems is solved by recursively applying a similar division into smaller subproblems followed by a combination of their results
- A natural concurrency exists as different subproblems can be solved concurrently.

## 6

The diagram shows a binary tree structure for Huffman coding. The root node is a box containing the sequence [5, 12, 11, 1, 10, 6, 8, 3, 7, 4, 9, 2]. It branches into two nodes: [1, 3, 4, 2] on the left and [5, 12, 11, 10, 6, 8, 7, 9] on the right. The left node branches into [1, 2] and [3, 4], which then branch into leaf nodes [1], [2], [3], and [4]. The right node branches into [5, 6, 8, 7] and [9, 12, 11, 10]. The [5, 6, 8, 7] node branches into [5, 6] and [7, 8], which then branch into leaf nodes [5], [6], [7], and [8]. The [9, 12, 11, 10] node branches into [9] and [10, 12, 11]. The [10, 12, 11] node branches into [10] and [11, 12], which then branches into leaf nodes [11] and [12]. Shaded boxes indicate the current step in the construction process.

Parallel and Distributed Computing (CS3006) - Spring 2024

# Decomposition Techniques

7

**Recursive  
Decomposition**  
(Modifying  
simple problem  
to support  
recursive  
decomposition)

---

```
1.  procedure SERIAL_MIN ( $A, n$ )
2.  begin
3.     $min = A[0]$ ;
4.    for  $i := 1$  to  $n - 1$  do
5.      if ( $A[i] < min$ )  $min := A[i]$ ;
6.    endfor;
7.    return  $min$ ;
8.  end SERIAL_MIN
```

---

**Algorithm 3.1** A serial program for finding the minimum in an array of numbers  $A$  of length  $n$ .

---

```
1.  procedure RECURSIVE_MIN ( $A, n$ )
2.  begin
3.    if ( $n = 1$ ) then
4.       $min := A[0]$ ;
5.    else
6.       $lmin := RECURSIVE\_MIN (A, n/2)$ ;
7.       $rmin := RECURSIVE\_MIN (&(A[n/2]), n - n/2)$ ;
8.      if ( $lmin < rmin$ ) then
9.         $min := lmin$ ;
10.     else
11.        $min := rmin$ ;
12.     endelse;
13.  endelse;
14.  return  $min$ ;
15. end RECURSIVE_MIN
```


---

**Algorithm 3.2** A recursive program for finding the minimum in an array of numbers  $A$  of length  $n$ .

# Decomposition Techniques

8

**Recursive Decomposition** (Modifying simple problem to support recursive decomposition)



**Figure 3.9** The task-dependency graph for finding the minimum number in the sequence {4, 9, 1, 7, 8, 11, 2, 12}. Each node in the tree represents the task of finding the minimum of a pair of numbers.



# Decomposition Techniques

9

## 2. Data Decomposition

- Powerful and commonly used method
- Two step procedure:
  1. Partition data on which computation is to be performed
  2. This data partitioning is used to induce a partitioning of the computations into tasks.
- **Partitioning output data**
  - Used where each element of the output can be computed independently of others as a function of the input.
  - Partitioning of the output data automatically induces a decomposition of the problems into tasks
  - each task is assigned the work of computing a portion of the output

# Decomposition Techniques

10

## Data Decomposition (Partitioning Output Data)

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

(a)

Task 1:  $C_{1,1} = A_{1,1}B_{1,1} + A_{1,2}B_{2,1}$

Task 2:  $C_{1,2} = A_{1,1}B_{1,2} + A_{1,2}B_{2,2}$

Task 3:  $C_{2,1} = A_{2,1}B_{1,1} + A_{2,2}B_{2,1}$

Task 4:  $C_{2,2} = A_{2,1}B_{1,2} + A_{2,2}B_{2,2}$

(b)

**Figure 3.10** (a) Partitioning of input and output matrices into  $2 \times 2$  submatrices. (b) A decomposition of matrix multiplication into four tasks based on the partitioning of the matrices in (a).

# Decomposition Techniques

11

## Data Decomposition (Partitioning Output Data)

Decomposition I	Decomposition II
Task 1: $C_{1,1} = A_{1,1} B_{1,1}$	Task 1: $C_{1,1} = A_{1,1} B_{1,1}$
Task 2: $C_{1,1} = C_{1,1} + A_{1,2} B_{2,1}$	Task 2: $C_{1,1} = C_{1,1} + A_{1,2} B_{2,1}$
Task 3: $C_{1,2} = A_{1,1} B_{1,2}$	Task 3: $C_{1,2} = A_{1,2} B_{2,2}$
Task 4: $C_{1,2} = C_{1,2} + A_{1,2} B_{2,2}$	Task 4: $C_{1,2} = C_{1,2} + A_{1,1} B_{1,2}$
Task 5: $C_{2,1} = A_{2,1} B_{1,1}$	Task 5: $C_{2,1} = A_{2,2} B_{2,1}$
Task 6: $C_{2,1} = C_{2,1} + A_{2,2} B_{2,1}$	Task 6: $C_{2,1} = C_{2,1} + A_{2,1} B_{1,1}$
Task 7: $C_{2,2} = A_{2,1} B_{1,2}$	Task 7: $C_{2,2} = A_{2,1} B_{1,2}$
Task 8: $C_{2,2} = C_{2,2} + A_{2,2} B_{2,2}$	Task 8: $C_{2,2} = C_{2,2} + A_{2,2} B_{2,2}$

**Figure 3.11** Two examples of decomposition of matrix multiplication into eight tasks.

# Decomposition Techniques

(a) Transactions (input), itemsets (input), and frequencies (output)

Database Transactions	A, B, C, E, G, H	Itemsets	A, B, C	Itemset Frequency	1
	B, D, E, F, K, L		D, E		3
	A, B, F, H, L		C, F, G		0
	D, E, F, H		A, E		2
	F, G, H, K,		C, D		1
	A, E, F, K, L		D, K		2
	B, C, D, G, H, L		B, C, F		0
	G, H, I,		C, D, K		0
	D, E, F, K, L				
	F, G, H, L				

(b) Partitioning the frequencies (and itemsets) among the tasks

Database Transactions	A, B, C, E, G, H	Itemsets	A, B, C	Itemset Frequency	1
	B, D, E, F, K, L		D, E		3
	A, B, F, H, L		C, F, G		0
	D, E, F, H		A, E		2
	F, G, H, K,				
	A, E, F, K, L				
	B, C, D, G, H, L				
	G, H, L				
	D, E, F, K, L				
	F, G, H, L				

task 1

Database Transactions	A, B, C, E, G, H	Itemsets	C, D	Itemset Frequency	1
	B, D, E, F, K, L		D, K		2
	A, B, F, H, L		B, C, F		0
	D, E, F, H		C, D, K		0
	F, G, H, K,				
	A, E, F, K, L				
	B, C, D, G, H, L				
	G, H, L				
	D, E, F, K, L				
	F, G, H, L				

task 2

**Figure 3.12** Computing itemset frequencies in a transaction database.

# Decomposition Techniques

13

## Data Decomposition

### Partitioning input data

- In many algorithms, it is not possible or desirable to partition the output data.
  - The output may be a single unknown value. Such as in case of finding sum, minimum, maximum or frequencies of a number.
- It is sometimes possible to partition the input data, and then use this partitioning to induce concurrency
- A task is created for each partition of the input data and this task performs as much computation as possible using these local data
- Then local solutions are combined to generate a global solution

# Decomposition Techniques

14

## Partitioning input data

(a) Partitioning the transactions among the tasks

Database Transactions	A, B, C, E, G, H	Itemsets	A, B, C	Itemset Frequency	1
	B, D, E, F, K, L		D, E		2
	A, B, F, H, L		C, F, G		0
	D, E, F, H		A, E		1
	F, G, H, K,		C, D		0
			D, K		1
			B, C, F		0
			C, D, K		0

task 1

Database Transactions		Itemsets	A, B, C	Itemset Frequency	0
			D, E		1
			C, F, G		0
	A, E, F, K, L		A, E		1
	B, C, D, G, H, L		C, D		1
	G, H, L		D, K		1
	D, E, F, K, L		B, C, F		0
	F, G, H, L		C, D, K		0

task 2

# Decomposition Techniques

15

## Data Decomposition

Partitioning both input and output data

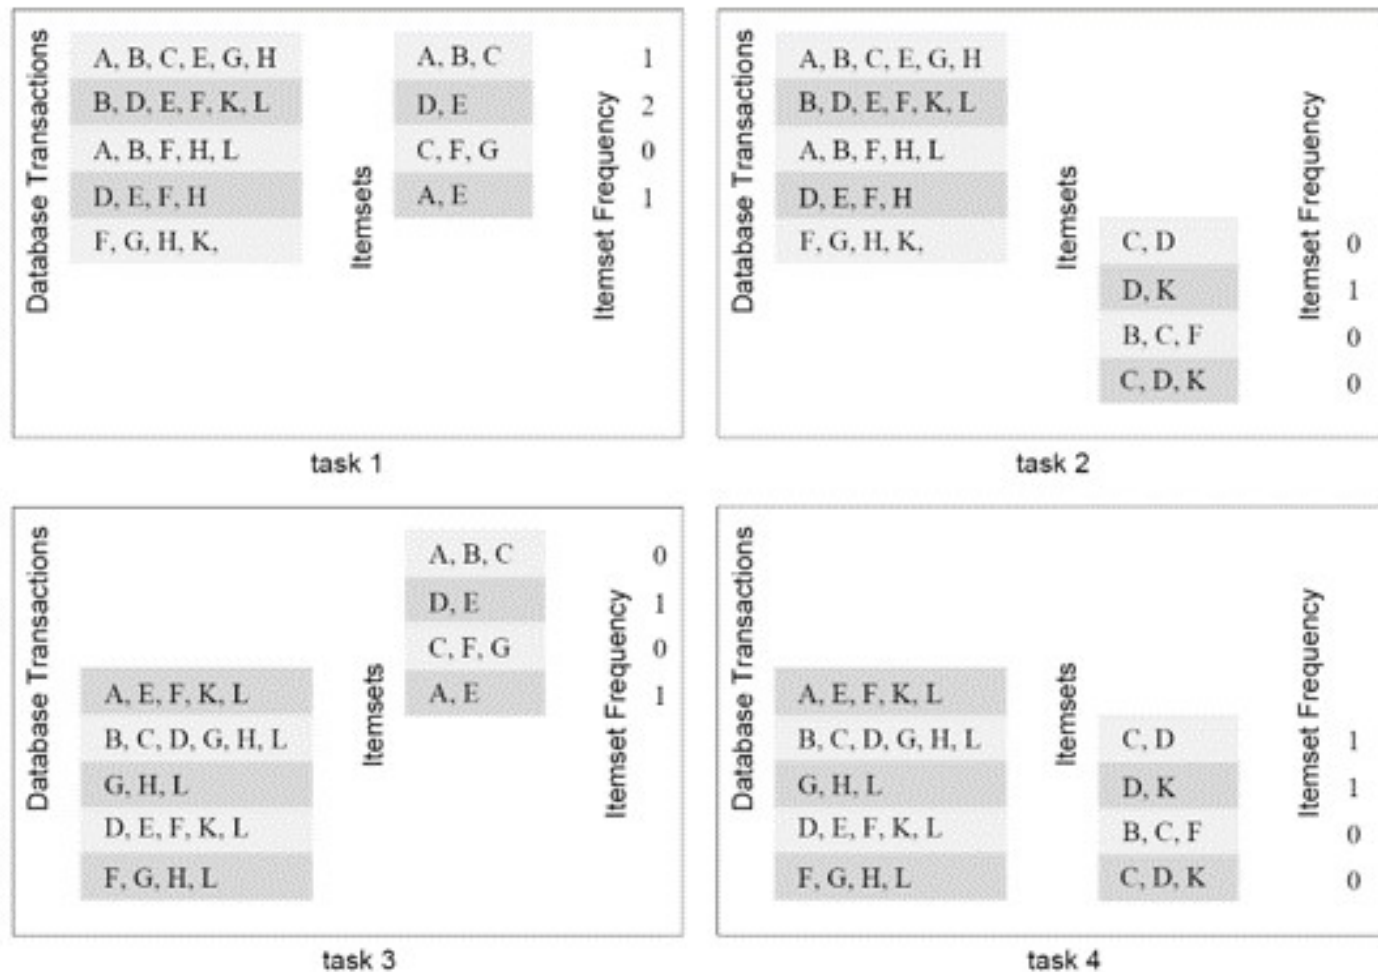
- Consider the problems where output data-partitioning is possible
- Here, partitioning the input also, can offer additional concurrency
- The next example shows 4-way decomposition of the previous example based on both input-output partitioning.

# Decomposition Techniques

16

Partitioning both input and output data

(b) Partitioning both transactions and frequencies among the tasks



**Figure 3.13** Some decompositions for computing itemset frequencies in a transaction database.



# Decomposition Techniques

17

Partitioning both intermediate data

Stage I

$$\begin{pmatrix} A_{1,1} & A_{1,2} \\ A_{2,1} & A_{2,2} \end{pmatrix} \cdot \begin{pmatrix} B_{1,1} & B_{1,2} \\ B_{2,1} & B_{2,2} \end{pmatrix} \rightarrow \left( \begin{pmatrix} D_{1,1,1} & D_{1,1,2} \\ D_{1,2,2} & D_{1,2,2} \end{pmatrix} \right)$$


Stage II

$$\begin{pmatrix} D_{1,1,1} & D_{1,1,2} \\ D_{1,2,2} & D_{1,2,2} \end{pmatrix} + \begin{pmatrix} D_{2,1,1} & D_{2,1,2} \\ D_{2,2,2} & D_{2,2,2} \end{pmatrix} \rightarrow \begin{pmatrix} C_{1,1} & C_{1,2} \\ C_{2,1} & C_{2,2} \end{pmatrix}$$

A decomposition induced by a partitioning of  $D$

- Task 01:  $D_{1,1,1} = A_{1,1} B_{1,1}$
- Task 02:  $D_{2,1,1} = A_{1,2} B_{2,1}$
- Task 03:  $D_{1,1,2} = A_{1,1} B_{1,2}$
- Task 04:  $D_{2,1,2} = A_{1,2} B_{2,2}$
- Task 05:  $D_{1,2,1} = A_{2,1} B_{1,1}$
- Task 06:  $D_{2,2,1} = A_{2,2} B_{2,1}$
- Task 07:  $D_{1,2,2} = A_{2,1} B_{1,2}$
- Task 08:  $D_{2,2,2} = A_{2,2} B_{2,2}$
- Task 09:  $C_{1,1} = D_{1,1,1} + D_{2,1,1}$
- Task 10:  $C_{1,2} = D_{1,1,2} + D_{2,1,2}$
- Task 11:  $C_{2,1} = D_{1,2,1} + D_{2,2,1}$
- Task 12:  $C_{2,2} = D_{1,2,2} + D_{2,2,2}$

**Figure 3.15** A decomposition of matrix multiplication based on partitioning the intermediate three-dimensional matrix.



**Figure 3.16** The task-dependency graph of the decomposition shown in Figure 3.15.

# Decomposition Techniques

19

## Owner Compute Rule

- Task decomposition based on data-partitioning is widely known as owner compute rule.
- Two types of partitioning hence, two definitions:
  1. If we assign partitions of the input data to tasks:
    - The rule means that a task performs all the computations that can be done using these data
  2. If we assign partition of output data to the tasks:
    - The rule means that a task computes all the data in the partition assigned to it (portion of the output).

# Decomposition Techniques

20

## 3. Exploratory Decomposition

- Specially used to decompose the problems having underlying computation like search-space exploration.
- Steps:
  1. Partition the search space into smaller parts
  2. Search each one of these parts concurrently, until the desired solutions are found.

# Decomposition Techniques

21

## 3. Exploratory Decomposition

1	2	3	4
5	6	7	8
9	10	7	11
13	14	15	12

(a)

1	2	3	4
5	6	7	8
9	10	11	11
13	14	15	12

(b)


1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	12

(c)

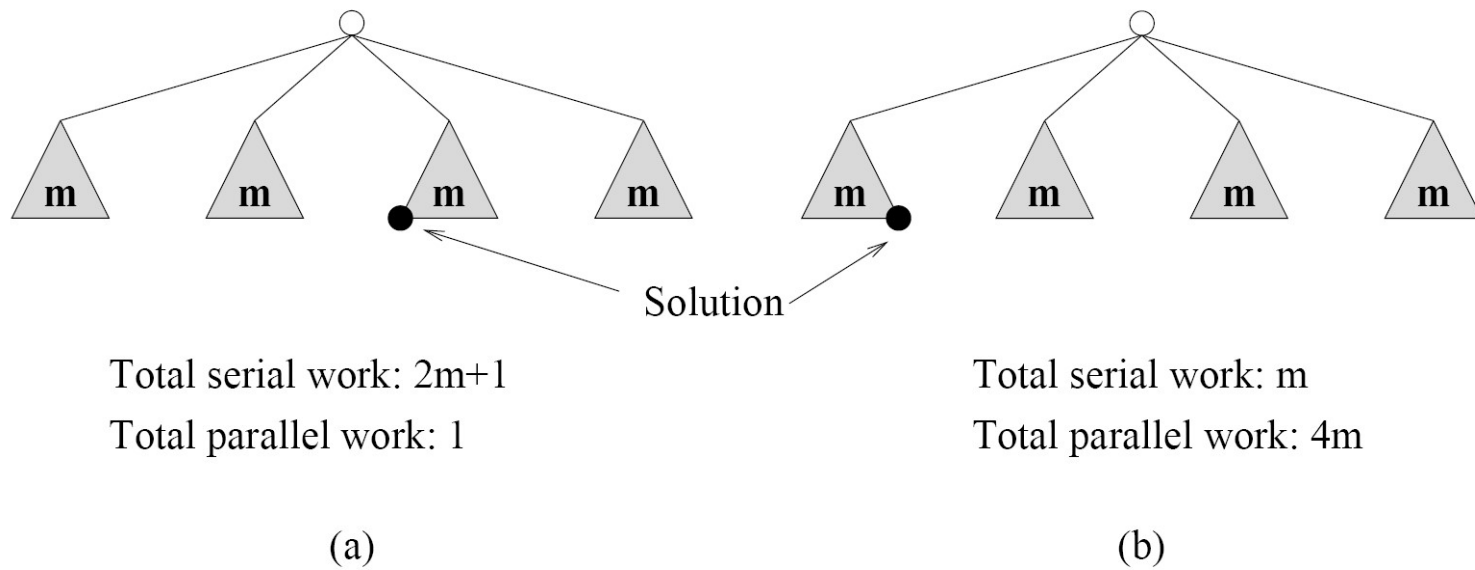
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	

(d)

**Figure 3.17** A 15-puzzle problem instance showing the initial configuration (a), the final configuration (d), and a sequence of moves leading from the initial to the final configuration.



**Figure 3.18** The states generated by an instance of the 15-puzzle problem.



**Figure 3.19** An illustration of anomalous speedups resulting from exploratory decomposition.

# Decomposition Techniques

24

## 4. Speculative Decomposition

- Usually used in the problems where different input values or output of previous stage causes many computationally intensive branches.
- Speculation is something like Gamble or Risk or preliminary guess.
- Steps:
  - Speculate(guess) the output of previous stage
  - Start performing computations in the next stage before even the completion of the previous stage.
  - After availability of the output of previous stage, if speculation was correct than most of the computation for next step would have already been done.



# Decomposition Techniques

25

## 4. Speculative Decomposition

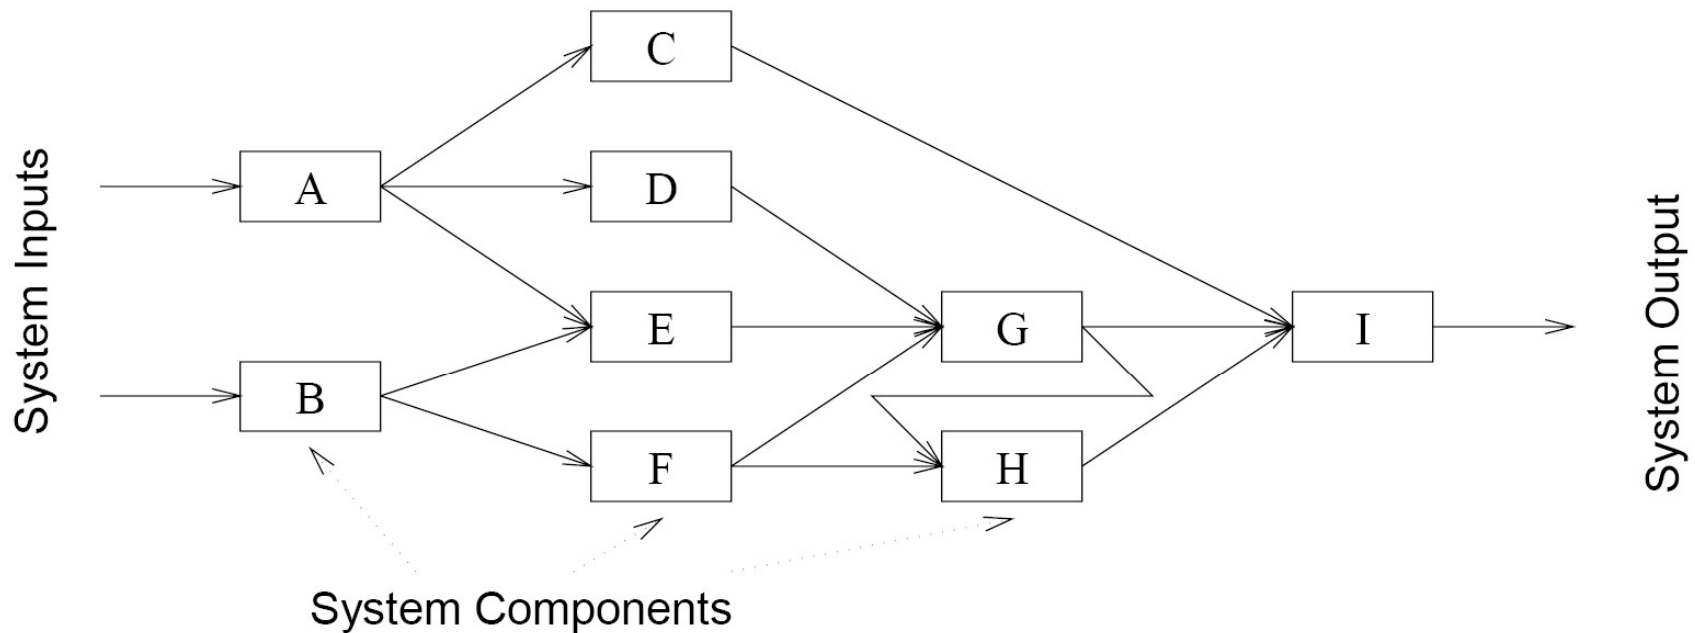
➡ Switch Example Algorithm:

- 1: Calculate expression for the switch condition → task 0
- 2:     Case 0: Multiply vector **b** with matrix **A** → task 1
- 3:     Case 1: Multiply vector **c** with matrix **A** → task 2
- 4:     Case 2: Multiply vector **d** with matrix **A** → task 3
- 5: display result → task 4

# Decomposition Techniques

26

## 4. Speculative Decomposition




**Figure 3.20** A simple network for discrete event simulation.

# Decomposition Techniques

27

## 5. Hybrid Decomposition

- Decomposition techniques are not exclusive
  - We often need to combine them together



**Figure 3.21** Hybrid decomposition for finding the minimum of an array of size 16 using four tasks.

# Questions

28

Parallel and Distributed Computing  
(CS3006) - Spring 2024



# References

29

1. Kumar, V., Grama, A., Gupta, A., & Karypis, G. (1994). *Introduction to parallel computing* (Vol. 110). Redwood City, CA: Benjamin/Cummings.
2. Quinn, M. J. *Parallel Programming in C with MPI and OpenMP*, (2003).