

An abstract wireframe sculpture of a human head, composed of numerous thin, white lines forming a complex, geometric mesh. The sculpture is positioned on the left side of the frame, facing right. The background is a solid, dark gray. The text "Generative AI" is centered over the sculpture.

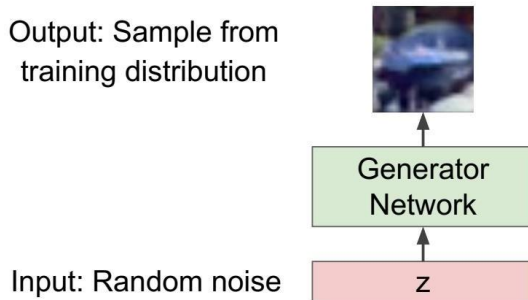
# Generative AI



# Generative Adversarial Networks (GANs)

# Generative Adversarial Networks

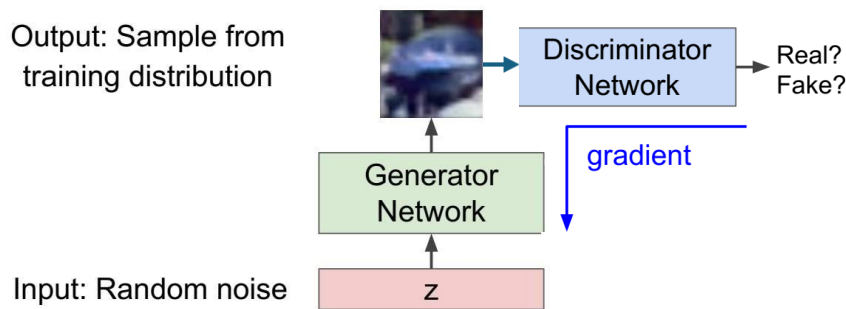
- Problem: Want to sample from complex, high-dimensional training distribution. No direct way to do this!
- Solution: Sample from a simple distribution we can easily sample from, e.g. random noise. Learn transformation to training distribution.
- Q: What can we use to represent this complex transformation?
- A: A neural network!



But we don't know which sample  $z$  maps to which training image -> can't learn by reconstructing training images

# Generative Adversarial Networks

- Problem: Want to sample from complex, high-dimensional training distribution. No direct way to do this!
- Solution: Sample from a simple distribution we can easily sample from, e.g. random noise. Learn transformation to training distribution.
- Q: What can we use to represent this complex transformation?
- A: A neural network! **Solution: Use a discriminator network to tell whether the generated image is within data distribution or not**

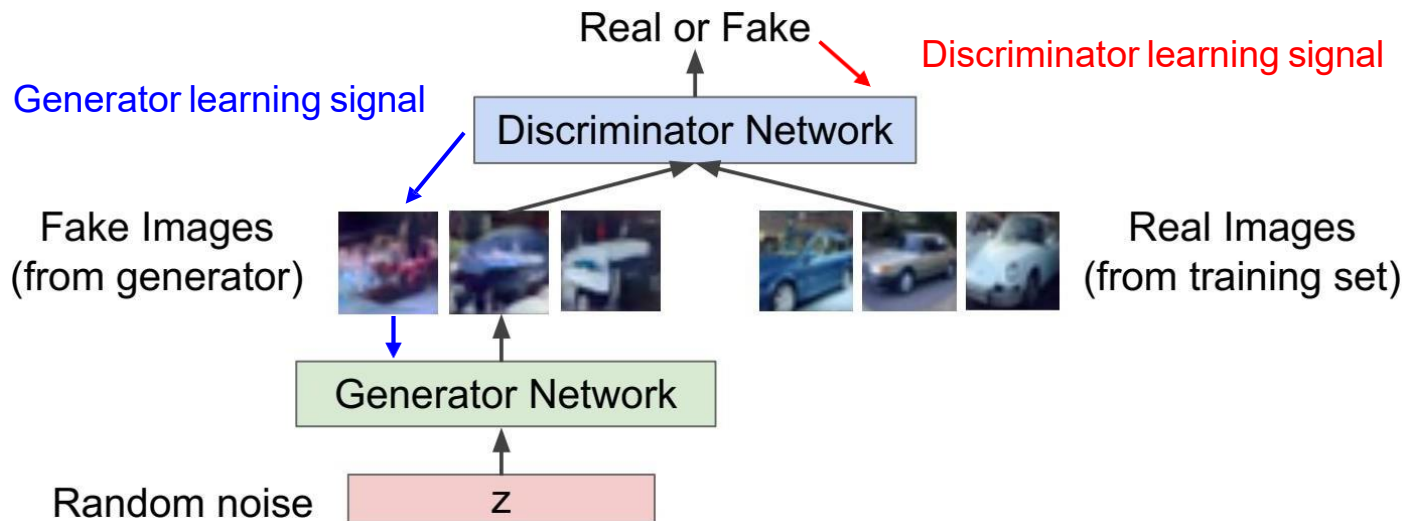


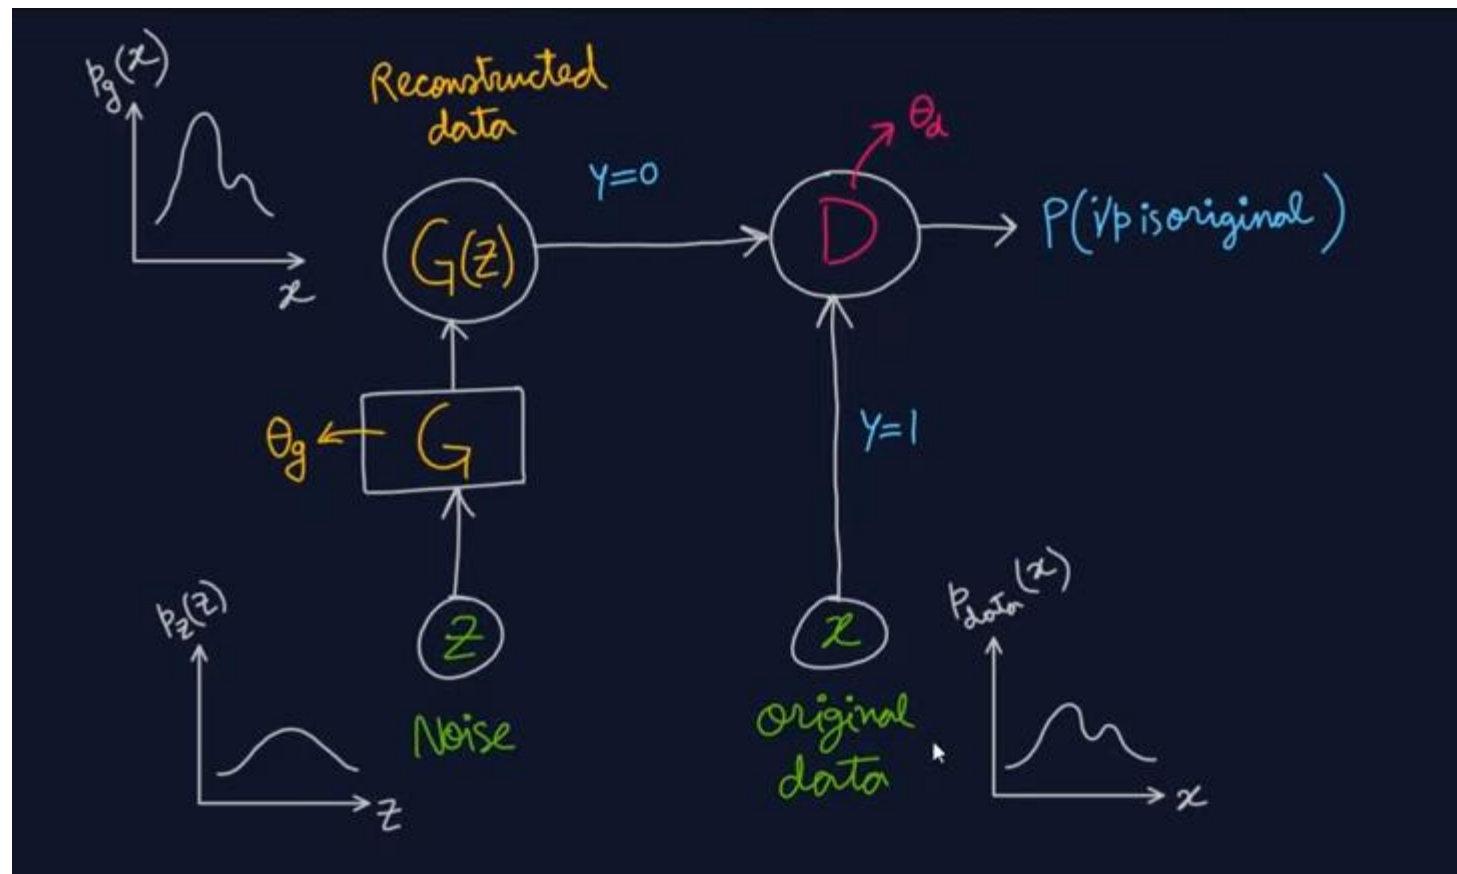
But we don't know which sample  $z$  maps to which training image -> can't learn by reconstructing training images

# Training GANs: Two Player Game

**Discriminator network:** try to distinguish between real and fake images

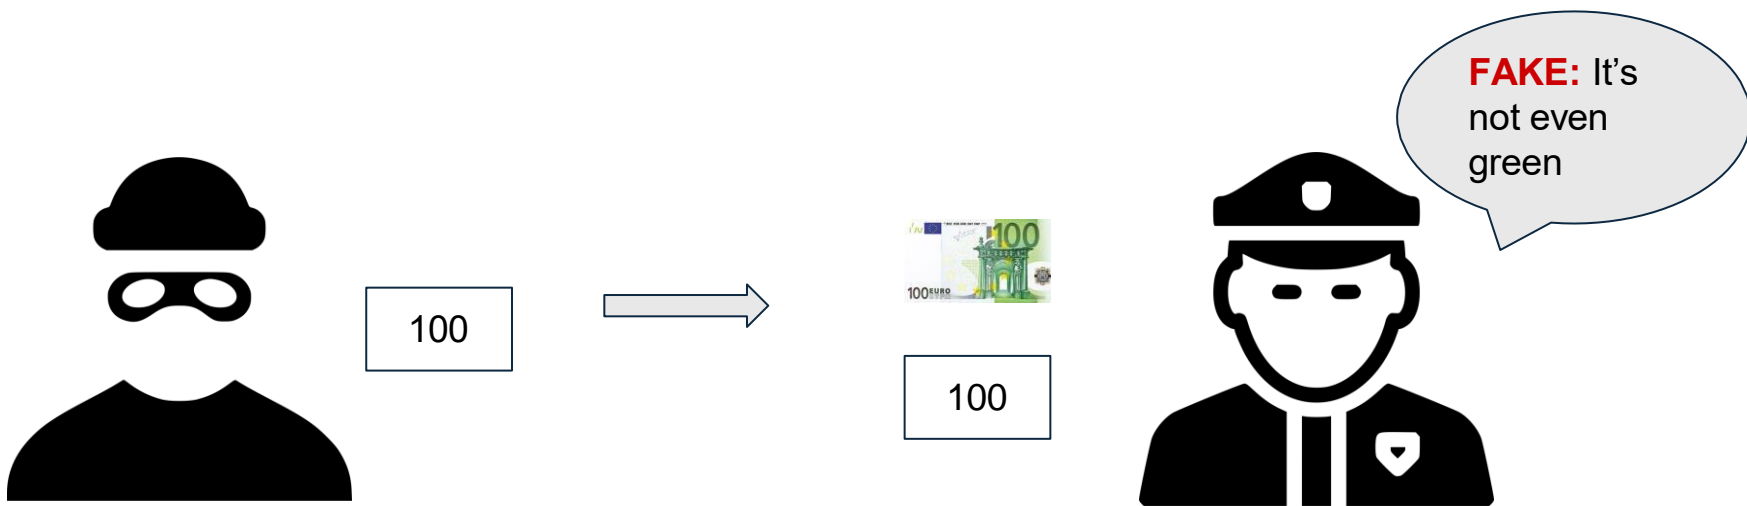
**Generator network:** try to fool the discriminator by generating real-looking images





# Adversarial Training Analogy: is it fake money?

Imagine we have a counterfeiter (**G**) trying to make fake money, and the police (**D**) has to detect whether money is real or fake.



# Adversarial Training Analogy: is it fake money?

Imagine we have a counterfeiter (**G**) trying to make fake money, and the police (**D**) has to detect whether money is real or fake.

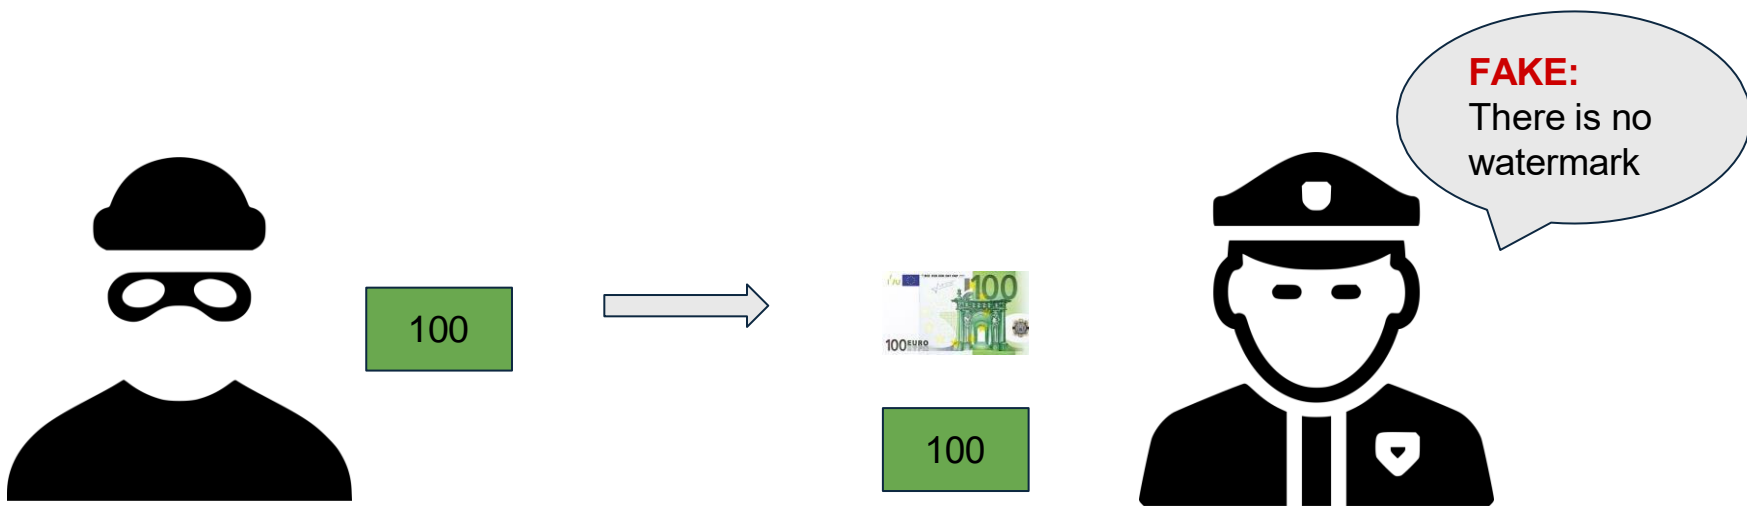


Figure: Santiago Pascual (UPC)



# Adversarial Training Analogy: is it fake money?

Imagine we have a counterfeiter (**G**) trying to make fake money, and the police (**D**) has to detect whether money is real or fake.

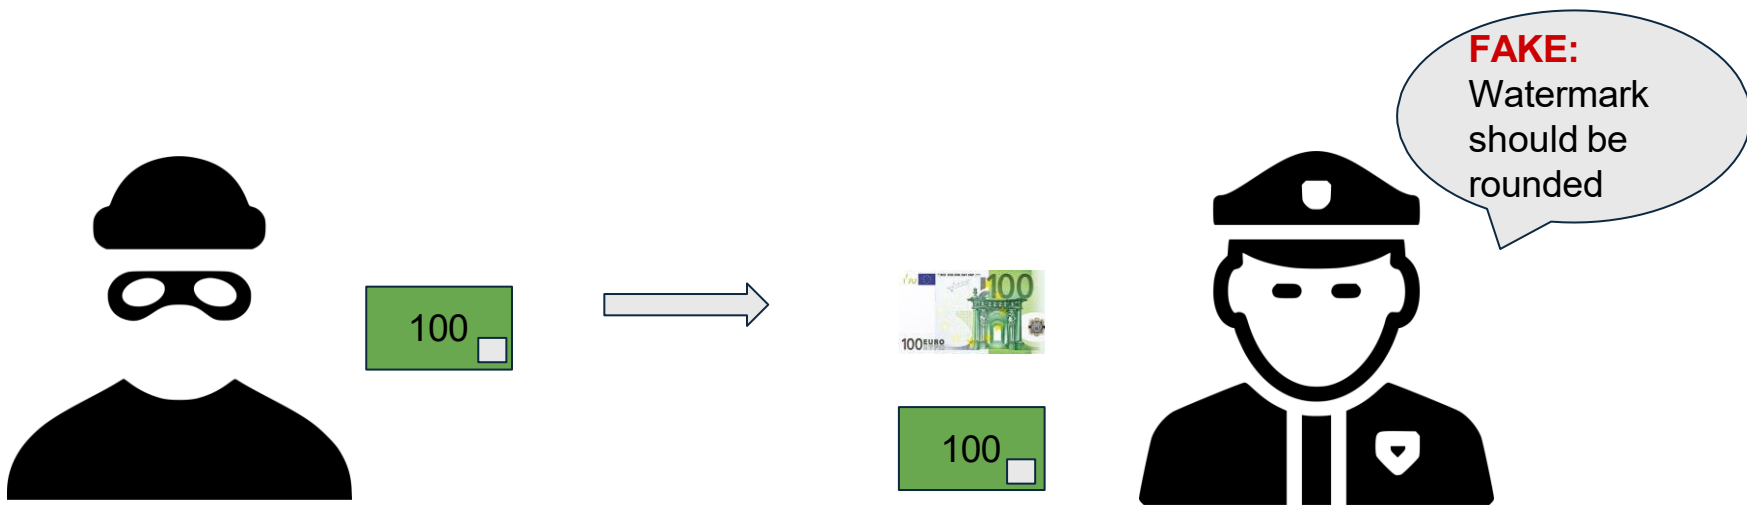


Figure: Santiago Pascual (UPC)

# Adversarial Training Analogy: is it fake money?

Imagine we have a counterfeiter (**G**) trying to make fake money, and the police (**D**) has to detect whether money is real or fake.

After enough iterations, and if the counterfeiter is good enough (in terms of **G** network it means “has enough parameters”), the police should be confused.

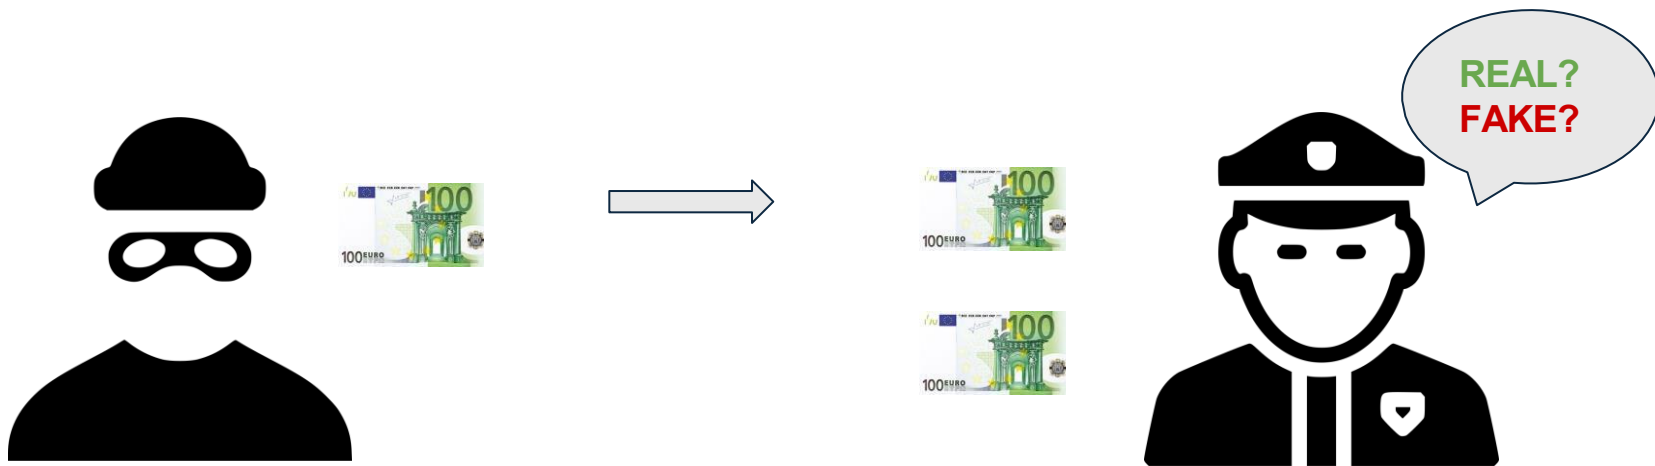


Figure: Santiago Pascual (UPC)

# Adversarial Training

Alternate between training the discriminator and generator

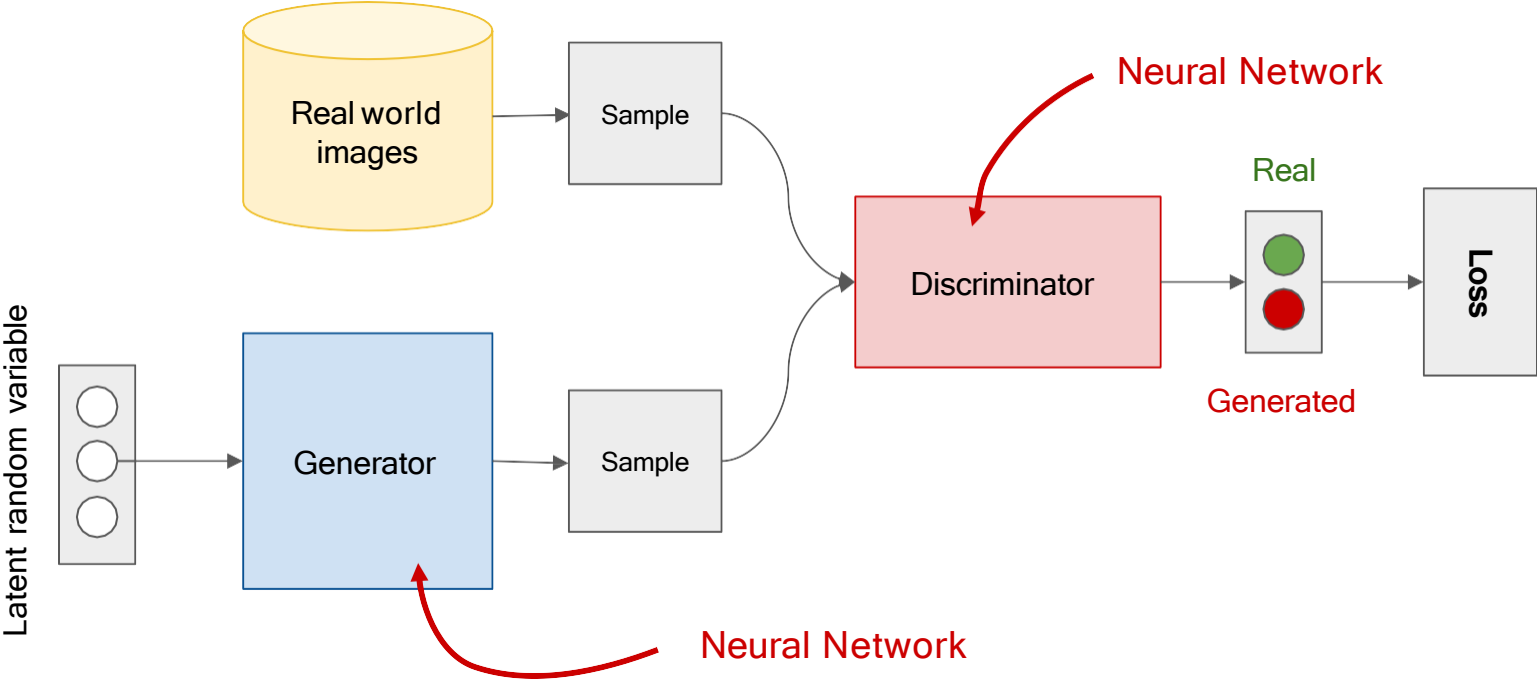


Figure: Kevin McGuinness (DCU)

BCE loss:  $\text{Loss}(\hat{y}, y) = y \log \hat{y} + (1-y) \log(1-\hat{y})$

discriminator: real data  $\hat{y} = D(x) \approx 1$

$y = 1$

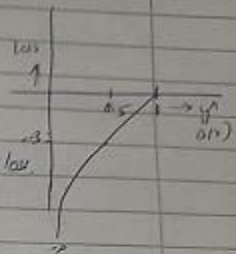
substituting  $y=1$  &  $\hat{y}=D(x)$   
 $1 \cdot \log D(x) + (1-1) \cdot \log(1-D(x))$   
 $\log D(x)$

$$L(\hat{y}, y) = L(1, 1) = \log(1) = 0$$

$$L(0.5, 1) = \log(0.5) = -0.3$$

$$L(0, 1) = \log(0) = -\infty$$

You get maximum value 1 at 0, hence you need to maximize this loss.



Generator Loss:  $1 - D(G(z))$ : Case B is applied  
 $y=0$ ,  $\hat{y} = \log(1 - D(G(z))) \approx 1$  i.e.  $G$  wants to fool  $D$ .

Graph will remain the same, and observe the values, if we want to reach 1 for  $\hat{y}$ , you need to minimize this function. Hence we get a compact representation of our loss.

$$\min_G \max_D \log D(x) + \log(1 - D(G(z)))$$

(B) discriminator: fake data

$\hat{y} = D(G(z))$ ,  $y = 0$

substituting  $\hat{y}$  in BCE eq.

$$0 \cdot \log \hat{y} + (1-0) \log(1 - \hat{y})$$

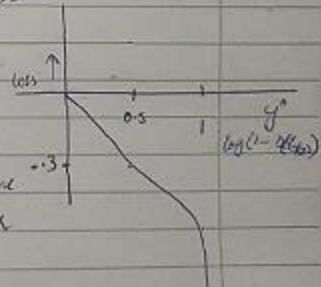
$$\Rightarrow \log(1 - D(G(z)))$$

$$L(\hat{y}, y) = L(0, 0) = \log(1 - 0)$$

$$\Rightarrow \log(1) = 0$$

$$L(0.5, 0) = \log(1 - 0.5) = -0.3$$

$$L(1, 0) = \log(1 - 1) = -\infty$$



For this case you want the  $D$  to classify this as 0, and you're getting this at 0, and increasing  $\hat{y}$  values is giving you lower values therefore we need to maximize this loss for discriminator.

This loss equation is valid for only one data point, but it must be done for the entire training dataset

To represent this, we use the expectation term

# Training GANs: Two Player Game

**Discriminator network:** try to distinguish between real and fake images

**Generator network:** try to fool the discriminator by generating real-looking images

Train jointly in minmax game

Discriminator outputs likelihood in (0,1) of real image

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log \underbrace{D_{\theta_d}(x)}_{\substack{\text{Discriminator output} \\ \text{for real data } x}} + \mathbb{E}_{z \sim p(z)} \log(1 - \underbrace{D_{\theta_d}(G_{\theta_g}(z))}_{\substack{\text{Discriminator output for} \\ \text{generated fake data } G(z)}}) \right]$$

- Discriminator ( $\theta_d$ ) wants to **maximize objective** such that  $D(x)$  is close to 1 (real) and  $D(G(z))$  is close to 0 (fake)
- Generator ( $\theta_g$ ) wants to **minimize objective** such that  $D(G(z))$  is close to 1 (discriminator is fooled into thinking generated  $G(z)$  is real)

# Training GANs: Two Player Game

Minimax objective function:

$$\min_{\theta_g} \max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

Alternate between:

1. **Gradient ascent** on discriminator

$$\max_{\theta_d} \left[ \mathbb{E}_{x \sim p_{data}} \log D_{\theta_d}(x) + \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z))) \right]$$

2. **Gradient descent** on generator

$$\min_{\theta_g} \mathbb{E}_{z \sim p(z)} \log(1 - D_{\theta_d}(G_{\theta_g}(z)))$$

# GANs

---

Training loop:

\* fix the learning of  $G$  \*

Inner loop for  $D$ :

- take  $m$  data samples &  $m$  fake data samples
- update  $\theta_d$  by grad. ascent

$$\frac{\partial}{\partial \theta_d} \frac{1}{m} [\ln[D(x)] + \ln[1 - D(G(z))]]$$

\* fix the learning of  $D$  \*

take  $m$  fake data samples

update  $\theta_g$  by grad. descent

$$\frac{\partial}{\partial \theta_g} \frac{1}{m} [\ln[1 - D(G(z))]]$$

# GANs

---

Binary Crossentropy Function

$$\mathcal{L} = -\sum y \ln \hat{y} + (1-y) \ln(1-\hat{y})$$

when  $y=1$ ,  $\hat{y}=D(x) \Rightarrow \mathcal{L} = \ln [D(x)]$

when  $y=0$ ,  $\hat{y}=D(G(z)) \Rightarrow \mathcal{L} = \ln [1-D(G(z))]$

Adding,  $\mathcal{L} = \ln [D(x)] + \ln [1-D(G(z))]$

This expression is valid for only one data point, but it must be done for the entire training dataset

To represent this, we use the expectation term



# GANs

---

$$E(\mathcal{L}) = E(\ln[D(x)]) + E(\ln[1 - D(G(z))])$$

$$\sum p_{\text{data}}(x) \ln[D(x)] + \sum p_z(z) \ln[1 - D(G(z))]$$

$$V(G, D) = E_{x \sim p_{\text{data}}} [\ln(D(x))] + E_{z \sim p_z} [\ln(1 - D(G(z)))]$$

# Adversarial Training: Discriminator

1. Fix generator weights, draw samples from both real world and generated images
2. Train discriminator to distinguish between real world and generated images

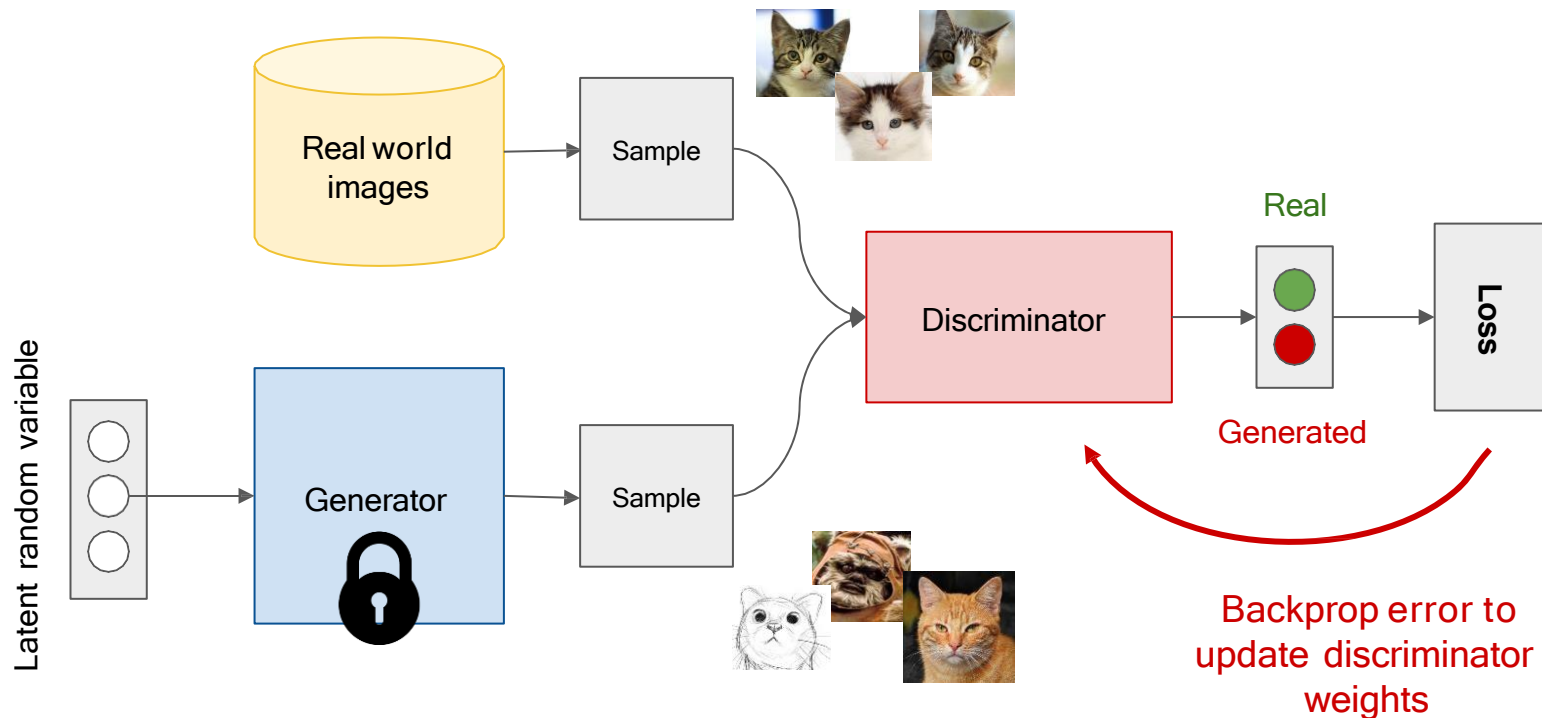


Figure: Kevin McGuinness (DCU)

# Adversarial Training: Discriminator

Consider a binary encoding of “1” (Real) and “0” (Fake). Which of these two values should the discriminator predict when we train the discriminator with a generated image ?

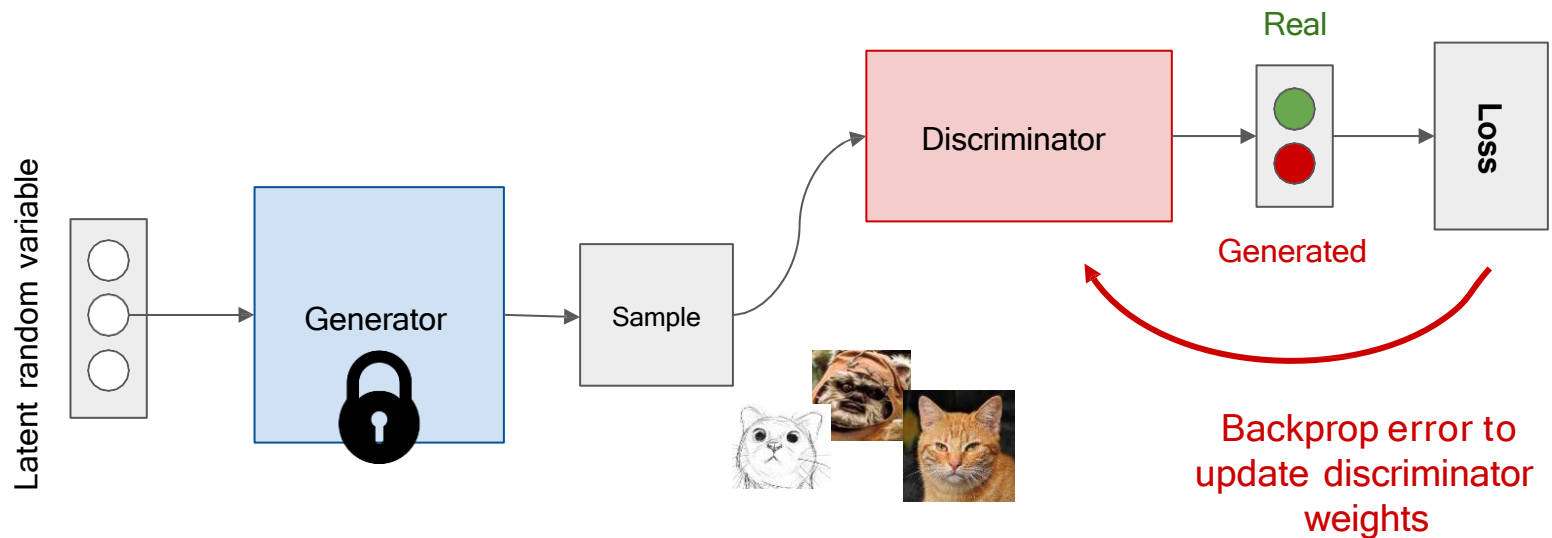
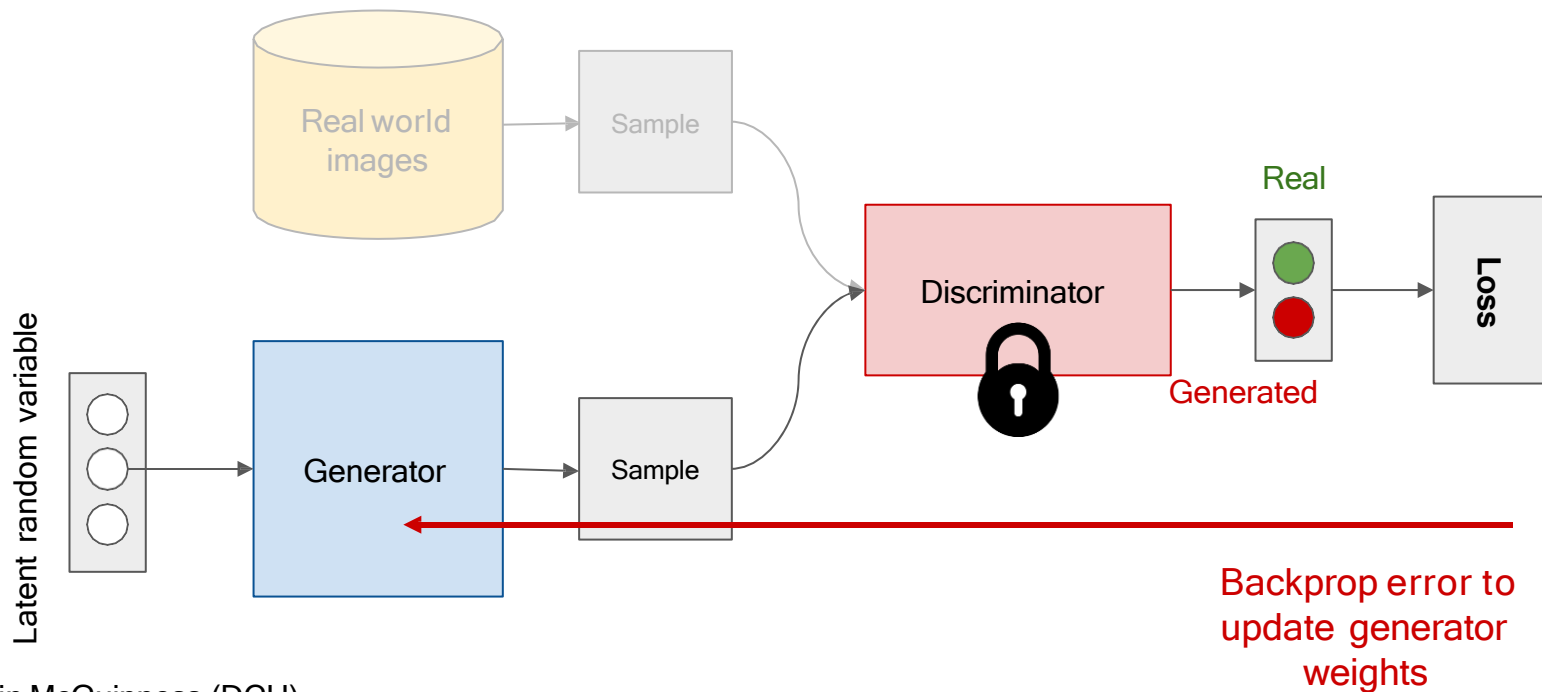


Figure: Kevin McGuinness (DCU)

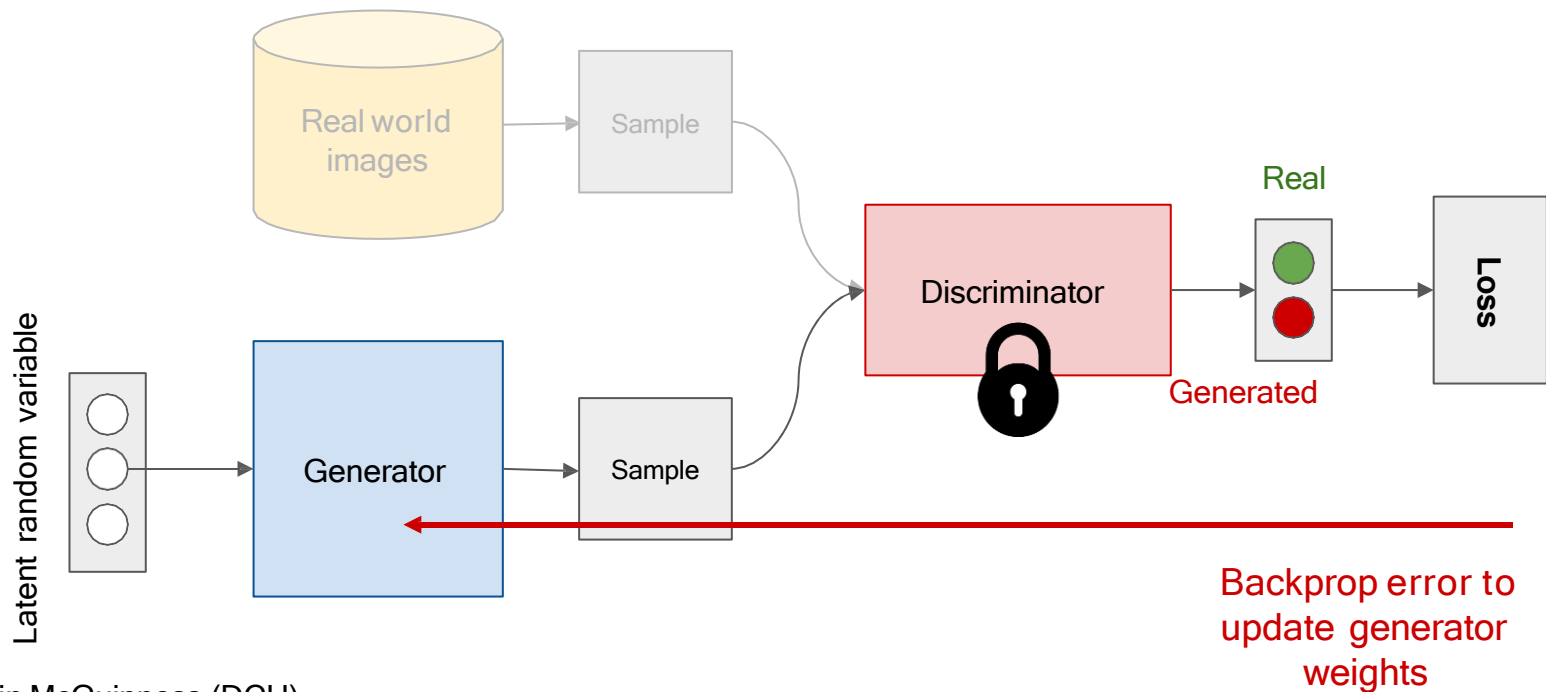
# Adversarial Training: Generator

1. Fix discriminator weights
2. Sample from generator by injecting noise.
3. Backprop error through discriminator to update generator weights ( $D(G(z))$ )



# Adversarial Training: Generator

Consider a binary encoding of “1” (Real) and “0” (Fake). Which of these two values should the discriminator predict when we train the generator with a generated image ?



# Noise Dimensions

## 1. For Images (e.g., MNIST):

- It is common to use a lower-dimensional  $z$  (e.g.,  $z = 64$ ) to generate higher-dimensional outputs (e.g., 784 for MNIST or even  $1024 \times 1024$  in high-res image GANs).
- This is **normal practice** because the generator learns how to expand the latent space into complex high-dimensional image representations.

## 2. For Tabular Data:

1. **Small Dataset or Simple Structure:** Use  $z$  equal to or slightly smaller than the output dimension.
2. **Complex Dataset or High Variability:** Slightly increase  $z$  (e.g.,  $z = 2 \times$  output dimension) to give the generator more flexibility.
3. **Empirical Approach:** The best choice of  $z$  often depends on experimentation—try different sizes of  $z$  and evaluate performance metrics.

# Generation Phase (After Training)

- Only the **generator**  $G$  is used after training. The discriminator is discarded.
- To generate new samples:
  1. **Sample noise**  $z \sim p(z)$  (e.g., from a Gaussian  $\mathcal{N}(0, 1)$  or uniform distribution).
  2. **Feed  $z$  into the trained generator:**

$$x_{\text{generated}} = G(z)$$

3. **Output samples** are obtained based on the architecture (images, tabular data, etc.).

## Controlling Sample Size:

- The number of samples generated is determined by how many  $z$  vectors are sampled.
- Example: If we want **100** generated samples, we sample **100 noise vectors** and pass them to  $G$ .