

Diffusion Models

Diffusion-based

Forward diffusion process (fixed) 😊



Data



x_0




x_1




x_2




x_3



x_4






...



x_T

Noise

Data



Noise

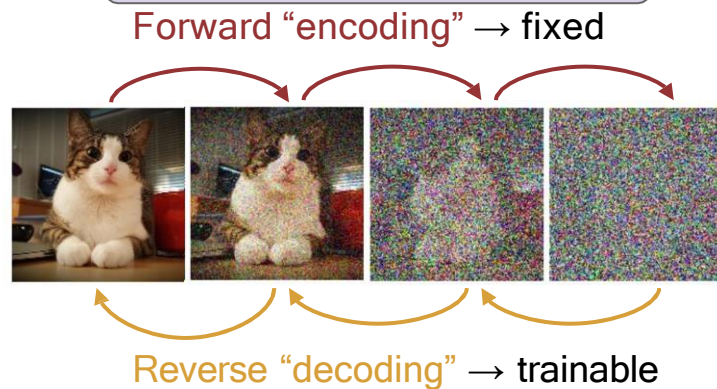




Reverse diffusion process (must learn) 😬



Relation with VAEs

Diffusion models can be considered as a special form of VAEs. However in diffusion models:

- The encoder is fixed: is a **predefined Markovian process**, meaning no learning happens here.
- The latent variables have the same dimension as the data: The noisy versions of the image at different timesteps retain the same size as the original image.
- The decoder is run multiple times in an autoregressive fashion: it runs in **multiple iterations** reconstructing the image from pure noise




Forward Diffusion Process



Reverse Diffusion Process

Forward Diffusion Process



Distribution of the
noised images

Output

Mean μ_t

Variance Σ_t

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

Notations:

t : time step (from 1 to T)

x_0 : a data sampled from the real data distribution $q(x)$ (i.e. $x_0 \sim q(x)$)

β_t : variance schedule ($0 \leq \beta_t \leq 1$, and $\beta_0 = \text{small number}$, $\beta_T = \text{large number}$)

I : identity matrix

The forward diffusion process gradually adds Gaussian noise to the input image x_0 step by step, and there will be T steps in total. the process will produce a sequence of noisy image samples $x_1, ..., x_T$.

When $T \rightarrow \infty$, the final result will become a completely noisy image as if it is sampled from an isotropic Gaussian distribution.

But instead of designing an algorithm to iteratively add noise to the image, we can use a closed-form formula to directly sample a noisy image at a specific time step t .

Ancestral Sampling (One Shot)

Ancestral sampling (One Shot)

$$x_t \sim p(x_t|x_0) = \mathcal{N}\left(x_t; \sqrt{\hat{\alpha}_t} \cdot x_0, (1 - \hat{\alpha}_t)I\right)$$

$$\hat{\alpha}_t = \prod_{i=1}^t \alpha_i$$
$$\alpha_t = 1 - \beta_t$$

- Decompose the image generation (sampling) process in many small steps “denosing”
- Small step help to remove noise easily and accurately in backward process
- **Constant Variance:**
- Variance is kept constant, and mean is changed at each time stamp
- Changing both will collapse the structure of the image, and we need to reconstruct from noise.

How we get $\sqrt{1 - \beta_t}$

Assume we start from a random variable x_0 that follows a normal distribution

$$x_0 \sim \mathcal{N}(0, 1)$$

Noise Addition Formula in Diffusion Models

Step 1: Adding Noise

In the first step of the diffusion process ($t = 1$), we scale the original image x_0 by a factor α and add Gaussian noise:

$$x_1 = \alpha x_0 + \sqrt{\beta_1} \epsilon_1$$

where:

- $\epsilon_1 \sim \mathcal{N}(0, 1)$ is a random variable representing Gaussian noise with mean 0 and variance 1.

Calculating Variance

To compute the variance of x_1 , we use the property that the variance of the sum of independent random variables is the sum of their variances:

$$\text{Var}(x_1) = \text{Var}(\alpha x_0 + \sqrt{\beta_1} \epsilon_1)$$

Expanding this:

$$\text{Var}(x_1) = \alpha^2 \text{Var}(x_0) + \beta_1 \text{Var}(\epsilon_1)$$

Since:

- x_0 has unit variance: $\text{Var}(x_0) = 1$,
- ϵ_1 also has unit variance: $\text{Var}(\epsilon_1) = 1$,

we get:

$$\text{Var}(x_1) = \alpha^2 + \beta_1$$

Maintaining Unit Variance

To ensure that variance remains equal to 1 at every step, we enforce:

$$\alpha^2 + \beta_1 = 1$$

which gives:

$$\alpha_1 = \sqrt{1 - \beta_1}$$

Thus, the noise-added formulation can be rewritten as:

$$x_1 = \sqrt{1 - \beta_1} x_0 + \sqrt{\beta_1} \epsilon_1$$

$$x_1 = \alpha x_0 + \sqrt{\beta_1} \epsilon_1$$

This ensures that the variance remains normalized throughout the diffusion process.

Sample

$$x_t = \sqrt{1 - \beta_t} x_{t-1} + \sqrt{\beta_t} \epsilon_{t-1}$$

Distribution

$$q(x_t | x_{t-1}) = \mathcal{N}(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$




$$q(x_t | x_{t-1}) = N(x_t; \sqrt{1 - \beta_t} x_{t-1}, \beta_t I)$$

$$\rightarrow q(x_t | x_0) = N(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) I)$$

$T=200$

Closed form forward process



$$q(x_t | x_0) = N(x_t; \sqrt{\bar{\alpha}_t} x_0, (1 - \bar{\alpha}_t) I)$$

Closed-Form Formula

The closed-form sampling formula can be derived using the Reparameterization Trick.

If $z \sim \mathcal{N}(\mu, \sigma^2)$ then

$$z = \mu + \sigma \varepsilon \quad \text{where } \varepsilon \sim \mathcal{N}(0, 1)$$

Reparameterization trick

With this trick, we can express the sampled image \mathbf{x}_t as follows:

$$\mathbf{x}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1} + \sqrt{\beta_t} \varepsilon_{t-1}$$

Express \mathbf{x}_t using the reparameterization trick

Then we can expand it recursively to get the closed-form formula:

$$\begin{aligned}
 x_t &= \sqrt{1 - \beta_t} x_{t-1} + \sqrt{\beta_t} \varepsilon_{t-1} & \varepsilon_0, \dots, \varepsilon_{t-2}, \varepsilon_{t-1} &\sim \mathcal{N}(0, I) \\
 &= \sqrt{\alpha_t} \boxed{x_{t-1}} + \sqrt{1 - \alpha_t} \varepsilon_{t-1} & \bar{\varepsilon}_0, \dots, \bar{\varepsilon}_{t-2}, \bar{\varepsilon}_{t-1} &\sim \mathcal{N}(0, I) \\
 &= \sqrt{\alpha_t} \left(\sqrt{\alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_{t-1}} \varepsilon_{t-2} \right) + \sqrt{1 - \alpha_t} \varepsilon_{t-1} & \varepsilon &\sim \mathcal{N}(0, I) \\
 &= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \boxed{\sqrt{\alpha_t (1 - \alpha_{t-1})} \varepsilon_{t-2}} + \sqrt{1 - \alpha_t} \varepsilon_{t-1} & \alpha_t &= 1 - \beta_t \\
 &= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \boxed{\sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\varepsilon}_{t-2}} & \bar{\alpha}_t &= \prod_{i=1}^t \alpha_i \\
 &\vdots & & \\
 &= \sqrt{\alpha_t \alpha_{t-1} \dots \alpha_1} x_0 + \sqrt{1 - \alpha_t \alpha_{t-1} \dots \alpha_1} \varepsilon \\
 &= \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \varepsilon
 \end{aligned}$$

How?

But how do we jump from line 4 to line 5?

$$\sqrt{\alpha_t(1 - \alpha_{t-1})} \varepsilon_{t-2} + \sqrt{1 - \alpha_t} \varepsilon_{t-1}$$
$$\sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\varepsilon}_{t-2} \leftarrow \text{How?}$$

$$x_t = \sqrt{1 - \beta_t} x_{t-1} + \sqrt{\beta_t} \varepsilon_{t-1}$$

$$\varepsilon_0, \dots, \varepsilon_{t-2}, \varepsilon_{t-1} \sim \mathcal{N}(0, I)$$

$$= \sqrt{\alpha_t} \boxed{x_{t-1}} + \sqrt{1 - \alpha_t} \varepsilon_{t-1}$$

$$\bar{\varepsilon}_0, \dots, \bar{\varepsilon}_{t-2}, \bar{\varepsilon}_{t-1} \sim \mathcal{N}(0, I)$$

$$\varepsilon \sim \mathcal{N}(0, I)$$

$$= \sqrt{\alpha_t} \left(\sqrt{\alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_{t-1}} \varepsilon_{t-2} \right) + \sqrt{1 - \alpha_t} \varepsilon_{t-1}$$

$$\alpha_t = 1 - \beta_t$$

$$= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \boxed{\sqrt{\alpha_t (1 - \alpha_{t-1})} \varepsilon_{t-2}} + \boxed{\sqrt{1 - \alpha_t} \varepsilon_{t-1}}$$

$$\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

$X + Y$

Reparameterization Trick

Underlying Normal Distribution

$$0 + \sqrt{\alpha_t (1 - \alpha_{t-1})} \varepsilon_{t-2} \quad \text{-----} \quad X \sim \mathcal{N}(0, \alpha_t (1 - \alpha_{t-1}) I)$$

$$0 + \sqrt{1 - \alpha_t} \varepsilon_{t-1} \quad \text{-----} \quad Y \sim \mathcal{N}(0, (1 - \alpha_t) I)$$

Recall:

$$\text{If } X \sim \mathcal{N}(\mu_X, \sigma_X^2) \quad Y \sim \mathcal{N}(\mu_Y, \sigma_Y^2)$$

$$Z = X + Y$$

$$\text{Then } Z \sim \mathcal{N}(\mu_X + \mu_Y, \sigma_X^2 + \sigma_Y^2)$$

$$Z \sim \mathcal{N}(0, \sigma_X^2 + \sigma_Y^2)$$

$$Z \sim \mathcal{N}(0, (1 - \alpha_t \alpha_{t-1}) I)$$



Reparameterization Trick

$$0 + \sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\epsilon}_{t-2}$$

$$\begin{aligned} \sigma_X^2 + \sigma_Y^2 &= \alpha_t(1 - \alpha_{t-1}) + (1 - \alpha_t) \\ &= \cancel{\alpha_t} - \alpha_t \alpha_{t-1} + 1 - \cancel{\alpha_t} \\ &= 1 - \alpha_t \alpha_{t-1} \end{aligned}$$

$$= \sqrt{\alpha_t \alpha_{t-1}} x_{t-2} + \sqrt{1 - \alpha_t \alpha_{t-1}} \bar{\epsilon}_{t-2}$$

\vdots

$$= \sqrt{\alpha_t \alpha_{t-1} \dots \alpha_1} x_0 + \sqrt{1 - \alpha_t \alpha_{t-1} \dots \alpha_1} \epsilon$$

$$= \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

Detailed derivation from line 4 to line 5

Repeating these steps will give us the following formula which depends only on the input image x_0 :

$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

The closed-form formula

Now we can directly sample x_t at any time step using this formula, and this makes the forward process much faster.

where:

To extract ϵ_t (the noise at step t) from a given noisy sample x_t :

$$\epsilon_t = \frac{x_t - \sqrt{\bar{\alpha}_t} x_0}{\sqrt{1 - \bar{\alpha}_t}}$$

- x_0 is the original clean image,
- x_t is the noisy image at timestep t ,
- $\bar{\alpha}_t$ is the accumulated product of all previous noise scales,
- $\epsilon \sim \mathcal{N}(0, I)$ is standard Gaussian noise.

How Do We Compute $\bar{\alpha}_t$?

$\bar{\alpha}_t$ is the cumulative product of a sequence of predefined noise scales α_t . These α_t values are usually determined by a noise schedule, which is a function that controls how noise is added across timesteps. The most common schedules include:

- **Linear schedule:** β_t increases linearly from a small to a large value.
- **Cosine schedule:** Inspired by cosine functions for smoother noise distribution.

Once we have a noise schedule, we compute:

$$\alpha_t = 1 - \beta_t$$

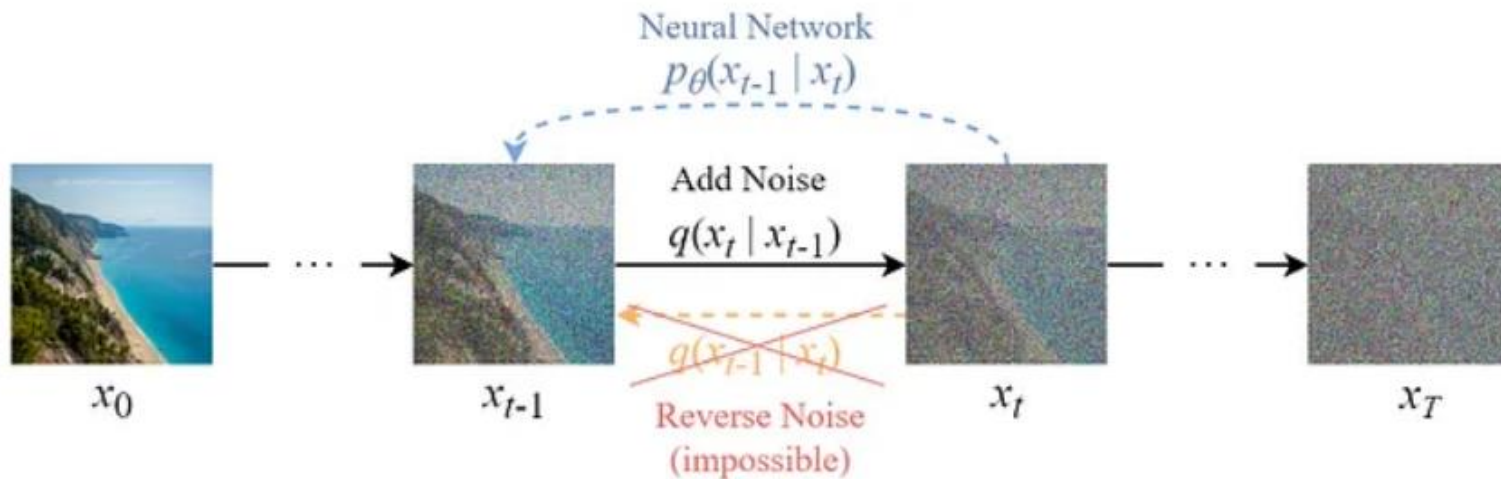
$$\bar{\alpha}_t = \prod_{s=1}^t \alpha_s$$

the noise schedule provides all the α_t values upfront, allowing you to calculate any $\bar{\alpha}_t$ directly without needing to step through the diffusion process.

This cumulative product accumulates the effect of all past timesteps.

If you only have the original image x_0 , you must define a noise schedule first (e.g., a linear or cosine schedule). Then, compute $\bar{\alpha}_t$ as the cumulative product of α_t up to timestep t . This allows you to generate noisy samples in one step using the equation.

Reverse Diffusion Process



Target Distribution

$$q(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \tilde{\mu}_t(x_t, x_0), \tilde{\beta}_t I)$$

Approximated Distribution

$$p_\theta(x_{t-1}|x_t) = \mathcal{N}(x_{t-1}; \mu_\theta(x_t, t), \Sigma_\theta(x_t, t))$$

Learnable parameters
(Neural Network)

Unlike the forward process, we cannot use $q(x_{t-1}|x_t)$ to reverse the noise since it is intractable (uncomputable).

Thus we need to train a neural network $p_\theta(x_{t-1}|x_t)$ to approximate $q(x_{t-1}|x_t)$. The approximation $p_\theta(x_{t-1}|x_t)$ follows a normal distribution and its mean and variance are set as follows:

$$\begin{cases} \mu_\theta(x_t, t) &:= \tilde{\mu}_t(x_t, x_0) \\ \Sigma_\theta(x_t, t) &:= \tilde{\beta}_t I \end{cases}$$

mean and variance of p_θ

Loss Function

We can define our loss as a Negative Log-Likelihood:

$$\text{Loss} = -\log(p_{\theta}(x_0))$$


Depends on x_1, x_2, \dots, x_T
Therefore it is intractable!

Negative log-likelihood

This setup is very similar to the one in VAE. instead of optimizing the intractable loss function itself, we can optimize the Variational Lower Bound.

To approximate the target denoising step q , we only need to approximate its mean using a neural network. So we set the approximated mean μ_θ to be in the same form as the target mean $\tilde{\mu}_t$ (with a learnable neural network ϵ_θ):

$$\tilde{\mu}_t(x_t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_t \right)$$


$$\text{Set } \mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right)$$

Approximated mean

Simplified Loss

So the final simplified training objective is as follows:

$$x_t = \sqrt{\bar{a}_t} x_0 + \sqrt{1 - \bar{a}_t} \epsilon$$

$$L_{\text{simple}} = \mathbb{E}_{t, x_0, \epsilon} \left[\|\epsilon - \epsilon_{\theta}(x_t, t)\|^2 \right]$$

Simplified training objective

Loss Function: The UNet is trained to minimize the difference between its predicted noise $\hat{\epsilon}_{\theta}$ and the actual noise ϵ (which was added during forward diffusion):

$$L = \mathbb{E}_{x_0, \epsilon, t} \left[\|\epsilon - \hat{\epsilon}_{\theta}(x_t, t)\|^2 \right]$$

Reverse inference/sampling part

1. Input Noisy Image x_t

- Given an image x_t at time t , we want to move backward in the diffusion process to reconstruct x_{t-1} .

2. UNet Predicts Noise $\epsilon_\theta(x_t, t)$

- The model (usually a UNet) takes in x_t and estimates the noise component $\epsilon_\theta(x_t, t)$ that was added in the forward process.

3. Estimate the Original Image x_0 (Intermediate Step)

- Using the closed-form forward process equation, we estimate the original clean image x_0 :

$$\hat{x}_0 = \frac{x_t - \sqrt{1 - \bar{\alpha}_t} \epsilon_\theta(x_t, t)}{\sqrt{\bar{\alpha}_t}} \qquad x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

- This equation removes the predicted noise from x_t and estimates what the original image might have been before the noise was added.

4. Compute the Mean $\mu_\theta(x_t, t)$ for Transition to x_{t-1}

- Now that we have an estimate of x_0 , we compute the mean for the next step:

$$\mu_\theta(x_t, t) = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right)$$

- This tells us where x_{t-1} should be centered.

5. Sample x_{t-1} Using a Stochastic Term

- Instead of directly setting $x_{t-1} = \mu_\theta(x_t, t)$, we add **Gaussian noise** to make the process probabilistic:

$$x_{t-1} = \mu_\theta(x_t, t) + \sigma_t z, \quad \text{where } z \sim \mathcal{N}(0, I)$$

- σ_t is a predefined noise scale (depends on the variance schedule).
- z is standard Gaussian noise, ensuring some randomness in the process.

6. Repeat Until x_0 is Reached

- This step-by-step denoising process continues until we reach x_0 , the final clean image.

Why do we add noise when denoising?

- The forward diffusion process was **stochastic** (random noise added at each step).
- To reverse it, we also need a **stochastic process** rather than a deterministic one.
- The added noise helps maintain diversity in generated images.

- Creativity and Variation:** It allows the model to generate creative and varied outputs, even when starting from the same latent noise. Imagine generating different variations of a "cat" image – different poses, colors, fur patterns, etc.
- Avoiding Overfitting:** Stochasticity helps prevent the model from memorizing specific training examples and encourages it to learn the underlying distribution of the data.
- Exploring the Latent Space:** The added noise allows the reverse process to explore different regions of the latent space, leading to a wider range of generated samples.

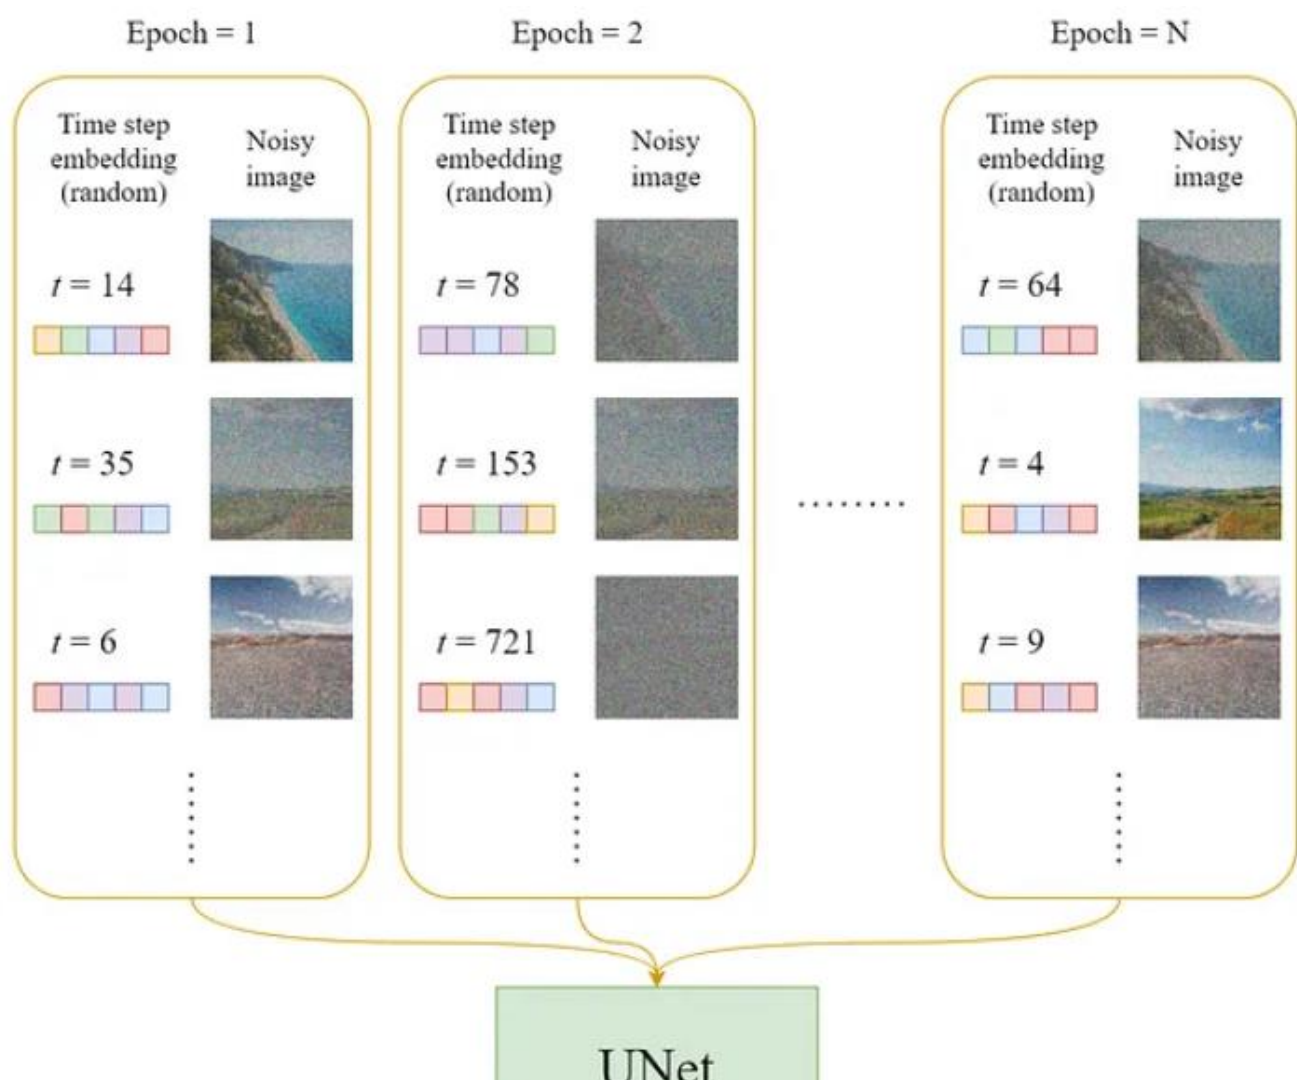
The U-Net Model

Dataset

In each epoch:

1. A random time step t will be selected for each training sample (image).
2. Apply the Gaussian noise (corresponding to t) to each image.
3. Convert the time steps to embeddings (vectors).

Why do we pass random time steps in Uet during training?



Algorithm 1 Training

1: **repeat**

2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$

3: $t \sim \text{Uniform}(\{1, \dots, T\})$

4: $\epsilon \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$

5: Take gradient descent step on

$$\nabla_{\theta} \|\epsilon - \epsilon_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon, t)\|^2$$

6: **until** converged


During inference, the forward process is **skipped**—we directly sample noise and denoise step by step.

For each training step:

1. Randomly select a time step & encode it



2. Add noise to image




$$\epsilon \sim \mathcal{N}(0, 1)$$

$$\alpha_t = 1 - \beta_t$$

$$\bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

3. Train the UNet



Algorithm 2 Sampling

- 1: $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
- 2: **for** $t = T, \dots, 1$ **do**
- 3: $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ if $t > 1$, else $\mathbf{z} = \mathbf{0}$
- 4: $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(\mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\alpha_t}} \epsilon_\theta(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$
- 5: **end for**
- 6: **return** \mathbf{x}_0

1. Sample a Gaussian noise




$$\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$

E.g. $T = 1000$

$$\mathbf{x}_{1000} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$$



2. Iteratively denoise the image



3. Output the denoised image

Denoised image \mathbf{x}_0




Summary


Here are some main takeaways from this article:

- The Diffusion model is divided into two parts: forward diffusion and reverse diffusion.
- The forward diffusion can be done using the closed-form formula.
- The backward diffusion can be done using a trained neural network.
- To approximate the desired denoising step q , we just need to approximate the noise ϵ_t using a neural network $\epsilon\theta$.
- Training on the simplified loss function yields better sample quality.



UNET model (architecture)



Generative AI Timeline



Recommended books



Interview of David Foster for Machine Learning Street Talk (2023)